



SEARCH



SIGRED – RESOLVING YOUR WAY INTO DOMAIN ADMIN: EXPLOITING A 17 YEAR-OLD BUG IN WINDOWS DNS SERVERS

...

July 14, 2020

Research by: Sagi Tzadik

Introduction

DNS, which is often described as the “phonebook of the internet”, is a network protocol for translating human-friendly computer hostnames into IP addresses. Because it is such a core component of the internet, there are many solutions and implementations of DNS servers out there, but only a few are extensively used.

“Windows DNS Server” is the Microsoft implementation and is an essential part of and a requirement for a Windows Domain environment.

SIGRed (CVE-2020-1350) is a wormable, critical vulnerability (CVSS base score of 10.0) in the Windows DNS server that affects Windows Server versions 2003 to 2019, and can be triggered by a malicious DNS response. As the service is running in elevated privileges (SYSTEM), if exploited successfully, an attacker is granted Domain Administrator rights, effectively compromising the entire corporate infrastructure.

<https://youtu.be/PULMmhD5it8>

Motivation

Our main goal was to find a vulnerability that would let an attacker compromise a Windows Domain environment, preferably unauthenticated. There is a lot of related research by various independent security researchers as well as those sponsored by nation-states. Most of the published and publicly available materials and exploits focus on Microsoft’s implementation of SMB ([EternalBlue](#)) and RDP ([BlueKeep](#)) protocols, as these targets affect both servers and endpoints. To obtain Domain Admin privileges, a straightforward approach is to directly exploit the Domain Controller. Therefore, we decided to focus our research on a less publicly explored attack surface that exists primarily on Windows Server and Domain Controllers. Enter WinDNS.

Windows DNS Overview

“Domain Name System (DNS) is one of the industry-standard suite of protocols that comprise TCP/IP, and together the DNS Client and DNS Server provide computer name-to-IP address mapping name resolution services to computers and users.” – [Microsoft](#).

DNS primarily uses the User Datagram Protocol (UDP) on port 53 to serve requests. DNS queries consist of a single UDP request from the client followed by a single UDP reply from the server.

In addition to translating names to IP addresses, DNS serves other purposes as well. For example, mail transfer agents use DNS to find the best mail server to deliver e-mail: An MX record provides a mapping between a domain and a mail exchanger, which can provide an additional layer of fault tolerance and load distribution. A list of available DNS record types and their corresponding purposes can be found on [Wikipedia](#).

But the point of this blog post is not to present a lengthy discourse on DNS features and history, so we encourage you to read more about DNS [here](#).

What you need to know:

- DNS operates over UDP/TCP port 53.
- A single DNS message (response / query) is limited to 512 bytes in UDP and 65,535 bytes in TCP.
- DNS is hierachal and decentralized in nature. This means when a DNS server doesn't know the answer to a query it receives, the query is forwarded to a DNS server above it in the hierarchy. At the top of the hierarchy there are 13 root DNS servers worldwide.

In Windows, the DNS client and DNS server are implemented in two different modules:

- **DNS Client** – `dnsapi.dll` is responsible for DNS resolving.
- **DNS Server** – `dns.exe` is responsible for answering DNS queries on Windows Server, in which the DNS role is installed.

Our research is centered around the `dns.exe` module.

Preparing the Environment

There are two main scenarios for our attack surface:

1. A bug in the way the DNS server parses an incoming query.
2. A bug in the way the DNS server parses a response (answer) for a forwarded query.

As DNS queries do not have a complex structure, there is a lower chance of finding parsing issues in the first scenario, so we decided to target functions that parse incoming responses for forwarded queries.

As mentioned previously, a forwarded query is the utilization of the DNS architecture to be able to forward queries it does not know the answer to – to the DNS server above it in the hierarchy.

However, most environments configure their forwarders to well-known, respectable DNS servers such as **8.8.8.8** (Google) or **1.1.1.1** (Cloudflare), or at the very least a server that is not under the attacker's control.

This means that even if we find an issue in the parsing of DNS responses, we need to establish a Man-in-the-Middle to exploit it. Obviously, that's not good enough.

NS Records to the Rescue

NS stands for 'name server' and this record indicates which DNS server is the authority for that domain (which server contains the actual DNS records). The NS record is usually in charge of resolving the subdomains of a given domain. A domain often has multiple NS records which can indicate primary and backup name servers for that domain.

To have the target Windows DNS Server parse responses from our malicious DNS NameServer, we do the following:

1. Configure our domain's (**deadbeef.fun**) NS Records to point at our malicious DNS Server (**ns1.41414141.club**).
2. Query the victim Windows DNS Server for NS Records of **deadbeef.fun**.
3. The victim DNS, not yet knowing the answer for this query, forwards the query to the DNS server above it (**8.8.8.8**).

4. The authoritative server (8.8.8.8) knows the answer, and responds that the NameServer of **deadbeef.fun** is **ns1.41414141.club**.

5. The victim Windows DNS Server processes and caches this response.

6. The next time we query for a subdomain of **deadbeef.fun**, the target Windows DNS Server will also query **ns1.41414141.club** for its response, as it is the NameServer for this domain.

Source	Destination	Protocol	Length	Info
192.168.147.1	192.168.147.149	DNS	185	Standard query 0xf5d9 A resolveme.deadbeef.fun OPT
192.168.147.149	8.8.8	DNS	93	Standard query 0xle45 A resolveme.deadbeef.fun OPT
8.8.8	192.168.147.149	DNS	93	Standard query response 0xle45 Server failure A resolveme.deadbeef.fun OPT
192.168.147.149	192.168.147.1	DNS	93	Standard query response 0xf5d9 Server failure A resolveme.deadbeef.fun OPT
192.168.147.1	192.168.147.149	DNS	95	Standard query 0x1535 NS deadbeef.fun OPT
192.168.147.149	8.8.8	DNS	72	Standard query 0x7e07 NS deadbeef.fun
8.8.8	192.168.147.149	DNS	121	Standard query response 0x7e07 NS deadbeef.fun NS ns3.41414141.club NS ns4.41414141.club
192.168.147.149	8.8.8	DNS	88	Standard query 0x356b A ns3.41414141.club OPT
8.8.8	192.168.147.149	DNS	104	Standard query response 0x356b A ns3.41414141.club A 35.238.100.241 OPT
192.168.147.149	8.8.8	DNS	88	Standard query 0xec49 AAAA ns3.41414141.club OPT
8.8.8	192.168.147.149	DNS	161	Standard query response 0xec49 AAAA ns3.41414141.club SOA dns1.registrar-servers.com OPT
192.168.147.149	8.8.8	DNS	88	Standard query 0x1551 A ns4.41414141.club OPT
8.8.8	192.168.147.149	DNS	104	Standard query response 0x1551 A ns4.41414141.club A 35.238.100.241 OPT
192.168.147.149	8.8.8	DNS	88	Standard query 0xdfa9 AAAA ns4.41414141.club OPT
8.8.8	192.168.147.149	DNS	161	Standard query response 0xdfa9 AAAA ns4.41414141.club SOA dns1.registrar-servers.com OPT
192.168.147.149	192.168.147.1	DNS	164	Standard query response 0x1535 NS deadbeef.fun NS ns3.41414141.club NS ns4.41414141.club A 35.238.100.243 A 35.238.100.241 OPT
192.168.147.1	192.168.147.149	DNS	185	Standard query 0xd9c4 A resolveme.deadbeef.fun OPT
192.168.147.149	8.8.8	DNS	93	Standard query 0x1597 A resolveme.deadbeef.fun OPT
192.168.147.149	35.238.100.241	DNS	93	Standard query 0x5d01 A resolveme.deadbeef.fun OPT

Figure 1: Packet capture of the victim DNS server querying our malicious server.

The Vulnerability – CVE-2020-1350

Function: `dns.exe!SigWireRead`

Vulnerability Type: Integer Overflow leading to Heap-Based Buffer Overflow

`dns.exe` implements a parsing function for every supported response type.

```
0fb7571e    movzx edx, word [rdi + 0x1e]
488d0def430b00  lea rcx, [rbx + 0x10]
e8266d0700  call pdb.RRWireReadTable
4c8bf0    mov r14, rax
4885c0    test rax, rax
752e      jne 0x7ff62e21f25c
```

Figure 2: `Wire_CreateRecordFromWire`: `RRWireReadTable` is passed to `RR_DispatchFunctionForType` to determine the handling function.

```
;-- pdb.RRWireReadTable:  
pdb.RRWireReadTable .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0x8 .qword 0x000000001400ae330 ; pdb.AWireRead  
pdb.RRWireReadTable+0x10 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x18 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x20 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x28 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x30 .qword 0x000000001400ae510 ; pdb.SoaWireRead  
pdb.RRWireReadTable+0x38 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x40 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x48 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x50 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0x58 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0x60 .qword 0x000000001400ae3d0 ; pdb.PtrWireRead  
pdb.RRWireReadTable+0x68 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0x70 .qword 0x000000001400ae740 ; pdb.MinfoWireRead  
pdb.RRWireReadTable+0x78 .qword 0x000000001400ae460 ; pdb.MxWireRead  
pdb.RRWireReadTable+0x80 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0x88 .qword 0x000000001400ae740 ; pdb.MinfoWireRead  
pdb.RRWireReadTable+0x90 .qword 0x000000001400ae460 ; pdb.MxWireRead  
pdb.RRWireReadTable+0x98 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0xa0 .qword 0x000000001400ae370 ; pdb.CopyWireRead  
pdb.RRWireReadTable+0xa8 .qword 0x000000001400ae460 ; pdb.MxWireRead
```

Figure 3: `RRWireReadTable` and some of its supported response types.

One of the supported response types is for a SIG query. According to Wikipedia, a SIG query is the “*signature record used in SIG(0) (RFC 2931) and TKEY (RFC 2930). RFC 3755 designated RRSIG as the replacement for SIG for use within DNSSEC.*”

Let's examine the disassembly generated by Cutter for `dns.exe!SigWireRead` – the handler function for the SIG response type:

```
8364242000 and dword [var_148h], 0  
488d4c2430 lea rcx, [var_138h]  
4c8bcf mov r9, rdi  
e87fdafcff call pdb.Name_PacketNameToCountNameEx  
488be8 mov rbp, rax  
4885c0 test rax, rax  
74df je 0x7ff62e21ed14  
  
0fb64c2430 movzx ecx, byte [var_138h]  
482bf8 sub rdi, rax  
33d2 xor edx, edx  
6683c114 add cx, 0x14 ; 20  
4533c0 xor r8d, r8d  
6603cf add cx, di  
e85e710700 call pdb.RR_AllocateEx  
488bf0 mov rsi, rax  
4885c0 test rax, rax  
74be je 0x7ff62e21ed14
```

Figure 4: Disassembly of `dns.exe!SigWireRead` as seen in Cutter.

The first parameter that is passed to `RR_AllocateEx` (the function responsible for allocating memory for the Resource Record) is calculated by the following formula:

`[Name_PacketNameToCountNameEx result] + [0x14] + [The Signature field's length (rdi - rax)]`

The signature field size may vary as it is the primary payload of the SIG response.

Figure 5: The structure of SIG Resource Record according to [RFC 2535](#).

As you can see in the image below, `RR_AllocateEx` expects its parameters to be passed in **16bit registers** as it only uses the `dx` part of `rdx` and part of `rcx`.

This means that if we can make the above formula output a result bigger than 65,535 bytes (the maximum value for a 16 bit integer), we have an integer overflow that leads to a much smaller allocation than expected, which hopefully leads to a heap based buffer overwrite.

Figure 6: `RR_AllocateEx` converts its parameters to their 16bit value.

Conveniently enough, this allocated memory address is then passed as a destination buffer for `memcpy`, leading to a Heap-Based buffer overflow.

Figure 7: The allocated buffer from `RR_AllocateEx` is passed into `memcpy`.

To summarize, by sending a DNS response that contains a large (bigger than 64KB) SIG record, we can cause a controlled heap-based buffer overflow of roughly 64KB over a small allocated buffer.

Triggering the Vulnerability

Now that we're able to get the victim DNS server to query our DNS server for various questions, we have effectively turned it into a client. We can make the victim DNS server ask our malicious DNS server specific types of queries, and respectively answer with matching malicious responses.

We thought that all we needed to trigger this vulnerability was to make the victim DNS server query us for a SIG record, and answer it a SIG response with a lengthy signature (length \geq 64KB). We were disappointed to find that DNS over UDP has a size limit of 512 bytes (or 4,096 bytes if the server supports EDNS0). In any case, that is not enough to trigger the vulnerability.

But what happens if there's a legitimate reason for a server to send a response larger than 4,096 bytes? For example, a lengthy TXT response or a hostname that can be resolved to multiple IP addresses.

DNS Truncation – But Wait, There's More!

According to the DNS [RFC 5966](#):

“In the absence of EDNS0 (Extension Mechanisms for DNS 0), the normal behavior of any DNS server needing to send a UDP response that would exceed the 512-byte limit is for the server to truncate the response so that it fits within that limit and then set the TC flag in the response header. When the client receives such a response, it takes the TC flag as an indication that it should retry over TCP instead.”

Great! So we can set the **TC** (truncation) flag in our response, which causes the target Windows DNS Server to initiate a new TCP connection to our malicious NameServer, and we can pass a message larger than 4,096 bytes. But how much larger?

According to DNS [RFC 7766](#):

“DNS clients and servers SHOULD pass the two-octet length field, and the message described by that length field, to the TCP layer at the same time (e.g., in a single “write” system call) to make it more likely that all the data will be transmitted in a single TCP segment.”

As the first two bytes of the message represent its length, the maximum size of a message in DNS over TCP is represented as 16 bits and is therefore limited to 64KB.

Figure 8: The first two bytes of a DNS over TCP message represent the message's length.

But even a message of length 65,535 is not large enough to trigger the vulnerability, as the message length includes the headers and the original query. This overhead is not taken into consideration when calculating the size that is passed to `RR_AllocateEx`.

DNS Pointer Compression – Less is More

Let's have another look at a legitimate DNS response (we chose a response of type A for convenience).

Figure 9: DNS response for `dig research.checkpoint.com A @8.8.8.8`, as seen in Wireshark.

You can see that Wireshark evaluated the bytes `0xc00c` in the answer's name field to `research.checkpoint.com`. The question is, why?

According to [A warm welcome to DNS, powerdns.org](#):

'To squeeze as much information as possible into the 512 bytes, DNS names can (and often MUST) be compressed... In this case, the DNS name of the answer is encoded as 0xc0 0x0c. The c0 part has the two most significant bits set, indicating that the following 6+8 bits are a pointer to somewhere earlier in the message. In this case, this points to position 12 (= 0x0c) within the packet, which is immediately after the DNS header."

What is at the offset 0x0c (12) from the beginning of the packet? It's `research.checkpoint.com`!

In this form of compression, the pointer points at the start of an encoded string. In DNS, strings are encoded as a (`<size><value>`) chain.

Figure 10: An illustration of a <size><value> chain.

So we can use the “magic” byte `0xc0` to reference strings from within the packet. Let’s once again examine the formula that calculates the size that is passed to `RR_AllocateEx`:

`[Name_PacketNameToCountNameEx result] + [0x14] + [The Signature field's length (rdi - rax)]`

Reversing `Name_PacketNameToCountNameEx` confirms the behavior we described above. The purpose of `Name_PacketNameToCountNameEx` is to calculate the size of a name field, taking pointer compression into consideration. Having a primitive that allows us to increase the size of the allocation by a large amount, when only representing it with two bytes, is exactly what we need.

Therefore, we can use the pointer compression in the SIG Signer’s Name field. However, simply specifying `0xc00c` as the Signer’s name would not cause the overflow, as the queried domain name is already present in the query, and the overhead size is subtracted from the allocated value. But what about `0xc00d`? The only constraint we have to satisfy is that our encoded string is valid (ending with `0x0000`), and we can do it

easily because we have a field without any character constraints – the signature value. For the domain `41414141.fun`, `0xc00d` points at the first character of the domain ('4'). The ordinal value of this character is then used as the size of the uncompressed string ('4' represents the value 0x34 (52)). Aggregation of the size of this uncompressed string, with the maximum amount of data we can fit in the Signature field (up to 65,535, depending on the original query), results in a value greater than 65,535 bytes, thus causing the overflow!

Let's test this with WinDBG attached to `dns.exe`:

We crashed!

Although it seems that we crashed because we were trying to write values to unmapped memory, the heap can be shaped in a way that allows us to overwrite some meaningful values.

It is also worth mentioning that since SIG records and RRSIG records have the same structure, Microsoft uses the same function (`SigWireRead`) to parse both record types. This can be seen when examining `RRWireReadTable`: both indexes 0x18 (24) and 46 (0x2e) are pointing to the function `SigWireRead`. This means that both record types SIG and RRSIG can be used to trigger this vulnerability, as they are parsed by the same vulnerable function – `SigWireRead`.

Previous exploitation attempts for `dns.exe` are available online. For example: [A deeper look at ms11-058.](#)

Triggering From the Browser

We know this bug can be triggered by a malicious actor who is present in the LAN environment. However, we thought it would be interesting to see if this bug can be triggered remotely without LAN access.

Smuggling DNS inside HTTP

By now you should be aware that DNS can be transported over TCP and that Windows DNS Server supports this connection type. You should also be familiar with the structure of DNS over TCP, but just in case, here's a quick review:

Figure 11: DNS over TCP message format.

Consider the following standard HTTP payload:

0000	50 4f 53 54 20 2f 70 77 6e 20 48 54 54 50 2f 31	POST /pwn HTTP/1
0010	2e 31 0d 0a 41 63 63 65 70 74 3a 20 2a 2f 2a 0d	.1..Accept: */*.
0020	0a 52 65 66 65 72 65 72 3a 20 68 74 74 70 3a 2f	.Referer: http:/

Even though this is an HTTP payload, sending it to our target DNS server on port 53 causes the Windows DNS Server to interpret this payload as if it was a DNS query. It does this using the following structure:

0000	50 4f 53 54 20 2f 70 77 6e 20 48 54 54 50 2f 31	POST /pwn HTTP/1
0010	2e 31 0d 0a 41 63 63 65 70 74 3a 20 2a 2f 2a 0d	.1..Accept: */*.
0020	0a 52 65 66 65 72 65 72 3a 20 68 74 74 70 3a 2f	.Referer: http:/

Message Length: 20559 (0x504f)

Transaction ID: 0x5354

Flags: 0x202f

Questions: 28791 (0x7077)

Answer RRs: 28192 (0x6e20)

Authority RRs: 18516 (0x4854)

Additional RRs: 21584 (0x5450)

Queries: [...]

Fortunately, Windows DNS Server supports both “Connection Reuse” and “Pipelining” of [RFC 7766](#), which means we can issue multiple queries over a single TCP session and we can do so without waiting for replies.

Why is this important?

We can use basic JavaScript to issue a POST request to the DNS Server from the browser when a victim visits a website we control. But as shown above, the POST request is interpreted in a manner we don’t really control.

However, we can abuse the “Connection Reuse” and “Pipelining” features by sending an HTTP POST request to the target DNS server (<http://target-dns:53/>) with binary data, containing another “smuggled” DNS query in the POST data, to be queried separately.

Our HTTP payload consists of the following:

- HTTP request headers that we do not control (User-Agent, Referer, etc).
- “Padding” so that the first DNS query has a proper length (0x504f) inside the POST data.
- Our “smuggled” DNS query inside the POST data.

Figure 12: Multiple queries over a single TCP session as seen in Wireshark.

<https://youtu.be/PULMmhD5it8>

In practice, most popular browsers (such as Google Chrome and Mozilla Firefox) do not allow HTTP requests to port 53, so this bug can only be exploited in a limited set of web browsers – including Internet Explorer and Microsoft Edge (non-Chromium based).

Variant Analysis

The primary reason that this bug exists is because the `RR_AllocateEx` API expects a size parameter of 16 bits. It is generally safe to assume that the size of a single DNS message does not exceed 64KB and thus this behavior should not present an issue. However, as we just saw, this assumption is wrong when the result of `Name_PacketNameToCountNameEx` is taken into consideration while calculating the size of the buffer. This

happens because the `Name_PacketNameToCountNameEx` function calculates the effective size of the uncompressed name and not the number of bytes it took to represent it in the packet.

To find other variants of this bug, we need to find a function that satisfies the following conditions:

- `RR_AllocateEx` is called with a variable size (and not a constant value).
- There is a call to `Name_PacketNameToCountNameEx` and its result is used to calculate the size passed to `RR_AllocateEx`.
- The value that is passed to `RR_AllocateEx` is calculated using values in the range of 16bits or more.

The only other function in `dns.exe` that satisfied these three conditions is `NsecWireRead`. Let's examine the following simplified code snippet we deduced from decompiling the function:

```
1. RESOURCE_RECORD* NsecWireRead(PARSED_WIRE_RECORD *pParsedWireRecord,
2.                               DNS_PACKET *pPacket, BYTE *pRecordData, WORD wRecordDataLength)
3. {
4.     DNS_RESOURCE_RECORD *pResourceRecord;
5.     unsigned BYTE *pCurrentPos;
6.     unsigned int dwRemainingDataLength;
7.     unsigned int dwBytesRead;
8.     unsigned int dwAllocationSize;
9.     DNS_COUNT_NAME countName;
10.    pResourceRecord = NULL;
11.    pCurrentPos = Name_PacketNameToCountNameEx(&countName, pPacket,
12.                                              pRecordData, pRecordData + wRecordDataLength, 0);
13.    if (pCurrentPos)
14.    {
15.        if
16.            (pCurrentPos >= pRecordData //)
17.            <- Check #1 - Bounds check
18.            && pCurrentPos - pRecordData <= 0xFFFFFFFF //)
19.            <- Check #2 - Same bounds check (?)
20.            && wRecordDataLength >= (unsigned int)(pCurrentPos - pRecordData)) //)
21.            <- Check #3 - Bounds check
22.        {
23.            dwRemainingDataLength = wRecordDataLength - (pCurrentPos -
24.                pRecordData);
25.            dwBytesRead = countName.bNameLength + 2;
```

```
20.         // size := len(countName) + 2 + len(payload)
21.         dwAllocationSize = dwBytesRead + dwRemainingDataLength;
22.         if (dwBytesRead + dwRemainingDataLength >= dwBytesRead) // 
23.             <-- Check #4 - Integer Overflow check (32 bits)
24.             && dwAllocationSize <= 0xFFFF) // 
25.             <-- Check #5 - Integer Overflow check (16 bits)
26.         {
27.             pResourceRecord = RR_AllocateEx(dwAllocationSize, 0, 0);
28.             if (pResourceRecord)
29.             {
30.                 Name_CopyCountName(&pResourceRecord->data, &countName);
31.                 memcpy(&pResourceRecord->data + pResourceRecord->data->bOffset + 2,
32.                         pCurrentPos, dwRemainingDataLength);
33.             }
34.             return pResourceRecord;
35.         }
```

As you can see, this function contains many security checks. One of them (Check #5) is a 16 bit overflow check that prevents the variant of our vulnerability in this function. We would also like to mention that this function has many more security checks than the average function in `dns.exe`, which makes us wonder if this bug was already noticed and fixed, but only in that specific function.

As we mentioned previously, Microsoft implemented the DNS client and DNS server in two different modules. While our vulnerability definitely exists in the DNS server, we wanted to see if it exists in the DNS client as well.

Figure 13: Disassembly snippet of `Sig_RecordRead` from `dnsapi.dll`.

It appears that, unlike `dns.exe!SigWireRead`, `dnsapi.dll!Sig_RecordRead` *does* validate at `Sig_RecordRead+D0` that the value that is passed to `dnsapi.dll!Dns_AllocateRecordEx` is less than 0xFFFF bytes, thus preventing the overflow.

The fact that this vulnerability does not exist in `dnsapi.dll`, as well as having different naming conventions between the two modules, leads us to believe that Microsoft manages two completely different code bases for the DNS server and the DNS client, and does not synchronize bug patches between them.

Exploitation Plan

Per Microsoft's request, we decided to withhold information about the exploitation primitives in order to give users enough time to patch their DNS servers. Instead, we discuss our exploitation plan as it applies to Windows Server 2012R2. However, we do believe that this plan should apply to other versions of Windows Server as well.

The `dns.exe` binary was compiled with Control Flow Guard (CFG), which means the traditional approach of overwriting a function pointer in memory is not enough to exploit this bug. If this binary was not compiled with CFG, exploiting this bug would be pretty straight-forward, as quite early on we encountered the following crash:

Figure 14: Crash at `ntdll!LdrpValidateUserCallTarget`.

As you can see, we crashed at `ntdll!LdrpValidateUserCallTarget`. This is the function responsible for validating function pointer targets as part of CFG. We can see that the pointer to be validated (`rcx`) is fully controllable, which means that we successfully overwrote a function pointer somewhere along the way. The reason we saw a crash is that the function pointer is used as an index to a global bitmap table with “allowed” / “disallowed” bit per address, and our arbitrary address led to a read from an unmapped page in the table itself.

To exploit this bug to a full Remote Code Execution while defeating CFG, we need to find primitives that give us the following capabilities: write-what-where (to precisely overwrite a return address on the stack) and an info leak (to leak memory addresses, such as the stack).

Infoleak

In order to achieve an Infoleak primitive, we corrupted the metadata of a DNS resource record, while it is still in the cache, using our overflow. Then, when queried again from the cache, we were able to leak adjacent heap memory.

WinDNS' Heap Manager

WinDNS uses the function `Mem_Alloc` to dynamically allocate memory. This function manages its own memory pools to be used as an efficient cache. There are 4 memory pool buckets for different allocation sizes (up to 0x50, 0x68, 0x88, 0xA0). If the requested allocation size is greater than 0xA0 bytes, it defaults to `HeapAlloc`, which uses the native Windows heap. The heap manager allocates an additional 0x10 bytes for the memory pool header, which contains metadata including the buffer's type (allocated / free), a pointer to

the next available chunk of memory, a cookie for debug checks, and more. The heap manager implemented its allocation lists in a singly-linked-list fashion, meaning that chunks are allocated in the reverse order that they were freed (LIFO).

Write-What-Where

To achieve a write-what-where primitive, we attacked the WinDNS heap manager by corrupting a chunk's header (metadata), de-facto corrupting the freelist.

After the freelist is corrupted, the next time we try to allocate anything of the right size, the memory allocator assigns a memory region of our choice for us as a writable allocation – a “Malloc-Where” exploit primitive.

To bypass CFG, we want that memory region to be on the stack (whose location we hopefully know thanks to the infoleak). Once we have a write capability on the stack, we can overwrite a return address to an address we want to execute, effectively hijacking the execution flow.

It is important to mention that by default, the DNS service restarts in the first 3 crashes, increasing the chances for successful exploitation.

Conclusion

This high-severity vulnerability was acknowledged by Microsoft and was assigned CVE-2020-1350.

We believe that the likelihood of this vulnerability being exploited is high, as we internally found all of the primitives required to exploit this bug. Due to time constraints, we did not continue to pursue the exploitation of the bug (which includes chaining together all of the exploitation primitives), but we do believe that a determined attacker will be able to exploit it. Successful exploitation of this vulnerability would have a severe impact, as you can often find unpatched Windows Domain environments, especially Domain Controllers. In addition, some Internet Service Providers (ISPs) may even have set up their public DNS servers as WinDNS.

We strongly recommend users to patch their affected Windows DNS Servers in order to prevent the exploitation of this vulnerability.

As a temporary workaround, until the patch is applied, we suggest setting the maximum length of a DNS message (over TCP) to 0xFF00, which should eliminate the vulnerability. You can do so by executing the following commands:

```
1. reg add "HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Services\DNS\Parameters" /v "TcpReceivePacketSize" /t REG_DWORD /d 0xFF00 /f  
2. net stop DNS && net start DNS
```

Check Point [IPS](#) blade provides protection against this threat:

“Microsoft Windows DNS Server Remote Code Execution (CVE-2020-1350)”

Check Point [SandBlast Agent](#) E83.11 already protects against this threat

Disclosure Timeline

- 19 May 2020 – Initial report to Microsoft.
- 18 Jun 2020 – Microsoft issued CVE-2020-1350 to this vulnerability.
- 09 Jul 2020 – Microsoft acknowledged this issue as a wormable, critical vulnerability with a CVSS score of 10.0.
- 14 Jul 2020 – Microsoft released a fix (Patch Tuesday).

References

- https://en.wikipedia.org/wiki/Domain_Name_System
- <https://blog.skullsecurity.org/2011/a-deeper-look-at-ms11-058>
- <https://know.bishopfox.com/blog/2017/10/a-bug-has-no-name-multiple-heap-buffer-overflows-in-the-windows-dns-client>
- <https://powerdns.org/hello-dns/basic.md.html>

- <https://www.cloudflare.com/learning/dns/what-is-dns/>
- <https://tools.ietf.org/html/rfc7766>
- <https://tools.ietf.org/html/rfc5966>
- <https://tools.ietf.org/html/rfc2535>

Many thanks to my colleagues Eyal Itkin ([@EyalItkin](#)) and Omri Herscovici ([@omriher](#)) for their help in this research.



[GO UP](#)

[**BACK TO ALL POSTS**](#)

POPULAR POSTS



ARTIFICIAL INTELLIGENCE

CHATGPT

CHECK POINT RESEARCH PUBLICATIONS

OPWNAI : Cybercriminals Starting to Use ChatGPT

BLOGS AND PUBLICATIONS





Let's get in touch

Subscribe for cpr blogs, news and more

[**Subscribe Now**](#)

© 1994-2024 Check Point Software Technologies LTD. All rights reserved.

Property of CheckPoint.com

[Privacy Policy](#)