man7.org > Linux > man-pages   **Linux/UNIX system programming training**

bpf-helpers(7) — Linux manual page

NAME | DESCRIPTION | HELPERS | EXAMPLES | LICENSE | IMPLEMENTATION | SEE ALSO |
COLOPHON

| Search online pages |

BPF-HELPERS(7)   Miscellaneous Information Manual   *BPF-HELPERS*(7)

NAME   top

BPF-HELPERS - list of eBPF helper functions

DESCRIPTION   top

The extended Berkeley Packet Filter (eBPF) subsystem consists in
programs written in a pseudo-assembly language, then attached to
one of the several kernel hooks and run in reaction of specific
events. This framework differs from the older, "classic" BPF (or
"cBPF") in several aspects, one of them being the ability to call
special functions (or "helpers") from within a program. These
functions are restricted to a white-list of helpers defined in
the kernel.

These helpers are used by eBPF programs to interact with the
system, or with the context in which they work. For instance,
they can be used to print debugging messages, to get the time
since the system was booted, to interact with eBPF maps, or to
manipulate network packets. Since there are several eBPF program
types, and that they do not run in the same context, each program
type can only call a subset of those helpers.

Due to eBPF conventions, a helper can not have more than five
arguments.

Internally, eBPF programs call directly into the compiled helper
functions without requiring any foreign-function interface. As a
result, calling helpers introduces no overhead, thus offering
excellent performance.

This document is an attempt to list and document the helpers

available to eBPF developers. They are sorted by chronological
order (the oldest helpers in the kernel at the top).

HELPERS   top

void *bpf_map_lookup_elem(struct bpf_map **map***, const void ****key***)**

> **Description**
> > Perform a lookup in *map* for an entry associated to
> > *key*.
>
> **Return** Map value associated to *key*, or **NULL** if no entry
> > was found.

long bpf_map_update_elem(struct bpf_map **map***, const void ****key***,
const void ****value***, u64** *flags***)**

> **Description**
> > Add or update the value of the entry associated to
> > *key* in *map* with *value*. *flags* is one of:
> >
> > **BPF_NOEXIST**
> > > The entry for *key* must not exist in the map.
> >
> > **BPF_EXIST**
> > > The entry for *key* must already exist in the
> > > map.
> >
> > **BPF_ANY**
> > > No condition on the existence of the entry
> > > for *key*.
> >
> > Flag value **BPF_NOEXIST** cannot be used for maps of
> > types **BPF_MAP_TYPE_ARRAY** or
> > **BPF_MAP_TYPE_PERCPU_ARRAY** (all elements always
> > exist), the helper would return an error.
>
> **Return** 0 on success, or a negative error in case of
> > failure.

long bpf_map_delete_elem(struct bpf_map **map***, const void ****key***)**

> **Description**
> > Delete entry with *key* from *map*.
>
> **Return** 0 on success, or a negative error in case of
> > failure.

long bpf_probe_read(void *_dst_, u32 _size_, const void *_unsafe_ptr_)

>   **Description**
>   For tracing programs, safely attempt to read _size_
>   bytes from kernel space address _unsafe_ptr_ and
>   store the data in _dst_.
>
>   Generally, use **bpf_probe_read_user**() or
>   **bpf_probe_read_kernel**() instead.
>
>   **Return** 0 on success, or a negative error in case of
>   failure.

u64 bpf_ktime_get_ns(void)

>   **Description**
>   Return the time elapsed since system boot, in
>   nanoseconds. Does not include time the system was
>   suspended. See: **clock_gettime**(**CLOCK_MONOTONIC**)
>
>   **Return** Current _ktime_.

long bpf_trace_printk(const char *_fmt_, u32 _fmt_size_, ...)

>   **Description**
>   This helper is a "printk()-like" facility for
>   debugging. It prints a message defined by format
>   _fmt_ (of size _fmt_size_) to file
>   _/sys/kernel/tracing/trace_ from TraceFS, if
>   available. It can take up to three additional **u64**
>   arguments (as an eBPF helpers, the total number of
>   arguments is limited to five).
>
>   Each time the helper is called, it appends a line
>   to the trace. Lines are discarded while
>   _/sys/kernel/tracing/trace_ is open, use
>   _/sys/kernel/tracing/trace_pipe_ to avoid this. The
>   format of the trace is customizable, and the exact
>   output one will get depends on the options set in
>   _/sys/kernel/tracing/trace_options_ (see also the
>   _README_ file under the same directory). However, it
>   usually defaults to something like:
>
>   telnet-470   [001] .N.. 419421.045894: 0x00000001: <formatted msg>
>
>   In the above:

- **telnet** is the name of the current task.

- **470** is the PID of the current task.

- **001** is the CPU number on which the task is running.

- In **.N..**, each character refers to a set of options (whether irqs are enabled, scheduling options, whether hard/softirqs are running, level of preempt_disabled respectively). **N** means that **TIF_NEED_RESCHED** and **PREEMPT_NEED_RESCHED** are set.

- **419421.045894** is a timestamp.

- **0x00000001** is a fake value used by BPF for the instruction pointer register.

- **<formatted msg>** is the message formatted with *fmt*.

The conversion specifiers supported by *fmt* are similar, but more limited than for printk(). They are **%d, %i, %u, %x, %ld, %li, %lu, %lx, %lld, %lli, %llu, %llx, %p, %s**. No modifier (size of field, padding with zeroes, etc.) is available, and the helper will return **-EINVAL** (but print nothing) if it encounters an unknown specifier.

Also, note that **bpf_trace_printk**() is slow, and should only be used for debugging purposes. For this reason, a notice block (spanning several lines) is printed to kernel logs and states that the helper should not be used "for production use" the first time this helper is used (or more precisely, when **trace_printk**() buffers are allocated). For passing values to user space, perf events should be preferred.

Return The number of bytes written to the buffer, or a negative error in case of failure.

u32 bpf_get_prandom_u32(void)

Description
Get a pseudo-random number.

From a security point of view, this helper uses its own pseudo-random internal state, and cannot be used to infer the seed of other random functions in the kernel. However, it is essential to note that the generator used by the helper is not cryptographically secure.

Return A random 32-bit unsigned value.

u32 bpf_get_smp_processor_id(void)

> **Description**
> > Get the SMP (symmetric multiprocessing) processor id. Note that all programs run with migration disabled, which means that the SMP processor id is stable during all the execution of the program.
>
> **Return** The SMP id of the processor running the program.

long bpf_skb_store_bytes(struct sk_buff *_skb_, u32 _offset_, const void *_from_, u32 _len_, u64 _flags_)

> **Description**
> > Store _len_ bytes from address _from_ into the packet associated to _skb_, at _offset_. _flags_ are a combination of **BPF_F_RECOMPUTE_CSUM** (automatically recompute the checksum for the packet after storing the bytes) and **BPF_F_INVALIDATE_HASH** (set _skb->_**hash**, _skb->_**swhash** and _skb->_**l4hash** to 0).
> >
> > A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_l3_csum_replace(struct sk_buff *_skb_, u32 _offset_, u64 _from_, u64 _to_, u64 _size_)

> **Description**
> > Recompute the layer 3 (e.g. IP) checksum for the packet associated to _skb_. Computation is incremental, so the helper must know the former value of the header field that was modified (_from_),

the new value of this field (*to*), and the number of
bytes (2 or 4) for this field, stored in *size*.
Alternatively, it is possible to store the
difference between the previous and the new values
of the header field in *to*, by setting *from* and *size*
to 0. For both methods, *offset* indicates the
location of the IP checksum within the packet.

This helper works in combination with
bpf_csum_diff(), which does not update the checksum
in-place, but offers more flexibility and can
handle sizes larger than 2 or 4 for the checksum to
update.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

long bpf_l4_csum_replace(struct sk_buff **skb***, u32** *offset***, u64**
*from***, u64** *to***, u64** *flags***)**

Description
Recompute the layer 4 (e.g. TCP, UDP or ICMP)
checksum for the packet associated to *skb*.
Computation is incremental, so the helper must know
the former value of the header field that was
modified (*from*), the new value of this field (*to*),
and the number of bytes (2 or 4) for this field,
stored on the lowest four bits of *flags*.
Alternatively, it is possible to store the
difference between the previous and the new values
of the header field in *to*, by setting *from* and the
four lowest bits of *flags* to 0. For both methods,
offset indicates the location of the IP checksum
within the packet. In addition to the size of the
field, *flags* can be added (bitwise OR) actual
flags. With **BPF_F_MARK_MANGLED_0**, a null checksum
is left untouched (unless **BPF_F_MARK_ENFORCE** is
added as well), and for updates resulting in a null
checksum the value is set to **CSUM_MANGLED_0**
instead. Flag **BPF_F_PSEUDO_HDR** indicates the
checksum is to be computed against a pseudo-header.

This helper works in combination with
bpf_csum_diff(), which does not update the checksum
in-place, but offers more flexibility and can
handle sizes larger than 2 or 4 for the checksum to
update.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

**long bpf_tail_call(void *_ctx_, struct bpf_map *_prog_array_map_, u32
index)**

Description
This special helper is used to trigger a "tail
call", or in other words, to jump into another eBPF
program. The same stack frame is used (but values
on stack and in registers for the caller are not
accessible to the callee). This mechanism allows
for program chaining, either for raising the
maximum number of available eBPF instructions, or
to execute given programs in conditional blocks.
For security reasons, there is an upper limit to
the number of successive tail calls that can be
performed.

Upon call of this helper, the program attempts to
jump into a program referenced at index _index_ in
_prog_array_map_, a special map of type
BPF_MAP_TYPE_PROG_ARRAY, and passes _ctx_, a pointer
to the context.

If the call succeeds, the kernel immediately runs
the first instruction of the new program. This is
not a function call, and it never returns to the
previous program. If the call fails, then the
helper has no effect, and the caller continues to
run its subsequent instructions. A call can fail if
the destination program for the jump does not exist
(i.e. _index_ is superior to the number of entries in
_prog_array_map_), or if the maximum number of tail

calls has been reached for this chain of programs. This limit is defined in the kernel by the macro **MAX_TAIL_CALL_CNT** (not accessible to user space), which is currently set to 33.

Return 0 on success, or a negative error in case of failure.

long bpf_clone_redirect(struct sk_buff *_skb_, u32 _ifindex_, u64 _flags_)

> **Description**
>> Clone and redirect the packet associated to _skb_ to another net device of index _ifindex_. Both ingress and egress interfaces can be used for redirection. The **BPF_F_INGRESS** value in _flags_ is used to make the distinction (ingress path is selected if the flag is present, egress path otherwise). This is the only flag supported for now.
>>
>> In comparison with **bpf_redirect**() helper, **bpf_clone_redirect**() has the associated cost of duplicating the packet buffer, but this can be executed out of the eBPF program. Conversely, **bpf_redirect**() is more efficient, but it is handled through an action code where the redirection happens only after the eBPF program has returned.
>>
>> A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

> **Return** 0 on success, or a negative error in case of failure. Positive error indicates a potential drop or congestion in the target device. The particular positive error codes are not defined.

u64 bpf_get_current_pid_tgid(void)

> **Description**
>> Get the current pid and tgid.

> **Return** A 64-bit integer containing the current tgid and pid, and created as such: _current_task_->**tgid** << **32** | _current_task_->**pid**.

u64 bpf_get_current_uid_gid(void)

> **Description**
> > Get the current uid and gid.
>
> **Return** A 64-bit integer containing the current GID and
> > UID, and created as such: *current_gid* **<< 32 |**
> > *current_uid*.

long bpf_get_current_comm(void *_buf_, u32 _size_of_buf_)

> **Description**
> > Copy the **comm** attribute of the current task into
> > *buf* of *size_of_buf*. The **comm** attribute contains the
> > name of the executable (excluding the path) for the
> > current task. The *size_of_buf* must be strictly
> > positive. On success, the helper makes sure that
> > the *buf* is NUL-terminated. On failure, it is filled
> > with zeroes.
>
> **Return** 0 on success, or a negative error in case of
> > failure.

u32 bpf_get_cgroup_classid(struct sk_buff *_skb_)

> **Description**
> > Retrieve the classid for the current task, i.e. for
> > the net_cls cgroup to which *skb* belongs.
> >
> > This helper can be used on TC egress path, but not
> > on ingress.
> >
> > The net_cls cgroup provides an interface to tag
> > network packets based on a user-provided identifier
> > for all traffic coming from the tasks belonging to
> > the related cgroup. See also the related kernel
> > documentation, available from the Linux sources in
> > file
> > *Documentation/admin-guide/cgroup-v1/net_cls.rst*.
> >
> > The Linux kernel has two versions for cgroups:
> > there are cgroups v1 and cgroups v2. Both are
> > available to users, who can use a mixture of them,
> > but note that the net_cls cgroup is for cgroup v1
> > only. This makes it incompatible with BPF programs
> > run on cgroups, which is a cgroup-v2-only feature
> > (a socket can only hold data for one version of

cgroups at a time).

This helper is only available is the kernel was compiled with the **CONFIG_CGROUP_NET_CLASSID** configuration option set to "**y**" or to "**m**".

Return The classid, or 0 for the default unconfigured classid.

long bpf_skb_vlan_push(struct sk_buff **skb***, __be16 *vlan_proto*, u16 *vlan_tci*)**

Description
Push a *vlan_tci* (VLAN tag control information) of protocol *vlan_proto* to the packet associated to *skb*, then update the checksum. Note that if *vlan_proto* is different from **ETH_P_8021Q** and **ETH_P_8021AD**, it is considered to be **ETH_P_8021Q**.

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

Return 0 on success, or a negative error in case of failure.

long bpf_skb_vlan_pop(struct sk_buff **skb***)**

Description
Pop a VLAN header from the packet associated to *skb*.

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

Return 0 on success, or a negative error in case of failure.

long bpf_skb_get_tunnel_key(struct sk_buff **skb***, struct bpf_tunnel_key ****key***, u32 *size*, u64 *flags*)**

Description

Get tunnel metadata. This helper takes a pointer
key to an empty **struct bpf_tunnel_key** of **size**, that
will be filled with tunnel metadata for the packet
associated to *skb*. The *flags* can be set to
BPF_F_TUNINFO_IPV6, which indicates that the tunnel
is based on IPv6 protocol instead of IPv4.

The **struct bpf_tunnel_key** is an object that
generalizes the principal parameters used by
various tunneling protocols into a single struct.
This way, it can be used to easily make a decision
based on the contents of the encapsulation header,
"summarized" in this struct. In particular, it
holds the IP address of the remote end (IPv4 or
IPv6, depending on the case) in *key*->**remote_ipv4** or
key->**remote_ipv6**. Also, this struct exposes the
key->**tunnel_id**, which is generally mapped to a VNI
(Virtual Network Identifier), making it
programmable together with the
bpf_skb_set_tunnel_key() helper.

Let's imagine that the following code is part of a
program attached to the TC ingress interface, on
one end of a GRE tunnel, and is supposed to filter
out all messages coming from remote ends with IPv4
address other than 10.0.0.1:

```
int ret;
struct bpf_tunnel_key key = {};

ret = bpf_skb_get_tunnel_key(skb, &key, sizeof(key), 0);
if (ret < 0)
        return TC_ACT_SHOT;     // drop packet

if (key.remote_ipv4 != 0x0a000001)
        return TC_ACT_SHOT;     // drop packet

return TC_ACT_OK;               // accept packet
```

This interface can also be used with all
encapsulation devices that can operate in "collect
metadata" mode: instead of having one network
device per specific configuration, the "collect
metadata" mode only requires a single device where
the configuration can be extracted from this
helper.

This can be used together with various tunnels such
as VXLan, Geneve, GRE or IP in IP (IPIP).

Return 0 on success, or a negative error in case of
failure.

long bpf_skb_set_tunnel_key(struct sk_buff *_skb_**, struct
bpf_tunnel_key ***_key_**, u32** _size_**, u64** _flags_**)**

Description
Populate tunnel metadata for packet associated to
skb. The tunnel metadata is set to the contents of
key, of _size_. The _flags_ can be set to a combination
of the following values:

BPF_F_TUNINFO_IPV6
Indicate that the tunnel is based on IPv6
protocol instead of IPv4.

BPF_F_ZERO_CSUM_TX
For IPv4 packets, add a flag to tunnel
metadata indicating that checksum
computation should be skipped and checksum
set to zeroes.

BPF_F_DONT_FRAGMENT
Add a flag to tunnel metadata indicating
that the packet should not be fragmented.

BPF_F_SEQ_NUMBER
Add a flag to tunnel metadata indicating
that a sequence number should be added to
tunnel header before sending the packet.
This flag was added for GRE encapsulation,
but might be used with other protocols as
well in the future.

BPF_F_NO_TUNNEL_KEY
Add a flag to tunnel metadata indicating
that no tunnel key should be set in the
resulting tunnel header.

Here is a typical usage on the transmit path:

```
struct bpf_tunnel_key key;
    populate key ...
bpf_skb_set_tunnel_key(skb, &key, sizeof(key), 0);
bpf_clone_redirect(skb, vxlan_dev_ifindex, 0);
```

See also the description of the
bpf_skb_get_tunnel_key() helper for additional
information.

Return 0 on success, or a negative error in case of
failure.

u64 bpf_perf_event_read(struct bpf_map *_map_, u64 _flags_)

Description
Read the value of a perf event counter. This helper
relies on a _map_ of type
BPF_MAP_TYPE_PERF_EVENT_ARRAY. The nature of the
perf event counter is selected when _map_ is updated
with perf event file descriptors. The _map_ is an
array whose size is the number of available CPUs,
and each cell contains a value relative to one CPU.
The value to retrieve is indicated by _flags_, that
contains the index of the CPU to look up, masked
with **BPF_F_INDEX_MASK**. Alternatively, _flags_ can be
set to **BPF_F_CURRENT_CPU** to indicate that the value
for the current CPU should be retrieved.

Note that before Linux 4.13, only hardware perf
event can be retrieved.

Also, be aware that the newer helper
bpf_perf_event_read_value() is recommended over
bpf_perf_event_read() in general. The latter has
some ABI quirks where error and counter value are
used as a return code (which is wrong to do since
ranges may overlap). This issue is fixed with
bpf_perf_event_read_value(), which at the same time
provides more features over the
bpf_perf_event_read() interface. Please refer to
the description of **bpf_perf_event_read_value**() for
details.

Return The value of the perf event counter read from the
map, or a negative error code in case of failure.

long bpf_redirect(u32 _ifindex_, u64 _flags_)

Description
Redirect the packet to another net device of index
ifindex. This helper is somewhat similar to
bpf_clone_redirect(), except that the packet is not

cloned, which provides increased performance.

Except for XDP, both ingress and egress interfaces can be used for redirection. The **BPF_F_INGRESS** value in *flags* is used to make the distinction (ingress path is selected if the flag is present, egress path otherwise). Currently, XDP only supports redirection to the egress interface, and accepts no flag at all.

The same effect can also be attained with the more generic **bpf_redirect_map**(), which uses a BPF map to store the redirect target instead of providing it directly to the helper.

Return For XDP, the helper returns **XDP_REDIRECT** on success or **XDP_ABORTED** on error. For other program types, the values are **TC_ACT_REDIRECT** on success or **TC_ACT_SHOT** on error.

u32 bpf_get_route_realm(struct sk_buff *skb)

Description
Retrieve the realm or the route, that is to say the **tclassid** field of the destination for the *skb*. The identifier retrieved is a user-provided tag, similar to the one used with the net_cls cgroup (see description for **bpf_get_cgroup_classid**() helper), but here this tag is held by a route (a destination entry), not by a task.

Retrieving this identifier works with the clsact TC egress hook (see also tc-bpf(8)), or alternatively on conventional classful egress qdiscs, but not on TC ingress path. In case of clsact TC egress hook, this has the advantage that, internally, the destination entry has not been dropped yet in the transmit path. Therefore, the destination entry does not need to be artificially held via **netif_keep_dst**() for a classful qdisc until the *skb* is freed.

This helper is available only if the kernel was compiled with **CONFIG_IP_ROUTE_CLASSID** configuration option.

Return The realm of the route for the packet associated to *skb*, or 0 if none was found.

long bpf_perf_event_output(void *_ctx_**, struct bpf_map ***_map_**, u64** _flags_**, void ***_data_**, u64** _size_**)**

Description

Write raw _data_ blob into a special BPF perf event
held by _map_ of type **BPF_MAP_TYPE_PERF_EVENT_ARRAY**.
This perf event must have the following attributes:
PERF_SAMPLE_RAW as **sample_type**, **PERF_TYPE_SOFTWARE**
as **type**, and **PERF_COUNT_SW_BPF_OUTPUT** as **config**.

The _flags_ are used to indicate the index in _map_ for
which the value must be put, masked with
BPF_F_INDEX_MASK. Alternatively, _flags_ can be set
to **BPF_F_CURRENT_CPU** to indicate that the index of
the current CPU core should be used.

The value to write, of _size_, is passed through eBPF
stack and pointed by _data_.

The context of the program _ctx_ needs also be passed
to the helper.

On user space, a program willing to read the values
needs to call **perf_event_open**() on the perf event
(either for one or for all CPUs) and to store the
file descriptor into the _map_. This must be done
before the eBPF program can send data into it. An
example is available in file
_samples/bpf/trace_output_user.c_ in the Linux kernel
source tree (the eBPF program counterpart is in
_samples/bpf/trace_output_kern.c_).

bpf_perf_event_output() achieves better performance
than **bpf_trace_printk**() for sharing data with user
space, and is much better suitable for streaming
data from eBPF programs.

Note that this helper is not restricted to tracing
use cases and can be used with programs attached to
TC or XDP as well, where it allows for passing data
to user space listeners. Data can be:

• Only custom structs,

• Only the packet payload, or

• A combination of both.

> Return 0 on success, or a negative error in case of
> failure.

long bpf_skb_load_bytes(const void **skb***, u32** *offset***, void ****to***,
u32** *len***)**

> **Description**
>> This helper was provided as an easy way to load
>> data from a packet. It can be used to load *len*
>> bytes from *offset* from the packet associated to
>> *skb*, into the buffer pointed by *to*.
>>
>> Since Linux 4.7, usage of this helper has mostly
>> been replaced by "direct packet access", enabling
>> packet data to be manipulated with *skb*->**data** and
>> *skb*->**data_end** pointing respectively to the first
>> byte of packet data and to the byte after the last
>> byte of packet data. However, it remains useful if
>> one wishes to read large quantities of data at once
>> from a packet into the eBPF stack.
>
> **Return** 0 on success, or a negative error in case of
> failure.

long bpf_get_stackid(void **ctx***, struct bpf_map ****map***, u64** *flags***)**

> **Description**
>> Walk a user or a kernel stack and return its id. To
>> achieve this, the helper needs *ctx*, which is a
>> pointer to the context on which the tracing program
>> is executed, and a pointer to a *map* of type
>> **BPF_MAP_TYPE_STACK_TRACE**.
>>
>> The last argument, *flags*, holds the number of stack
>> frames to skip (from 0 to 255), masked with
>> **BPF_F_SKIP_FIELD_MASK**. The next bits can be used to
>> set a combination of the following flags:
>>
>> **BPF_F_USER_STACK**
>>> Collect a user space stack instead of a
>>> kernel stack.
>>
>> **BPF_F_FAST_STACK_CMP**
>>> Compare stacks by hash only.
>>
>> **BPF_F_REUSE_STACKID**
>>> If two different stacks hash into the same

stackid, discard the old one.

The stack id retrieved is a 32 bit long integer handle which can be further combined with other data (including other stack ids) and used as a key into maps. This can be useful for generating a variety of graphs (such as flame graphs or off-cpu graphs).

For walking a stack, this helper is an improvement over **bpf_probe_read**(), which can be used with unrolled loops but is not efficient and consumes a lot of eBPF instructions. Instead, **bpf_get_stackid**() can collect up to **PERF_MAX_STACK_DEPTH** both kernel and user frames. Note that this limit can be controlled with the **sysctl** program, and that it should be manually increased in order to profile long user stacks (such as stacks for Java programs). To do so, use:

```
# sysctl kernel.perf_event_max_stack=<new value>
```

Return The positive or null stack id on success, or a negative error in case of failure.

s64 bpf_csum_diff(__be32 **from**, u32 *from_size***, __be32 ****to**, u32 *to_size***, __wsum *seed***)**

Description
Compute a checksum difference, from the raw buffer pointed by *from*, of length *from_size* (that must be a multiple of 4), towards the raw buffer pointed by *to*, of size *to_size* (same remark). An optional *seed* can be added to the value (this can be cascaded, the seed may come from a previous call to the helper).

This is flexible enough to be used in several ways:

- With *from_size* == 0, *to_size* > 0 and *seed* set to checksum, it can be used when pushing new data.

- With *from_size* > 0, *to_size* == 0 and *seed* set to checksum, it can be used when removing data from a packet.

- With *from_size* > 0, *to_size* > 0 and *seed* set to 0, it can be used to compute a diff. Note that

from_size and *to_size* do not need to be equal.

This helper can be used in combination with **bpf_l3_csum_replace**() and **bpf_l4_csum_replace**(), to which one can feed in the difference computed with **bpf_csum_diff**().

Return The checksum result, or a negative error code in case of failure.

long bpf_skb_get_tunnel_opt(struct sk_buff * *skb*, **void *** *opt*, **u32** *size*)

> **Description**
> > Retrieve tunnel options metadata for the packet associated to *skb*, and store the raw tunnel option data to the buffer *opt* of *size*.
> >
> > This helper can be used with encapsulation devices that can operate in "collect metadata" mode (please refer to the related note in the description of **bpf_skb_get_tunnel_key**() for more details). A particular example where this can be used is in combination with the Geneve encapsulation protocol, where it allows for pushing (with **bpf_skb_get_tunnel_opt**() helper) and retrieving arbitrary TLVs (Type-Length-Value headers) from the eBPF program. This allows for full customization of these headers.
>
> Return The size of the option data retrieved.

long bpf_skb_set_tunnel_opt(struct sk_buff * *skb*, **void *** *opt*, **u32** *size*)

> **Description**
> > Set tunnel options metadata for the packet associated to *skb* to the option data contained in the raw buffer *opt* of *size*.
> >
> > See also the description of the **bpf_skb_get_tunnel_opt**() helper for additional information.
>
> Return 0 on success, or a negative error in case of failure.

long bpf_skb_change_proto(struct sk_buff * *skb*, **__be16** *proto*, **u64**

flags)

>> **Description**
>>
>> Change the protocol of the *skb* to *proto*. Currently
>> supported are transition from IPv4 to IPv6, and
>> from IPv6 to IPv4. The helper takes care of the
>> groundwork for the transition, including resizing
>> the socket buffer. The eBPF program is expected to
>> fill the new headers, if any, via **skb_store_bytes**()
>> and to recompute the checksums with
>> **bpf_l3_csum_replace**() and **bpf_l4_csum_replace**().
>> The main case for this helper is to perform NAT64
>> operations out of an eBPF program.
>>
>> Internally, the GSO type is marked as dodgy so that
>> headers are checked and segments are recalculated
>> by the GSO/GRO engine. The size for GSO target is
>> adapted as well.
>>
>> All values for *flags* are reserved for future usage,
>> and must be left at zero.
>>
>> A call to this helper is susceptible to change the
>> underlying packet buffer. Therefore, at load time,
>> all checks on pointers previously done by the
>> verifier are invalidated and must be performed
>> again, if the helper is used in combination with
>> direct packet access.

> **Return** 0 on success, or a negative error in case of
> failure.

long bpf_skb_change_type(struct sk_buff **skb***, u32 *type*)**

>> **Description**
>>
>> Change the packet type for the packet associated to
>> *skb*. This comes down to setting *skb-*>**pkt_type** to
>> *type*, except the eBPF program does not have a write
>> access to *skb-*>**pkt_type** beside this helper. Using a
>> helper here allows for graceful handling of errors.
>>
>> The major use case is to change incoming *skb*s to
>> **PACKET_HOST** in a programmatic way instead of
>> having to recirculate via **redirect**(...,
>> **BPF_F_INGRESS**), for example.
>>
>> Note that *type* only allows certain values. At this
>> time, they are:

> **PACKET_HOST**
> > Packet is for us.
>
> **PACKET_BROADCAST**
> > Send packet to all.
>
> **PACKET_MULTICAST**
> > Send packet to group.
>
> **PACKET_OTHERHOST**
> > Send packet to someone else.

> **Return** 0 on success, or a negative error in case of failure.

long bpf_skb_under_cgroup(struct sk_buff *skb**, struct bpf_map ***map**, u32** index**)**

> **Description**
> > Check whether *skb* is a descendant of the cgroup2 held by *map* of type **BPF_MAP_TYPE_CGROUP_ARRAY**, at *index*.

> **Return** The return value depends on the result of the test, and can be:
>
> - 0, if the *skb* failed the cgroup2 descendant test.
>
> - 1, if the *skb* succeeded the cgroup2 descendant test.
>
> - A negative error code, if an error occurred.

u32 bpf_get_hash_recalc(struct sk_buff *skb**)**

> **Description**
> > Retrieve the hash of the packet, *skb*->**hash**. If it is not set, in particular if the hash was cleared due to mangling, recompute this hash. Later accesses to the hash can be done directly with *skb*->**hash**.
> >
> > Calling **bpf_set_hash_invalid**(), changing a packet prototype with **bpf_skb_change_proto**(), or calling **bpf_skb_store_bytes**() with the **BPF_F_INVALIDATE_HASH** are actions susceptible to clear the hash and to trigger a new computation for

the next call to **bpf_get_hash_recalc**().

> **Return** The 32-bit hash.

u64 bpf_get_current_task(void)

> **Description**
>> Get the current task.

> **Return** A pointer to the current task struct.

long bpf_probe_write_user(void *_dst_, const void *_src_, u32 _len_)

> **Description**
>> Attempt in a safe way to write _len_ bytes from the
>> buffer _src_ to _dst_ in memory. It only works for
>> threads that are in user context, and _dst_ must be a
>> valid user space address.
>>
>> This helper should not be used to implement any
>> kind of security mechanism because of TOC-TOU
>> attacks, but rather to debug, divert, and
>> manipulate execution of semi-cooperative processes.
>>
>> Keep in mind that this feature is meant for
>> experiments, and it has a risk of crashing the
>> system and running programs. Therefore, when an
>> eBPF program using this helper is attached, a
>> warning including PID and process name is printed
>> to kernel logs.

> **Return** 0 on success, or a negative error in case of
>> failure.

long bpf_current_task_under_cgroup(struct bpf_map *_map_, u32 _index_)

> **Description**
>> Check whether the probe is being run is the context
>> of a given subset of the cgroup2 hierarchy. The
>> cgroup2 to test is held by _map_ of type
>> **BPF_MAP_TYPE_CGROUP_ARRAY**, at _index_.

> **Return** The return value depends on the result of the test,
>> and can be:
>>
>> - 1, if current task belongs to the cgroup2.

- 0, if current task does not belong to the cgroup2.

- A negative error code, if an error occurred.

long bpf_skb_change_tail(struct sk_buff *skb, **u32** *len*, **u64** *flags***)**

      Description

Resize (trim or grow) the packet associated to *skb* to the new *len*. The *flags* are reserved for future usage, and must be left at zero.

The basic idea is that the helper performs the needed work to change the size of the packet, then the eBPF program rewrites the rest via helpers like **bpf_skb_store_bytes**(), **bpf_l3_csum_replace**(), **bpf_l3_csum_replace**() and others. This helper is a slow path utility intended for replies with control messages. And because it is targeted for slow path, the helper itself can afford to be slow: it implicitly linearizes, unclones and drops offloads from the *skb*.

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

      Return 0 on success, or a negative error in case of failure.

long bpf_skb_pull_data(struct sk_buff *skb, **u32** *len***)**

      Description

Pull in non-linear data in case the *skb* is non-linear and not all of *len* are part of the linear section. Make *len* bytes from *skb* readable and writable. If a zero value is passed for *len*, then all bytes in the linear part of *skb* will be made readable and writable.

This helper is only needed for reading and writing with direct packet access.

For direct packet access, testing that offsets to access are within packet boundaries (test on

skb->**data_end**) is susceptible to fail if offsets
are invalid, or if the requested data is in
non-linear parts of the *skb*. On failure the program
can just bail out, or in the case of a non-linear
buffer, use a helper to make the data available.
The **bpf_skb_load_bytes**() helper is a first solution
to access the data. Another one consists in using
bpf_skb_pull_data to pull in once the non-linear
parts, then retesting and eventually access the
data.

At the same time, this also makes sure the *skb* is
uncloned, which is a necessary condition for direct
write. As this needs to be an invariant for the
write part only, the verifier detects writes and
adds a prologue that is calling **bpf_skb_pull_data()**
to effectively unclone the *skb* from the very
beginning in case it is indeed cloned.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

> **Return** 0 on success, or a negative error in case of
> failure.

s64 bpf_csum_update(struct sk_buff **skb***, __wsum** *csum***)**

> **Description**
> Add the checksum *csum* into *skb*->**csum** in case the
> driver has supplied a checksum for the entire
> packet into that field. Return an error otherwise.
> This helper is intended to be used in combination
> with **bpf_csum_diff**(), in particular when the
> checksum needs to be updated after data has been
> written into the packet through direct packet
> access.

> **Return** The checksum on success, or a negative error code
> in case of failure.

void bpf_set_hash_invalid(struct sk_buff **skb***)**

> **Description**
> Invalidate the current *skb*->**hash**. It can be used

after mangling on headers through direct packet
access, in order to indicate that the hash is
outdated and to trigger a recalculation the next
time the kernel tries to access this hash or when
the **bpf_get_hash_recalc**() helper is called.

Return void.

long bpf_get_numa_node_id(void)

Description
Return the id of the current NUMA node. The primary
use case for this helper is the selection of
sockets for the local NUMA node, when the program
is attached to sockets using the
SO_ATTACH_REUSEPORT_EBPF option (see also
socket(7)), but the helper is also available to
other eBPF program types, similarly to
bpf_get_smp_processor_id().

Return The id of current NUMA node.

long bpf_skb_change_head(struct sk_buff *_skb_, u32 _len_, u64 _flags_)

Description
Grows headroom of packet associated to _skb_ and
adjusts the offset of the MAC header accordingly,
adding _len_ bytes of space. It automatically extends
and reallocates memory as required.

This helper can be used on a layer 3 _skb_ to push a
MAC header for redirection into a layer 2 device.

All values for _flags_ are reserved for future usage,
and must be left at zero.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

long bpf_xdp_adjust_head(struct xdp_buff *_xdp_md_, int _delta_)

Description

        Adjust (move) *xdp_md*->**data** by *delta* bytes. Note
        that it is possible to use a negative value for
        delta. This helper can be used to prepare the
        packet for pushing or popping headers.

        A call to this helper is susceptible to change the
        underlying packet buffer. Therefore, at load time,
        all checks on pointers previously done by the
        verifier are invalidated and must be performed
        again, if the helper is used in combination with
        direct packet access.

Return 0 on success, or a negative error in case of
        failure.

long bpf_probe_read_str(void **dst***, u32** *size***, const void**
unsafe_ptr*)**

Description

        Copy a NUL terminated string from an unsafe kernel
        address *unsafe_ptr* to *dst*. See
        bpf_probe_read_kernel_str() for more details.

        Generally, use **bpf_probe_read_user_str**() or
        bpf_probe_read_kernel_str() instead.

Return On success, the strictly positive length of the
        string, including the trailing NUL character. On
        error, a negative value.

u64 bpf_get_socket_cookie(struct sk_buff **skb***)**

Description

        If the **struct sk_buff** pointed by *skb* has a known
        socket, retrieve the cookie (generated by the
        kernel) of this socket. If no cookie has been set
        yet, generate a new cookie. Once generated, the
        socket cookie remains stable for the life of the
        socket. This helper can be useful for monitoring
        per socket networking traffic statistics as it
        provides a global socket identifier that can be
        assumed unique.

Return A 8-byte long unique number on success, or 0 if the
        socket field is missing inside *skb*.

u64 bpf_get_socket_cookie(struct bpf_sock_addr **ctx***)**

> **Description**
>> Equivalent to bpf_get_socket_cookie() helper that
>> accepts *skb*, but gets socket from **struct
>> bpf_sock_addr** context.
>
> **Return** A 8-byte long unique number.

u64 bpf_get_socket_cookie(struct bpf_sock_ops *_ctx_)

> **Description**
>> Equivalent to **bpf_get_socket_cookie**() helper that
>> accepts *skb*, but gets socket from **struct
>> bpf_sock_ops** context.
>
> **Return** A 8-byte long unique number.

u64 bpf_get_socket_cookie(struct sock *_sk_)

> **Description**
>> Equivalent to **bpf_get_socket_cookie**() helper that
>> accepts *sk*, but gets socket from a BTF **struct sock**.
>> This helper also works for sleepable programs.
>
> **Return** A 8-byte long unique number or 0 if *sk* is NULL.

u32 bpf_get_socket_uid(struct sk_buff *_skb_)

> **Description**
>> Get the owner UID of the socked associated to *skb*.
>
> **Return** The owner UID of the socket associated to *skb*. If
>> the socket is **NULL**, or if it is not a full socket
>> (i.e. if it is a time-wait or a request socket
>> instead), **overflowuid** value is returned (note that
>> **overflowuid** might also be the actual UID value for
>> the socket).

long bpf_set_hash(struct sk_buff *_skb_, u32 _hash_)

> **Description**
>> Set the full hash for *skb* (set the field *skb*->**hash**)
>> to value *hash*.
>
> **Return** 0

**long bpf_setsockopt(void *_bpf_socket_, int _level_, int _optname_,
void *_optval_, int _optlen_)**

Description

Emulate a call to **setsockopt()** on the socket associated to *bpf_socket*, which must be a full socket. The *level* at which the option resides and the name *optname* of the option must be specified, see setsockopt(2) for more information. The option value of length *optlen* is pointed by *optval*.

bpf_socket should be one of the following:

- **struct bpf_sock_ops** for **BPF_PROG_TYPE_SOCK_OPS**.

- **struct bpf_sock_addr** for **BPF_CGROUP_INET4_CONNECT**, **BPF_CGROUP_INET6_CONNECT** and **BPF_CGROUP_UNIX_CONNECT**.

This helper actually implements a subset of **setsockopt()**. It supports the following *level*s:

- **SOL_SOCKET**, which supports the following *optname*s: **SO_RCVBUF**, **SO_SNDBUF**, **SO_MAX_PACING_RATE**, **SO_PRIORITY**, **SO_RCVLOWAT**, **SO_MARK**, **SO_BINDTODEVICE**, **SO_KEEPALIVE**, **SO_REUSEADDR**, **SO_REUSEPORT**, **SO_BINDTOIFINDEX**, **SO_TXREHASH**.

- **IPPROTO_TCP**, which supports the following *optname*s: **TCP_CONGESTION**, **TCP_BPF_IW**, **TCP_BPF_SNDCWND_CLAMP**, **TCP_SAVE_SYN**, **TCP_KEEPIDLE**, **TCP_KEEPINTVL**, **TCP_KEEPCNT**, **TCP_SYNCNT**, **TCP_USER_TIMEOUT**, **TCP_NOTSENT_LOWAT**, **TCP_NODELAY**, **TCP_MAXSEG**, **TCP_WINDOW_CLAMP**, **TCP_THIN_LINEAR_TIMEOUTS**, **TCP_BPF_DELACK_MAX**, **TCP_BPF_RTO_MIN**.

- **IPPROTO_IP**, which supports *optname* **IP_TOS**.

- **IPPROTO_IPV6**, which supports the following *optname*s: **IPV6_TCLASS**, **IPV6_AUTOFLOWLABEL**.

Return 0 on success, or a negative error in case of failure.

long bpf_skb_adjust_room(struct sk_buff **skb**, s32** *len_diff***, u32** *mode***, u64** *flags***)**

Description

Grow or shrink the room for data in the packet associated to *skb* by *len_diff*, and according to the selected *mode*.

By default, the helper will reset any offloaded checksum indicator of the skb to CHECKSUM_NONE. This can be avoided by the following flag:

- **BPF_F_ADJ_ROOM_NO_CSUM_RESET**: Do not reset offloaded checksum data of the skb to CHECKSUM_NONE.

There are two supported modes at this time:

- **BPF_ADJ_ROOM_MAC**: Adjust room at the mac layer (room space is added or removed between the layer 2 and layer 3 headers).

- **BPF_ADJ_ROOM_NET**: Adjust room at the network layer (room space is added or removed between the layer 3 and layer 4 headers).

The following flags are supported at this time:

- **BPF_F_ADJ_ROOM_FIXED_GSO**: Do not adjust gso_size. Adjusting mss in this way is not allowed for datagrams.

- **BPF_F_ADJ_ROOM_ENCAP_L3_IPV4**, **BPF_F_ADJ_ROOM_ENCAP_L3_IPV6**: Any new space is reserved to hold a tunnel header. Configure skb offsets and other fields accordingly.

- **BPF_F_ADJ_ROOM_ENCAP_L4_GRE**, **BPF_F_ADJ_ROOM_ENCAP_L4_UDP**: Use with ENCAP_L3 flags to further specify the tunnel type.

- **BPF_F_ADJ_ROOM_ENCAP_L2**(*len*): Use with ENCAP_L3/L4 flags to further specify the tunnel type; *len* is the length of the inner MAC header.

- **BPF_F_ADJ_ROOM_ENCAP_L2_ETH**: Use with BPF_F_ADJ_ROOM_ENCAP_L2 flag to further specify the L2 type as Ethernet.

- **BPF_F_ADJ_ROOM_DECAP_L3_IPV4**, **BPF_F_ADJ_ROOM_DECAP_L3_IPV6**: Indicate the new IP

header version after decapsulating the outer IP
header. Used when the inner and outer IP versions
are different.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

long bpf_redirect_map(struct bpf_map *_map_, u64 _key_, u64 _flags_)

> **Description**
>> Redirect the packet to the endpoint referenced by
>> _map_ at index _key_. Depending on its type, this _map_
>> can contain references to net devices (for
>> forwarding packets through other ports), or to CPUs
>> (for redirecting XDP frames to another CPU; but
>> this is only implemented for native XDP (with
>> driver support) as of this writing).
>>
>> The lower two bits of _flags_ are used as the return
>> code if the map lookup fails. This is so that the
>> return value can be one of the XDP program return
>> codes up to **XDP_TX**, as chosen by the caller. The
>> higher bits of _flags_ can be set to BPF_F_BROADCAST
>> or BPF_F_EXCLUDE_INGRESS as defined below.
>>
>> With BPF_F_BROADCAST the packet will be broadcasted
>> to all the interfaces in the map, with
>> BPF_F_EXCLUDE_INGRESS the ingress interface will be
>> excluded when do broadcasting.
>>
>> See also **bpf_redirect**(), which only supports
>> redirecting to an ifindex, but doesn't require a
>> map to do so.

> **Return** **XDP_REDIRECT** on success, or the value of the two
> lower bits of the _flags_ argument on error.

**long bpf_sk_redirect_map(struct sk_buff *_skb_, struct bpf_map
*_map_, u32 _key_, u64 _flags_)**

> **Description**

Redirect the packet to the socket referenced by *map*
(of type **BPF_MAP_TYPE_SOCKMAP**) at index *key*. Both
ingress and egress interfaces can be used for
redirection. The **BPF_F_INGRESS** value in *flags* is
used to make the distinction (ingress path is
selected if the flag is present, egress path
otherwise). This is the only flag supported for
now.

Return **SK_PASS** on success, or **SK_DROP** on error.

long bpf_sock_map_update(struct bpf_sock_ops * *skops* **, struct**
bpf_map * *map* **, void *** *key* **, u64** *flags* **)**

> **Description**
>> Add an entry to, or update a *map* referencing
>> sockets. The *skops* is used as a new value for the
>> entry associated to *key*. *flags* is one of:
>>
>> **BPF_NOEXIST**
>>> The entry for *key* must not exist in the map.
>>
>> **BPF_EXIST**
>>> The entry for *key* must already exist in the
>>> map.
>>
>> **BPF_ANY**
>>> No condition on the existence of the entry
>>> for *key*.
>>
>> If the *map* has eBPF programs (parser and verdict),
>> those will be inherited by the socket being added.
>> If the socket is already attached to eBPF programs,
>> this results in an error.
>
> **Return** 0 on success, or a negative error in case of
>> failure.

long bpf_xdp_adjust_meta(struct xdp_buff * *xdp_md* **, int** *delta* **)**

> **Description**
>> Adjust the address pointed by *xdp_md* **->data_meta** by
>> *delta* (which can be positive or negative). Note
>> that this operation modifies the address stored in
>> *xdp_md* **->data**, so the latter must be loaded only
>> after the helper has been called.
>>
>> The use of *xdp_md* **->data_meta** is optional and

programs are not required to use it. The rationale
is that when the packet is processed with XDP (e.g.
as DoS filter), it is possible to push further meta
data along with it before passing to the stack, and
to give the guarantee that an ingress eBPF program
attached as a TC classifier on the same device can
pick this up for further post-processing. Since TC
works with socket buffers, it remains possible to
set from XDP the **mark** or **priority** pointers, or
other pointers for the socket buffer. Having this
scratch space generic and programmable allows for
more flexibility as the user is free to store
whatever meta data they need.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

long bpf_perf_event_read_value(struct bpf_map *_map_**, u64** _flags_**,
struct bpf_perf_event_value ***_buf_**, u32** _buf_size_**)**

Description
Read the value of a perf event counter, and store
it into _buf_ of size _buf_size_. This helper relies on
a _map_ of type **BPF_MAP_TYPE_PERF_EVENT_ARRAY**. The
nature of the perf event counter is selected when
map is updated with perf event file descriptors.
The _map_ is an array whose size is the number of
available CPUs, and each cell contains a value
relative to one CPU. The value to retrieve is
indicated by _flags_, that contains the index of the
CPU to look up, masked with **BPF_F_INDEX_MASK**.
Alternatively, _flags_ can be set to
BPF_F_CURRENT_CPU to indicate that the value for
the current CPU should be retrieved.

This helper behaves in a way close to
bpf_perf_event_read() helper, save that instead of
just returning the value observed, it fills the _buf_
structure. This allows for additional data to be
retrieved: in particular, the enabled and running
times (in _buf_**->enabled** and _buf_**->running,**

respectively) are copied. In general,
bpf_perf_event_read_value() is recommended over
bpf_perf_event_read(), which has some ABI issues
and provides fewer functionalities.

These values are interesting, because hardware PMU
(Performance Monitoring Unit) counters are limited
resources. When there are more PMU based perf
events opened than available counters, kernel will
multiplex these events so each event gets certain
percentage (but not all) of the PMU time. In case
that multiplexing happens, the number of samples or
counter value will not reflect the case compared to
when no multiplexing occurs. This makes comparison
between different runs difficult. Typically, the
counter value should be normalized before comparing
to other experiments. The usual normalization is
done as follows.

$$normalized\_counter = counter * t\_enabled / t\_running$$

Where t_enabled is the time enabled for event and
t_running is the time running for event since last
normalization. The enabled and running times are
accumulated since the perf event open. To achieve
scaling factor between two invocations of an eBPF
program, users can use CPU id as the key (which is
typical for perf array usage model) to remember the
previous value and do the calculation inside the
eBPF program.

Return 0 on success, or a negative error in case of
failure.

**long bpf_perf_prog_read_value(struct bpf_perf_event_data *_ctx_,
struct bpf_perf_event_value *_buf_, u32 _buf_size_)**

Description
For an eBPF program attached to a perf event,
retrieve the value of the event counter associated
to _ctx_ and store it in the structure pointed by _buf_
and of size _buf_size_. Enabled and running times are
also stored in the structure (see description of
helper **bpf_perf_event_read_value**() for more
details).

Return 0 on success, or a negative error in case of
failure.

long bpf_getsockopt(void *_bpf_socket_, int _level_, int _optname_, void *_optval_, int _optlen_)

> **Description**
>> Emulate a call to **getsockopt()** on the socket associated to _bpf_socket_, which must be a full socket. The _level_ at which the option resides and the name _optname_ of the option must be specified, see getsockopt(2) for more information. The retrieved value is stored in the structure pointed by _opval_ and of length _optlen_.
>>
>> _bpf_socket_ should be one of the following:
>>
>> - **struct bpf_sock_ops** for **BPF_PROG_TYPE_SOCK_OPS**.
>>
>> - **struct bpf_sock_addr** for **BPF_CGROUP_INET4_CONNECT**, **BPF_CGROUP_INET6_CONNECT** and **BPF_CGROUP_UNIX_CONNECT**.
>>
>> This helper actually implements a subset of **getsockopt()**. It supports the same set of _optname_s that is supported by the **bpf_setsockopt**() helper. The exceptions are **TCP_BPF_*** is **bpf_setsockopt**() only and **TCP_SAVED_SYN** is **bpf_getsockopt**() only.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_override_return(struct pt_regs *_regs_, u64 _rc_)

> **Description**
>> Used for error injection, this helper uses kprobes to override the return value of the probed function, and to set it to _rc_. The first argument is the context _regs_ on which the kprobe works.
>>
>> This helper works by setting the PC (program counter) to an override function which is run in place of the original probed function. This means the probed function is not run at all. The replacement function just returns with the required value.
>>
>> This helper has security implications, and thus is subject to restrictions. It is only available if

the kernel was compiled with the
CONFIG_BPF_KPROBE_OVERRIDE configuration option,
and in this case it only works on functions tagged
with **ALLOW_ERROR_INJECTION** in the kernel code.

Also, the helper is only available for the
architectures having the
CONFIG_FUNCTION_ERROR_INJECTION option. As of this
writing, x86 architecture is the only one to
support this feature.

Return 0

long bpf_sock_ops_cb_flags_set(struct bpf_sock_ops *_bpf_sock_, int _argval_)

Description
Attempt to set the value of the
bpf_sock_ops_cb_flags field for the full TCP socket
associated to _bpf_sock_ops_ to _argval_.

The primary use of this field is to determine if
there should be calls to eBPF programs of type
BPF_PROG_TYPE_SOCK_OPS at various points in the TCP
code. A program of the same type can change its
value, per connection and as necessary, when the
connection is established. This field is directly
accessible for reading, but this helper must be
used for updates in order to return an error if an
eBPF program tries to set a callback that is not
supported in the current kernel.

argval is a flag array which can combine these
flags:

- **BPF_SOCK_OPS_RTO_CB_FLAG** (retransmission time
  out)

- **BPF_SOCK_OPS_RETRANS_CB_FLAG** (retransmission)

- **BPF_SOCK_OPS_STATE_CB_FLAG** (TCP state change)

- **BPF_SOCK_OPS_RTT_CB_FLAG** (every RTT)

Therefore, this function can be used to clear a
callback flag by setting the appropriate bit to
zero. e.g. to disable the RTO callback:

```
bpf_sock_ops_cb_flags_set(bpf_sock,
        bpf_sock->bpf_sock_ops_cb_flags &
        ~BPF_SOCK_OPS_RTO_CB_FLAG)
```

Here are some examples of where one could call such eBPF program:

- When RTO fires.

- When a packet is retransmitted.

- When the connection terminates.

- When a packet is sent.

- When a packet is received.

Return Code **-EINVAL** if the socket is not a full TCP socket; otherwise, a positive number containing the bits that could not be set is returned (which comes down to 0 if all bits were set as required).

long bpf_msg_redirect_map(struct sk_msg_buff *msg, struct bpf_map *map, u32 key, u64 flags)

Description
This helper is used in programs implementing policies at the socket level. If the message *msg* is allowed to pass (i.e. if the verdict eBPF program returns **SK_PASS**), redirect it to the socket referenced by *map* (of type **BPF_MAP_TYPE_SOCKMAP**) at index *key*. Both ingress and egress interfaces can be used for redirection. The **BPF_F_INGRESS** value in *flags* is used to make the distinction (ingress path is selected if the flag is present, egress path otherwise). This is the only flag supported for now.

Return **SK_PASS** on success, or **SK_DROP** on error.

long bpf_msg_apply_bytes(struct sk_msg_buff *msg, u32 bytes)

Description
For socket policies, apply the verdict of the eBPF program to the next *bytes* (number of bytes) of message *msg*.

For example, this helper can be used in the

following cases:

- A single **sendmsg**() or **sendfile**() system call contains multiple logical messages that the eBPF program is supposed to read and for which it should apply a verdict.

- An eBPF program only cares to read the first *bytes* of a *msg*. If the message has a large payload, then setting up and calling the eBPF program repeatedly for all bytes, even though the verdict is already known, would create unnecessary overhead.

When called from within an eBPF program, the helper sets a counter internal to the BPF infrastructure, that is used to apply the last verdict to the next *bytes*. If *bytes* is smaller than the current data being processed from a **sendmsg**() or **sendfile**() system call, the first *bytes* will be sent and the eBPF program will be re-run with the pointer for start of data pointing to byte number *bytes* **+ 1**. If *bytes* is larger than the current data being processed, then the eBPF verdict will be applied to multiple **sendmsg**() or **sendfile**() calls until *bytes* are consumed.

Note that if a socket closes with the internal counter holding a non-zero value, this is not a problem because data is not being buffered for *bytes* and is sent as it is received.

Return 0

long bpf_msg_cork_bytes(struct sk_msg_buff **msg***, u32** *bytes***)**

Description

For socket policies, prevent the execution of the verdict eBPF program for message *msg* until *bytes* (byte number) have been accumulated.

This can be used when one needs a specific number of bytes before a verdict can be assigned, even if the data spans multiple **sendmsg**() or **sendfile**() calls. The extreme case would be a user calling **sendmsg**() repeatedly with 1-byte long message segments. Obviously, this is bad for performance, but it is still valid. If the eBPF program needs

bytes bytes to validate a header, this helper can be used to prevent the eBPF program to be called again until *bytes* have been accumulated.

Return 0

long bpf_msg_pull_data(struct sk_msg_buff **msg***, u32** *start***, u32** *end***, u64** *flags***)**

Description

For socket policies, pull in non-linear data from user space for *msg* and set pointers *msg*->**data** and *msg*->**data_end** to *start* and *end* bytes offsets into *msg*, respectively.

If a program of type **BPF_PROG_TYPE_SK_MSG** is run on a *msg* it can only parse data that the (**data**, **data_end**) pointers have already consumed. For **sendmsg**() hooks this is likely the first scatterlist element. But for calls relying on the **sendpage** handler (e.g. **sendfile**()) this will be the range (**0**, **0**) because the data is shared with user space and by default the objective is to avoid allowing user space to modify data while (or after) eBPF verdict is being decided. This helper can be used to pull in data and to set the start and end pointer to given values. Data will be copied if necessary (i.e. if data was not linear and if start and end pointers do not point to the same chunk).

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

All values for *flags* are reserved for future usage, and must be left at zero.

Return 0 on success, or a negative error in case of failure.

long bpf_bind(struct bpf_sock_addr **ctx***, struct sockaddr ****addr***, int** *addr_len***)**

Description

Bind the socket associated to *ctx* to the address

pointed by *addr*, of length *addr_len*. This allows
for making outgoing connection from the desired IP
address, which can be useful for example when all
processes inside a cgroup should use one single IP
address on a host that has multiple IP configured.

This helper works for IPv4 and IPv6, TCP and UDP
sockets. The domain (*addr*->**sa_family**) must be
AF_INET (or **AF_INET6**). It's advised to pass zero
port (**sin_port** or **sin6_port**) which triggers
IP_BIND_ADDRESS_NO_PORT-like behavior and lets the
kernel efficiently pick up an unused port as long
as 4-tuple is unique. Passing non-zero port might
lead to degraded performance.

Return 0 on success, or a negative error in case of
failure.

long bpf_xdp_adjust_tail(struct xdp_buff *_xdp_md_, int _delta_)

Description
Adjust (move) *xdp_md*->**data_end** by *delta* bytes. It
is possible to both shrink and grow the packet
tail. Shrink done via *delta* being a negative
integer.

A call to this helper is susceptible to change the
underlying packet buffer. Therefore, at load time,
all checks on pointers previously done by the
verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

**long bpf_skb_get_xfrm_state(struct sk_buff *_skb_, u32 _index_,
struct bpf_xfrm_state *_xfrm_state_, u32 _size_, u64 _flags_)**

Description
Retrieve the XFRM state (IP transform framework,
see also ip-xfrm(8)) at *index* in XFRM "security
path" for *skb*.

The retrieved value is stored in the **struct
bpf_xfrm_state** pointed by *xfrm_state* and of length
size.

All values for *flags* are reserved for future usage, and must be left at zero.

This helper is available only if the kernel was compiled with **CONFIG_XFRM** configuration option.

Return 0 on success, or a negative error in case of failure.

long bpf_get_stack(void *ctx, void *buf, u32 size, u64 flags)

Description

Return a user or a kernel stack in bpf program provided buffer. To achieve this, the helper needs *ctx*, which is a pointer to the context on which the tracing program is executed. To store the stacktrace, the bpf program provides *buf* with a nonnegative *size*.

The last argument, *flags*, holds the number of stack frames to skip (from 0 to 255), masked with **BPF_F_SKIP_FIELD_MASK**. The next bits can be used to set the following flags:

BPF_F_USER_STACK
Collect a user space stack instead of a kernel stack.

BPF_F_USER_BUILD_ID
Collect (build_id, file_offset) instead of ips for user stack, only valid if **BPF_F_USER_STACK** is also specified.

file_offset is an offset relative to the beginning of the executable or shared object file backing the vma which the *ip* falls in. It is *not* an offset relative to that object's base address. Accordingly, it must be adjusted by adding (sh_addr - sh_offset), where sh_{addr,offset} correspond to the executable section containing *file_offset* in the object, for comparisons to symbols' st_value to be valid.

bpf_get_stack() can collect up to **PERF_MAX_STACK_DEPTH** both kernel and user frames, subject to sufficient large buffer size. Note that this limit can be controlled with the **sysctl**

program, and that it should be manually increased
in order to profile long user stacks (such as
stacks for Java programs). To do so, use:

```
# sysctl kernel.perf_event_max_stack=<new value>
```

Return The non-negative copied *buf* length equal to or less
than *size* on success, or a negative error in case
of failure.

long bpf_skb_load_bytes_relative(const void **skb***, u32** *offset***,
void ****to***, u32** *len***, u32** *start_header***)**

Description
This helper is similar to **bpf_skb_load_bytes**() in
that it provides an easy way to load *len* bytes from
offset from the packet associated to *skb*, into the
buffer pointed by *to*. The difference to
bpf_skb_load_bytes() is that a fifth argument
start_header exists in order to select a base
offset to start from. *start_header* can be one of:

BPF_HDR_START_MAC
Base offset to load data from is *skb*'s mac
header.

BPF_HDR_START_NET
Base offset to load data from is *skb*'s
network header.

In general, "direct packet access" is the preferred
method to access packet data, however, this helper
is in particular useful in socket filters where
*skb***->data** does not always point to the start of the
mac header and where "direct packet access" is not
available.

Return 0 on success, or a negative error in case of
failure.

long bpf_fib_lookup(void **ctx***, struct bpf_fib_lookup ****params***, int**
*plen***, u32** *flags***)**

Description
Do FIB lookup in kernel tables using parameters in
params. If lookup is successful and result shows
packet is to be forwarded, the neighbor tables are
searched for the nexthop. If successful (ie., FIB

lookup shows forwarding and nexthop is resolved),
the nexthop address is returned in ipv4_dst or
ipv6_dst based on family, smac is set to mac
address of egress device, dmac is set to nexthop
mac address, rt_metric is set to metric from route
(IPv4/IPv6 only), and ifindex is set to the device
index of the nexthop from the FIB lookup.

plen argument is the size of the passed in struct.
flags argument can be a combination of one or more
of the following values:

BPF_FIB_LOOKUP_DIRECT
    Do a direct table lookup vs full lookup
    using FIB rules.

BPF_FIB_LOOKUP_TBID
    Used with BPF_FIB_LOOKUP_DIRECT. Use the
    routing table ID present in *params*->tbid for
    the fib lookup.

BPF_FIB_LOOKUP_OUTPUT
    Perform lookup from an egress perspective
    (default is ingress).

BPF_FIB_LOOKUP_SKIP_NEIGH
    Skip the neighbour table lookup.
    params->dmac and *params*->smac will not be
    set as output. A common use case is to call
    bpf_redirect_neigh() after doing
    bpf_fib_lookup().

BPF_FIB_LOOKUP_SRC
    Derive and set source IP addr in
    params->ipv{4,6}_src for the nexthop. If the
    src addr cannot be derived,
    BPF_FIB_LKUP_RET_NO_SRC_ADDR is returned. In
    this case, *params*->dmac and *params*->smac are
    not set either.

ctx is either **struct xdp_md** for XDP programs or
struct sk_buff tc cls_act programs.

Return

- < 0 if any input argument is invalid

- 0 on success (packet is forwarded, nexthop

neighbor exists)

- > 0 one of **BPF_FIB_LKUP_RET_** codes explaining why
  the packet is not forwarded or needs assist from
  full stack

  If lookup fails with BPF_FIB_LKUP_RET_FRAG_NEEDED,
  then the MTU was exceeded and output
  params->mtu_result contains the MTU.

long bpf_sock_hash_update(struct bpf_sock_ops *_skops_**, struct
bpf_map ***_map_**, void ***_key_**, u64** _flags_**)**

> **Description**
>> Add an entry to, or update a sockhash _map_
>> referencing sockets. The _skops_ is used as a new
>> value for the entry associated to _key_. _flags_ is one
>> of:
>>
>> **BPF_NOEXIST**
>>> The entry for _key_ must not exist in the map.
>>
>> **BPF_EXIST**
>>> The entry for _key_ must already exist in the
>>> map.
>>
>> **BPF_ANY**
>>> No condition on the existence of the entry
>>> for _key_.
>>
>> If the _map_ has eBPF programs (parser and verdict),
>> those will be inherited by the socket being added.
>> If the socket is already attached to eBPF programs,
>> this results in an error.
>
> **Return** 0 on success, or a negative error in case of
>> failure.

long bpf_msg_redirect_hash(struct sk_msg_buff *_msg_**, struct
bpf_map ***_map_**, void ***_key_**, u64** _flags_**)**

> **Description**
>> This helper is used in programs implementing
>> policies at the socket level. If the message _msg_ is
>> allowed to pass (i.e. if the verdict eBPF program
>> returns **SK_PASS**), redirect it to the socket
>> referenced by _map_ (of type **BPF_MAP_TYPE_SOCKHASH**)
>> using hash _key_. Both ingress and egress interfaces

can be used for redirection. The **BPF_F_INGRESS**
value in *flags* is used to make the distinction
(ingress path is selected if the flag is present,
egress path otherwise). This is the only flag
supported for now.

Return **SK_PASS** on success, or **SK_DROP** on error.

**long bpf_sk_redirect_hash(struct sk_buff *_skb_, struct bpf_map
*_map_, void *_key_, u64 _flags_)**

> **Description**
>> This helper is used in programs implementing
>> policies at the skb socket level. If the sk_buff
>> *skb* is allowed to pass (i.e. if the verdict eBPF
>> program returns **SK_PASS**), redirect it to the socket
>> referenced by *map* (of type **BPF_MAP_TYPE_SOCKHASH**)
>> using hash *key*. Both ingress and egress interfaces
>> can be used for redirection. The **BPF_F_INGRESS**
>> value in *flags* is used to make the distinction
>> (ingress path is selected if the flag is present,
>> egress otherwise). This is the only flag supported
>> for now.

> Return **SK_PASS** on success, or **SK_DROP** on error.

**long bpf_lwt_push_encap(struct sk_buff *_skb_, u32 _type_, void *_hdr_,
u32 _len_)**

> **Description**
>> Encapsulate the packet associated to *skb* within a
>> Layer 3 protocol header. This header is provided in
>> the buffer at address *hdr*, with *len* its size in
>> bytes. *type* indicates the protocol of the header
>> and can be one of:

>> **BPF_LWT_ENCAP_SEG6**
>>> IPv6 encapsulation with Segment Routing
>>> Header (**struct ipv6_sr_hdr**). *hdr* only
>>> contains the SRH, the IPv6 header is
>>> computed by the kernel.

>> **BPF_LWT_ENCAP_SEG6_INLINE**
>>> Only works if *skb* contains an IPv6 packet.
>>> Insert a Segment Routing Header (**struct
>>> ipv6_sr_hdr**) inside the IPv6 header.

>> **BPF_LWT_ENCAP_IP**

IP encapsulation (GRE/GUE/IPIP/etc). The outer header must be IPv4 or IPv6, followed by zero or more additional headers, up to **LWT_BPF_MAX_HEADROOM** total bytes in all prepended headers. Please note that if **skb_is_gso**(*skb*) is true, no more than two headers can be prepended, and the inner header, if present, should be either GRE or UDP/GUE.

BPF_LWT_ENCAP_SEG6* types can be called by BPF programs of type **BPF_PROG_TYPE_LWT_IN**; **BPF_LWT_ENCAP_IP** type can be called by bpf programs of types **BPF_PROG_TYPE_LWT_IN** and **BPF_PROG_TYPE_LWT_XMIT**.

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

Return 0 on success, or a negative error in case of failure.

long bpf_lwt_seg6_store_bytes(struct sk_buff **skb***, u32 ***offset***, const void ****from***, u32 ***len***)**

Description
Store *len* bytes from address *from* into the packet associated to *skb*, at *offset*. Only the flags, tag and TLVs inside the outermost IPv6 Segment Routing Header can be modified through this helper.

A call to this helper is susceptible to change the underlying packet buffer. Therefore, at load time, all checks on pointers previously done by the verifier are invalidated and must be performed again, if the helper is used in combination with direct packet access.

Return 0 on success, or a negative error in case of failure.

long bpf_lwt_seg6_adjust_srh(struct sk_buff **skb***, u32 ***offset***, s32 ***delta***)**

Description
> Adjust the size allocated to TLVs in the outermost
> IPv6 Segment Routing Header contained in the packet
> associated to *skb*, at position *offset* by *delta*
> bytes. Only offsets after the segments are
> accepted. *delta* can be as well positive (growing)
> as negative (shrinking).
>
> A call to this helper is susceptible to change the
> underlying packet buffer. Therefore, at load time,
> all checks on pointers previously done by the
> verifier are invalidated and must be performed
> again, if the helper is used in combination with
> direct packet access.

Return 0 on success, or a negative error in case of
> failure.

long bpf_lwt_seg6_action(struct sk_buff *skb**, u32** action**, void
***param**, u32** param_len**)**

Description
> Apply an IPv6 Segment Routing action of type *action*
> to the packet associated to *skb*. Each action takes
> a parameter contained at address *param*, and of
> length *param_len* bytes. *action* can be one of:
>
> **SEG6_LOCAL_ACTION_END_X**
>> End.X action: Endpoint with Layer-3
>> cross-connect. Type of *param*: **struct
>> in6_addr**.
>
> **SEG6_LOCAL_ACTION_END_T**
>> End.T action: Endpoint with specific IPv6
>> table lookup. Type of *param*: **int**.
>
> **SEG6_LOCAL_ACTION_END_B6**
>> End.B6 action: Endpoint bound to an SRv6
>> policy. Type of *param*: **struct ipv6_sr_hdr**.
>
> **SEG6_LOCAL_ACTION_END_B6_ENCAP**
>> End.B6.Encap action: Endpoint bound to an
>> SRv6 encapsulation policy. Type of *param*:
>> **struct ipv6_sr_hdr**.
>
> A call to this helper is susceptible to change the
> underlying packet buffer. Therefore, at load time,
> all checks on pointers previously done by the

verifier are invalidated and must be performed
again, if the helper is used in combination with
direct packet access.

Return 0 on success, or a negative error in case of
failure.

long bpf_rc_repeat(void *_ctx_)

Description
This helper is used in programs implementing IR
decoding, to report a successfully decoded repeat
key message. This delays the generation of a key up
event for previously generated key down event.

Some IR protocols like NEC have a special IR
message for repeating last button, for when a
button is held down.

The _ctx_ should point to the lirc sample as passed
into the program.

This helper is only available is the kernel was
compiled with the **CONFIG_BPF_LIRC_MODE2**
configuration option set to "**y**".

Return 0

**long bpf_rc_keydown(void *_ctx_, u32 _protocol_, u64 _scancode_, u32
toggle)**

Description
This helper is used in programs implementing IR
decoding, to report a successfully decoded key
press with _scancode_, _toggle_ value in the given
protocol. The scancode will be translated to a
keycode using the rc keymap, and reported as an
input key down event. After a period a key up event
is generated. This period can be extended by
calling either **bpf_rc_keydown**() again with the same
values, or calling **bpf_rc_repeat**().

Some protocols include a toggle bit, in case the
button was released and pressed again between
consecutive scancodes.

The _ctx_ should point to the lirc sample as passed
into the program.

The *protocol* is the decoded protocol number (see **enum rc_proto** for some predefined values).

This helper is only available is the kernel was compiled with the **CONFIG_BPF_LIRC_MODE2** configuration option set to "**y**".

Return 0

u64 bpf_skb_cgroup_id(struct sk_buff *_skb_)

Description
Return the cgroup v2 id of the socket associated with the *skb*. This is roughly similar to the **bpf_get_cgroup_classid**() helper for cgroup v1 by providing a tag resp. identifier that can be matched on or used for map lookups e.g. to implement policy. The cgroup v2 id of a given path in the hierarchy is exposed in user space through the f_handle API in order to get to the same 64-bit id.

This helper can be used on TC egress path, but not on ingress, and is available only if the kernel was compiled with the **CONFIG_SOCK_CGROUP_DATA** configuration option.

Return The id is returned or 0 in case the id could not be retrieved.

u64 bpf_get_current_cgroup_id(void)

Description
Get the current cgroup id based on the cgroup within which the current task is running.

Return A 64-bit integer containing the current cgroup id based on the cgroup within which the current task is running.

void *bpf_get_local_storage(void *_map_, u64 _flags_)

Description
Get the pointer to the local storage area. The type and the size of the local storage is defined by the *map* argument. The *flags* meaning is specific for each map type, and has to be 0 for cgroup local

storage.

Depending on the BPF program type, a local storage
area can be shared between multiple instances of
the BPF program, running simultaneously.

A user should care about the synchronization by
himself. For example, by using the **BPF_ATOMIC**
instructions to alter the shared data.

Return A pointer to the local storage area.

**long bpf_sk_select_reuseport(struct sk_reuseport_md *reuse,
struct bpf_map *map, void *key, u64 flags)**

> **Description**
>> Select a **SO_REUSEPORT** socket from a
>> **BPF_MAP_TYPE_REUSEPORT_SOCKARRAY** map. It checks
>> the selected socket is matching the incoming
>> request in the socket buffer.
>
> **Return** 0 on success, or a negative error in case of
>> failure.

**u64 bpf_skb_ancestor_cgroup_id(struct sk_buff *skb, int
ancestor_level)**

> **Description**
>> Return id of cgroup v2 that is ancestor of cgroup
>> associated with the *skb* at the *ancestor_level*. The
>> root cgroup is at *ancestor_level* zero and each step
>> down the hierarchy increments the level. If
>> *ancestor_level* == level of cgroup associated with
>> *skb*, then return value will be same as that of
>> **bpf_skb_cgroup_id()**.
>
>> The helper is useful to implement policies based on
>> cgroups that are upper in hierarchy than immediate
>> cgroup associated with *skb*.
>
>> The format of returned id and helper limitations
>> are same as in **bpf_skb_cgroup_id()**.
>
> **Return** The id is returned or 0 in case the id could not be
>> retrieved.

**struct bpf_sock *bpf_sk_lookup_tcp(void *ctx, struct
bpf_sock_tuple *tuple, u32 tuple_size, u64 netns, u64 flags)**

Description

Look for TCP socket matching *tuple*, optionally in a child network namespace *netns*. The return value must be checked, and if non-**NULL**, released via **bpf_sk_release**().

The *ctx* should point to the context of the program, such as the skb or socket (depending on the hook in use). This is used to determine the base network namespace for the lookup.

tuple_size must be one of:

sizeof(*tuple*->ipv4)
Look for an IPv4 socket.

sizeof(*tuple*->ipv6)
Look for an IPv6 socket.

If the *netns* is a negative signed 32-bit integer, then the socket lookup table in the netns associated with the *ctx* will be used. For the TC hooks, this is the netns of the device in the skb. For socket hooks, this is the netns of the socket. If *netns* is any other signed 32-bit value greater than or equal to zero then it specifies the ID of the netns relative to the netns associated with the *ctx*. *netns* values beyond the range of 32-bit integers are reserved for future use.

All values for *flags* are reserved for future usage, and must be left at zero.

This helper is available only if the kernel was compiled with **CONFIG_NET** configuration option.

Return Pointer to **struct bpf_sock**, or **NULL** in case of failure. For sockets with reuseport option, the **struct bpf_sock** result is from *reuse*->**socks**[] using the hash of the tuple.

struct bpf_sock *bpf_sk_lookup_udp(void **ctx***, struct bpf_sock_tuple ****tuple***, u32 ***tuple_size***, u64 ***netns***, u64 ***flags***)**

Description

Look for UDP socket matching *tuple*, optionally in a child network namespace *netns*. The return value

must be checked, and if non-**NULL**, released via
bpf_sk_release().

The *ctx* should point to the context of the program,
such as the skb or socket (depending on the hook in
use). This is used to determine the base network
namespace for the lookup.

tuple_size must be one of:

sizeof(*tuple***->ipv4)**
    Look for an IPv4 socket.

sizeof(*tuple***->ipv6)**
    Look for an IPv6 socket.

If the *netns* is a negative signed 32-bit integer,
then the socket lookup table in the netns
associated with the *ctx* will be used. For the TC
hooks, this is the netns of the device in the skb.
For socket hooks, this is the netns of the socket.
If *netns* is any other signed 32-bit value greater
than or equal to zero then it specifies the ID of
the netns relative to the netns associated with the
ctx. *netns* values beyond the range of 32-bit
integers are reserved for future use.

All values for *flags* are reserved for future usage,
and must be left at zero.

This helper is available only if the kernel was
compiled with **CONFIG_NET** configuration option.

Return Pointer to **struct bpf_sock**, or **NULL** in case of
failure. For sockets with reuseport option, the
struct bpf_sock result is from *reuse*->**socks**[] using
the hash of the tuple.

long bpf_sk_release(void * *sock* **)**

Description
    Release the reference held by *sock*. *sock* must be a
    non-**NULL** pointer that was returned from
    bpf_sk_lookup_xxx().

Return 0 on success, or a negative error in case of
failure.

long bpf_map_push_elem(struct bpf_map **map***, const void ****value***, u64** *flags***)**

> **Description**
>> Push an element *value* in *map*. *flags* is one of:
>>
>> **BPF_EXIST**
>>> If the queue/stack is full, the oldest element is removed to make room for this.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_map_pop_elem(struct bpf_map **map***, void ****value***)**

> **Description**
>> Pop an element from *map*.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_map_peek_elem(struct bpf_map **map***, void ****value***)**

> **Description**
>> Get an element from *map* without removing it.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_msg_push_data(struct sk_msg_buff **msg***, u32** *start***, u32** *len***, u64** *flags***)**

> **Description**
>> For socket policies, insert *len* bytes into *msg* at offset *start*.
>>
>> If a program of type **BPF_PROG_TYPE_SK_MSG** is run on a *msg* it may want to insert metadata or options into the *msg*. This can later be read and used by any of the lower layer BPF hooks.
>>
>> This helper may fail if under memory pressure (a malloc fails) in these cases BPF programs will get an appropriate error and BPF programs will need to handle them.
>
> **Return** 0 on success, or a negative error in case of failure.

long bpf_msg_pop_data(struct sk_msg_buff *msg*, **u32** *start*, **u32** *len*, **u64** *flags*)

> **Description**
> > Will remove *len* bytes from a *msg* starting at byte
> > *start*. This may result in **ENOMEM** errors under
> > certain situations if an allocation and copy are
> > required due to a full ring buffer. However, the
> > helper will try to avoid doing the allocation if
> > possible. Other errors can occur if input
> > parameters are invalid either due to *start* byte not
> > being valid part of *msg* payload and/or *pop* value
> > being to large.

> **Return** 0 on success, or a negative error in case of
> > failure.

long bpf_rc_pointer_rel(void *ctx*, **s32** *rel_x*, **s32** *rel_y*)

> **Description**
> > This helper is used in programs implementing IR
> > decoding, to report a successfully decoded pointer
> > movement.

> > The *ctx* should point to the lirc sample as passed
> > into the program.

> > This helper is only available is the kernel was
> > compiled with the **CONFIG_BPF_LIRC_MODE2**
> > configuration option set to "**y**".

> **Return** 0

long bpf_spin_lock(struct bpf_spin_lock *lock*)

> **Description**
> > Acquire a spinlock represented by the pointer *lock*,
> > which is stored as part of a value of a map. Taking
> > the lock allows to safely update the rest of the
> > fields in that value. The spinlock can (and must)
> > later be released with a call to
> > **bpf_spin_unlock**(*lock*).

> > Spinlocks in BPF programs come with a number of
> > restrictions and constraints:

> > • **bpf_spin_lock** objects are only allowed inside

maps of types **BPF_MAP_TYPE_HASH** and
BPF_MAP_TYPE_ARRAY (this list could be extended
in the future).

- BTF description of the map is mandatory.

- The BPF program can take ONE lock at a time,
  since taking two or more could cause dead locks.

- Only one **struct bpf_spin_lock** is allowed per map
  element.

- When the lock is taken, calls (either BPF to BPF
  or helpers) are not allowed.

- The **BPF_LD_ABS** and **BPF_LD_IND** instructions are
  not allowed inside a spinlock-ed region.

- The BPF program MUST call **bpf_spin_unlock**() to
  release the lock, on all execution paths, before
  it returns.

- The BPF program can access **struct bpf_spin_lock**
  only via the **bpf_spin_lock**() and
  bpf_spin_unlock() helpers. Loading or storing
  data into the **struct bpf_spin_lock** *lock*; field of
  a map is not allowed.

- To use the **bpf_spin_lock**() helper, the BTF
  description of the map value must be a struct and
  have **struct bpf_spin_lock** *anyname*; field at the
  top level. Nested lock inside another struct is
  not allowed.

- The **struct bpf_spin_lock** *lock* field in a map
  value must be aligned on a multiple of 4 bytes in
  that value.

- Syscall with command **BPF_MAP_LOOKUP_ELEM** does not
  copy the **bpf_spin_lock** field to user space.

- Syscall with command **BPF_MAP_UPDATE_ELEM**, or
  update from a BPF program, do not update the
  bpf_spin_lock field.

- **bpf_spin_lock** cannot be on the stack or inside a
  networking packet (it can only be inside of a map
  values).

- **bpf_spin_lock** is available to root only.

- Tracing programs and socket filter programs cannot use **bpf_spin_lock**() due to insufficient preemption checks (but this may change in the future).

- **bpf_spin_lock** is not allowed in inner maps of map-in-map.

    Return 0

long bpf_spin_unlock(struct bpf_spin_lock *lock**)**

    Description
        Release the *lock* previously locked by a call to **bpf_spin_lock**(*lock*).

    Return 0

struct bpf_sock *bpf_sk_fullsock(struct bpf_sock *sk**)**

    Description
        This helper gets a **struct bpf_sock** pointer such that all the fields in this **bpf_sock** can be accessed.

    Return A **struct bpf_sock** pointer on success, or **NULL** in case of failure.

struct bpf_tcp_sock *bpf_tcp_sock(struct bpf_sock *sk**)**

    Description
        This helper gets a **struct bpf_tcp_sock** pointer from a **struct bpf_sock** pointer.

    Return A **struct bpf_tcp_sock** pointer on success, or **NULL** in case of failure.

long bpf_skb_ecn_set_ce(struct sk_buff *skb**)**

    Description
        Set ECN (Explicit Congestion Notification) field of IP header to **CE** (Congestion Encountered) if current value is **ECT** (ECN Capable Transport). Otherwise, do nothing. Works with IPv6 and IPv4.

> **Return** 1 if the **CE** flag is set (either by the current
> helper call or because it was already present), 0
> if it is not set.

struct bpf_sock *bpf_get_listener_sock(struct bpf_sock **sk***)**

> **Description**
> Return a **struct bpf_sock** pointer in **TCP_LISTEN**
> state. **bpf_sk_release**() is unnecessary and not
> allowed.
>
> **Return** A **struct bpf_sock** pointer on success, or **NULL** in
> case of failure.

struct bpf_sock *bpf_skc_lookup_tcp(void **ctx***, struct
bpf_sock_tuple ****tuple***, u32** *tuple_size***, u64** *netns***, u64** *flags***)**

> **Description**
> Look for TCP socket matching *tuple*, optionally in a
> child network namespace *netns*. The return value
> must be checked, and if non-**NULL**, released via
> **bpf_sk_release**().
>
> This function is identical to **bpf_sk_lookup_tcp**(),
> except that it also returns timewait or request
> sockets. Use **bpf_sk_fullsock**() or **bpf_tcp_sock**() to
> access the full structure.
>
> This helper is available only if the kernel was
> compiled with **CONFIG_NET** configuration option.
>
> **Return** Pointer to **struct bpf_sock**, or **NULL** in case of
> failure. For sockets with reuseport option, the
> **struct bpf_sock** result is from *reuse*->**socks**[] using
> the hash of the tuple.

long bpf_tcp_check_syncookie(void **sk***, void ****iph***, u32** *iph_len***,
struct tcphdr ****th***, u32** *th_len***)**

> **Description**
> Check whether *iph* and *th* contain a valid SYN cookie
> ACK for the listening socket in *sk*.
>
> *iph* points to the start of the IPv4 or IPv6 header,
> while *iph_len* contains **sizeof**(**struct iphdr**) or
> **sizeof**(**struct ipv6hdr**).
>
> *th* points to the start of the TCP header, while

th_len contains the length of the TCP header (at least **sizeof**(**struct tcphdr**)).

Return 0 if *iph* and *th* are a valid SYN cookie ACK, or a negative error otherwise.

long bpf_sysctl_get_name(struct bpf_sysctl *ctx, char *buf, size_t *buf_len*, u64 *flags*)

Description
Get name of sysctl in /proc/sys/ and copy it into provided by program buffer *buf* of size *buf_len*.

The buffer is always NUL terminated, unless it's zero-sized.

If *flags* is zero, full name (e.g. "net/ipv4/tcp_mem") is copied. Use **BPF_F_SYSCTL_BASE_NAME** flag to copy base name only (e.g. "tcp_mem").

Return Number of character copied (not including the trailing NUL).

-E2BIG if the buffer wasn't big enough (*buf* will contain truncated name in this case).

long bpf_sysctl_get_current_value(struct bpf_sysctl *ctx, char *buf*, size_t *buf_len*)

Description
Get current value of sysctl as it is presented in /proc/sys (incl. newline, etc), and copy it as a string into provided by program buffer *buf* of size *buf_len*.

The whole value is copied, no matter what file position user space issued e.g. sys_read at.

The buffer is always NUL terminated, unless it's zero-sized.

Return Number of character copied (not including the trailing NUL).

-E2BIG if the buffer wasn't big enough (*buf* will contain truncated name in this case).

> **-EINVAL** if current value was unavailable, e.g.
> because sysctl is uninitialized and read returns
> -EIO for it.

**long bpf_sysctl_get_new_value(struct bpf_sysctl *_ctx_, char *_buf_,
size_t _buf_len_)**

> **Description**
> > Get new value being written by user space to sysctl
> > (before the actual write happens) and copy it as a
> > string into provided by program buffer _buf_ of size
> > _buf_len_.
> >
> > User space may write new value at file position >
> > 0.
> >
> > The buffer is always NUL terminated, unless it's
> > zero-sized.

> **Return** Number of character copied (not including the
> > trailing NUL).
> >
> > **-E2BIG** if the buffer wasn't big enough (_buf_ will
> > contain truncated name in this case).
> >
> > **-EINVAL** if sysctl is being read.

**long bpf_sysctl_set_new_value(struct bpf_sysctl *_ctx_, const char
*_buf_, size_t _buf_len_)**

> **Description**
> > Override new value being written by user space to
> > sysctl with value provided by program in buffer _buf_
> > of size _buf_len_.
> >
> > _buf_ should contain a string in same form as
> > provided by user space on sysctl write.
> >
> > User space may write new value at file position >
> > 0. To override the whole sysctl value file position
> > should be set to zero.

> **Return** 0 on success.
> >
> > **-E2BIG** if the _buf_len_ is too big.
> >
> > **-EINVAL** if sysctl is being read.

long bpf_strtol(const char **buf**, size_t** *buf_len***, u64** *flags***, long ****res**)**

       Description

              Convert the initial part of the string from buffer *buf* of size *buf_len* to a long integer according to the given base and save the result in *res*.

              The string may begin with an arbitrary amount of white space (as determined by isspace(3)) followed by a single optional '**-**' sign.

              Five least significant bits of *flags* encode base, other bits are currently unused.

              Base must be either 8, 10, 16 or 0 to detect it automatically similar to user space strtol(3).

       Return Number of characters consumed on success. Must be positive but no more than *buf_len*.

              -EINVAL if no valid digits were found or unsupported base was provided.

              -ERANGE if resulting value was out of range.

long bpf_strtoul(const char **buf**, size_t** *buf_len***, u64** *flags***, unsigned long ****res**)**

       Description

              Convert the initial part of the string from buffer *buf* of size *buf_len* to an unsigned long integer according to the given base and save the result in *res*.

              The string may begin with an arbitrary amount of white space (as determined by isspace(3)).

              Five least significant bits of *flags* encode base, other bits are currently unused.

              Base must be either 8, 10, 16 or 0 to detect it automatically similar to user space strtoul(3).

       Return Number of characters consumed on success. Must be positive but no more than *buf_len*.

              -EINVAL if no valid digits were found or

unsupported base was provided.

-ERANGE if resulting value was out of range.

void *bpf_sk_storage_get(struct bpf_map *_map_**, void ***_sk_**, void ***_value_**, u64** _flags_**)**

>   **Description**
>>   Get a bpf-local-storage from a _sk_.
>>
>>   Logically, it could be thought of getting the value from a _map_ with _sk_ as the **key**. From this perspective, the usage is not much different from **bpf_map_lookup_elem**(_map_, **&**_sk_) except this helper enforces the key must be a full socket and the map must be a **BPF_MAP_TYPE_SK_STORAGE** also.
>>
>>   Underneath, the value is stored locally at _sk_ instead of the _map_. The _map_ is used as the bpf-local-storage "type". The bpf-local-storage "type" (i.e. the _map_) is searched against all bpf-local-storages residing at _sk_.
>>
>>   _sk_ is a kernel **struct sock** pointer for LSM program. _sk_ is a **struct bpf_sock** pointer for other program types.
>>
>>   An optional _flags_ (**BPF_SK_STORAGE_GET_F_CREATE**) can be used such that a new bpf-local-storage will be created if one does not exist. _value_ can be used together with **BPF_SK_STORAGE_GET_F_CREATE** to specify the initial value of a bpf-local-storage. If _value_ is **NULL**, the new bpf-local-storage will be zero initialized.
>
>   **Return** A bpf-local-storage pointer is returned on success.
>>
>>   **NULL** if not found or there was an error in adding a new bpf-local-storage.

long bpf_sk_storage_delete(struct bpf_map *_map_**, void ***_sk_**)**

>   **Description**
>>   Delete a bpf-local-storage from a _sk_.
>
>   **Return** 0 on success.
>>
>>   **-ENOENT** if the bpf-local-storage cannot be found.

-EINVAL if sk is not a fullsock (e.g. a
request_sock).

long bpf_send_signal(u32 *sig*)

>   Description
>   Send signal *sig* to the process of the current task.
>   The signal may be delivered to any of this
>   process's threads.

>   Return 0 on success or successfully queued.

>   **-EBUSY** if work queue under nmi is full.

>   **-EINVAL** if *sig* is invalid.

>   **-EPERM** if no permission to send the *sig*.

>   **-EAGAIN** if bpf program can try again.

**s64 bpf_tcp_gen_syncookie(void *_sk_, void *_iph_, u32 _iph_len_,
struct tcphdr *_th_, u32 _th_len_)**

>   Description
>   Try to issue a SYN cookie for the packet with
>   corresponding IP/TCP headers, *iph* and *th*, on the
>   listening socket in *sk*.

>   *iph* points to the start of the IPv4 or IPv6 header,
>   while *iph_len* contains **sizeof**(**struct iphdr**) or
>   **sizeof**(**struct ipv6hdr**).

>   *th* points to the start of the TCP header, while
>   *th_len* contains the length of the TCP header with
>   options (at least **sizeof**(**struct tcphdr**)).

>   Return On success, lower 32 bits hold the generated SYN
>   cookie in followed by 16 bits which hold the MSS
>   value for that cookie, and the top 16 bits are
>   unused.

>   On failure, the returned value is one of the
>   following:

>   **-EINVAL** SYN cookie cannot be issued due to error

>   **-ENOENT** SYN cookie should not be issued (no SYN
>   flood)

                    -EOPNOTSUPP kernel configuration does not enable
                    SYN cookies

                    -EPROTONOSUPPORT IP packet version is not 4 or 6

long bpf_skb_output(void **ctx***, struct bpf_map ****map***, u64** *flags***,
void ****data***, u64** *size***)**

      Description
            Write raw *data* blob into a special BPF perf event
            held by *map* of type **BPF_MAP_TYPE_PERF_EVENT_ARRAY**.
            This perf event must have the following attributes:
            PERF_SAMPLE_RAW as **sample_type**, **PERF_TYPE_SOFTWARE**
            as **type**, and **PERF_COUNT_SW_BPF_OUTPUT** as **config**.

            The *flags* are used to indicate the index in *map* for
            which the value must be put, masked with
            BPF_F_INDEX_MASK. Alternatively, *flags* can be set
            to **BPF_F_CURRENT_CPU** to indicate that the index of
            the current CPU core should be used.

            The value to write, of *size*, is passed through eBPF
            stack and pointed by *data*.

            ctx is a pointer to in-kernel struct sk_buff.

            This helper is similar to **bpf_perf_event_output**()
            but restricted to raw_tracepoint bpf programs.

      Return 0 on success, or a negative error in case of
            failure.

long bpf_probe_read_user(void **dst***, u32** *size***, const void
****unsafe_ptr***)**

      Description
            Safely attempt to read *size* bytes from user space
            address *unsafe_ptr* and store the data in *dst*.

      Return 0 on success, or a negative error in case of
            failure.

long bpf_probe_read_kernel(void **dst***, u32** *size***, const void
****unsafe_ptr***)**

      Description
            Safely attempt to read *size* bytes from kernel space

address *unsafe_ptr* and store the data in *dst*.

Return 0 on success, or a negative error in case of
          failure.

long bpf_probe_read_user_str(void **dst***, u32** *size***, const void**
unsafe_ptr*)**

Description
          Copy a NUL terminated string from an unsafe user
          address *unsafe_ptr* to *dst*. The *size* should include
          the terminating NUL byte. In case the string length
          is smaller than *size*, the target is not padded with
          further NUL bytes. If the string length is larger
          than *size*, just *size*-1 bytes are copied and the
          last byte is set to NUL.

          On success, returns the number of bytes that were
          written, including the terminal NUL. This makes
          this helper useful in tracing programs for reading
          strings, and more importantly to get its length at
          runtime. See the following snippet:

```
SEC("kprobe/sys_open")
void bpf_sys_open(struct pt_regs *ctx)
{
        char buf[PATHLEN]; // PATHLEN is defined to 256
        int res = bpf_probe_read_user_str(buf, sizeof(buf),
                                          ctx->di);

        // Consume buf, for example push it to
        // userspace via bpf_perf_event_output(); we
        // can use res (the string length) as event
        // size, after checking its boundaries.
}
```

          In comparison, using **bpf_probe_read_user**() helper
          here instead to read the string would require to
          estimate the length at compile time, and would
          often result in copying more memory than necessary.

          Another useful use case is when parsing individual
          process arguments or individual environment
          variables navigating *current*->**mm**->**arg_start** and
          current->**mm**->**env_start**: using this helper and the
          return value, one can quickly iterate at the right
          offset of the memory area.

> **Return** On success, the strictly positive length of the
> output string, including the trailing NUL
> character. On error, a negative value.

long bpf_probe_read_kernel_str(void **dst***, u32** *size***, const void ****unsafe_ptr***)**

> **Description**
> Copy a NUL terminated string from an unsafe kernel
> address *unsafe_ptr* to *dst*. Same semantics as with
> **bpf_probe_read_user_str**() apply.
>
> **Return** On success, the strictly positive length of the
> string, including the trailing NUL character. On
> error, a negative value.

long bpf_tcp_send_ack(void **tp***, u32** *rcv_nxt***)**

> **Description**
> Send out a tcp-ack. *tp* is the in-kernel struct
> **tcp_sock**. *rcv_nxt* is the ack_seq to be sent out.
>
> **Return** 0 on success, or a negative error in case of
> failure.

long bpf_send_signal_thread(u32 *sig***)**

> **Description**
> Send signal *sig* to the thread corresponding to the
> current task.
>
> **Return** 0 on success or successfully queued.
>
> > **-EBUSY** if work queue under nmi is full.
> >
> > **-EINVAL** if *sig* is invalid.
> >
> > **-EPERM** if no permission to send the *sig*.
> >
> > **-EAGAIN** if bpf program can try again.

u64 bpf_jiffies64(void)

> **Description**
> Obtain the 64bit jiffies
>
> **Return** The 64 bit jiffies

long bpf_read_branch_records(struct bpf_perf_event_data **ctx***,
void ****buf***, u32** *size***, u64** *flags***)**

>   **Description**
>   For an eBPF program attached to a perf event,
>   retrieve the branch records (**struct
>   perf_branch_entry**) associated to *ctx* and store it
>   in the buffer pointed by *buf* up to size *size* bytes.
>
>   **Return** On success, number of bytes written to *buf*. On
>   error, a negative value.
>
>   The *flags* can be set to
>   **BPF_F_GET_BRANCH_RECORDS_SIZE** to instead return the
>   number of bytes required to store all the branch
>   entries. If this flag is set, *buf* may be NULL.
>
>   **-EINVAL** if arguments invalid or **size** not a multiple
>   of **sizeof**(**struct perf_branch_entry**).
>
>   **-ENOENT** if architecture does not support branch
>   records.

long bpf_get_ns_current_pid_tgid(u64 *dev***, u64** *ino***, struct
bpf_pidns_info ****nsdata***, u32** *size***)**

>   **Description**
>   Returns 0 on success, values for *pid* and *tgid* as
>   seen from the current *namespace* will be returned in
>   *nsdata*.
>
>   **Return** 0 on success, or one of the following in case of
>   failure:
>
>   **-EINVAL** if dev and inum supplied don't match dev_t
>   and inode number with nsfs of current task, or if
>   dev conversion to dev_t lost high bits.
>
>   **-ENOENT** if pidns does not exists for the current
>   task.

long bpf_xdp_output(void **ctx***, struct bpf_map ****map***, u64** *flags***,
void ****data***, u64** *size***)**

>   **Description**
>   Write raw *data* blob into a special BPF perf event
>   held by *map* of type **BPF_MAP_TYPE_PERF_EVENT_ARRAY**.
>   This perf event must have the following attributes:

PERF_SAMPLE_RAW as **sample_type**, **PERF_TYPE_SOFTWARE** as **type**, and **PERF_COUNT_SW_BPF_OUTPUT** as **config**.

The *flags* are used to indicate the index in *map* for which the value must be put, masked with **BPF_F_INDEX_MASK**. Alternatively, *flags* can be set to **BPF_F_CURRENT_CPU** to indicate that the index of the current CPU core should be used.

The value to write, of *size*, is passed through eBPF stack and pointed by *data*.

ctx is a pointer to in-kernel struct xdp_buff.

This helper is similar to **bpf_perf_eventoutput**() but restricted to raw_tracepoint bpf programs.

Return 0 on success, or a negative error in case of failure.

u64 bpf_get_netns_cookie(void **ctx***)**

Description

Retrieve the cookie (generated by the kernel) of the network namespace the input *ctx* is associated with. The network namespace cookie remains stable for its lifetime and provides a global identifier that can be assumed unique. If *ctx* is NULL, then the helper returns the cookie for the initial network namespace. The cookie itself is very similar to that of **bpf_get_socket_cookie**() helper, but for network namespaces instead of sockets.

Return A 8-byte long opaque number.

u64 bpf_get_current_ancestor_cgroup_id(int *ancestor_level***)**

Description

Return id of cgroup v2 that is ancestor of the cgroup associated with the current task at the *ancestor_level*. The root cgroup is at *ancestor_level* zero and each step down the hierarchy increments the level. If *ancestor_level* == level of cgroup associated with the current task, then return value will be the same as that of **bpf_get_current_cgroup_id**().

The helper is useful to implement policies based on

cgroups that are upper in hierarchy than immediate
cgroup associated with the current task.

The format of returned id and helper limitations
are same as in **bpf_get_current_cgroup_id**().

Return The id is returned or 0 in case the id could not be
retrieved.

long bpf_sk_assign(struct sk_buff *_skb_, void *_sk_, u64 _flags_)

> Description
> > Helper is overloaded depending on BPF program type.
> > This description applies to **BPF_PROG_TYPE_SCHED_CLS**
> > and **BPF_PROG_TYPE_SCHED_ACT** programs.
> >
> > Assign the _sk_ to the _skb_. When combined with
> > appropriate routing configuration to receive the
> > packet towards the socket, will cause _skb_ to be
> > delivered to the specified socket. Subsequent
> > redirection of _skb_ via **bpf_redirect**(),
> > **bpf_clone_redirect**() or other methods outside of
> > BPF may interfere with successful delivery to the
> > socket.
> >
> > This operation is only valid from TC ingress path.
> >
> > The _flags_ argument must be zero.

> Return 0 on success, or a negative error in case of
> failure:
> > **-EINVAL** if specified _flags_ are not supported.
> >
> > **-ENOENT** if the socket is unavailable for
> > assignment.
> >
> > **-ENETUNREACH** if the socket is unreachable (wrong
> > netns).
> >
> > **-EOPNOTSUPP** if the operation is not supported, for
> > example a call from outside of TC ingress.

**long bpf_sk_assign(struct bpf_sk_lookup *_ctx_, struct bpf_sock
*_sk_, u64 _flags_)**

> Description
> > Helper is overloaded depending on BPF program type.

This description applies to **BPF_PROG_TYPE_SK_LOOKUP** programs.

Select the *sk* as a result of a socket lookup.

For the operation to succeed passed socket must be compatible with the packet description provided by the *ctx* object.

L4 protocol (**IPPROTO_TCP** or **IPPROTO_UDP**) must be an exact match. While IP family (**AF_INET** or **AF_INET6**) must be compatible, that is IPv6 sockets that are not v6-only can be selected for IPv4 packets.

Only TCP listeners and UDP unconnected sockets can be selected. *sk* can also be NULL to reset any previous selection.

flags argument can combination of following values:

- **BPF_SK_LOOKUP_F_REPLACE** to override the previous socket selection, potentially done by a BPF program that ran before us.

- **BPF_SK_LOOKUP_F_NO_REUSEPORT** to skip load-balancing within reuseport group for the socket being selected.

On success *ctx->sk* will point to the selected socket.

Return 0 on success, or a negative errno in case of failure.

- **-EAFNOSUPPORT** if socket family (*sk->family*) is not compatible with packet family (*ctx->family*).

- **-EEXIST** if socket has been already selected, potentially by another program, and **BPF_SK_LOOKUP_F_REPLACE** flag was not specified.

- **-EINVAL** if unsupported flags were specified.

- **-EPROTOTYPE** if socket L4 protocol (*sk->protocol*) doesn't match packet protocol (*ctx->protocol*).

- **-ESOCKTNOSUPPORT** if socket is not in allowed state (TCP listening or UDP unconnected).

u64 bpf_ktime_get_boot_ns(void)

>   **Description**
>   Return the time elapsed since system boot, in
>   nanoseconds. Does include the time the system was
>   suspended. See: **clock_gettime**(**CLOCK_BOOTTIME**)
>
>   **Return** Current *ktime*.

long bpf_seq_printf(struct seq_file **m**, **const char** ****fmt**, **u32**
*fmt_size***, const void** ****data***, u32** *data_len***)**

>   **Description**
>   **bpf_seq_printf**() uses seq_file **seq_printf**() to
>   print out the format string. The *m* represents the
>   seq_file. The *fmt* and *fmt_size* are for the format
>   string itself. The *data* and *data_len* are format
>   string arguments. The *data* are a **u64** array and
>   corresponding format string values are stored in
>   the array. For strings and pointers where pointees
>   are accessed, only the pointer values are stored in
>   the *data* array. The *data_len* is the size of *data*
>   in bytes - must be a multiple of 8.
>
>   Formats **%s**, **%p{i,I}{4,6}** requires to read kernel
>   memory. Reading kernel memory may fail due to
>   either invalid address or valid address but
>   requiring a major memory fault. If reading kernel
>   memory fails, the string for **%s** will be an empty
>   string, and the ip address for **%p{i,I}{4,6}** will be
>   0. Not returning error to bpf program is consistent
>   with what **bpf_trace_printk**() does for now.
>
>   **Return** 0 on success, or a negative error in case of
>   failure:
>
>   **-EBUSY** if per-CPU memory copy buffer is busy, can
>   try again by returning 1 from bpf program.
>
>   **-EINVAL** if arguments are invalid, or if *fmt* is
>   invalid/unsupported.
>
>   **-E2BIG** if *fmt* contains too many format specifiers.
>
>   **-EOVERFLOW** if an overflow happened: The same object
>   will be tried again.

long bpf_seq_write(struct seq_file *m, **const void ****data*, **u32** *len***)**

>   **Description**
>   **bpf_seq_write**() uses seq_file **seq_write**() to write
>   the data. The *m* represents the seq_file. The *data*
>   and *len* represent the data to write in bytes.
>
>   **Return** 0 on success, or a negative error in case of
>   failure:
>
>   **-EOVERFLOW** if an overflow happened: The same object
>   will be tried again.

u64 bpf_sk_cgroup_id(void **sk***)**

>   **Description**
>   Return the cgroup v2 id of the socket *sk*.
>
>   *sk* must be a non-**NULL** pointer to a socket, e.g. one
>   returned from **bpf_sk_lookup_xxx**(),
>   **bpf_sk_fullsock**(), etc. The format of returned id
>   is same as in **bpf_skb_cgroup_id**().
>
>   This helper is available only if the kernel was
>   compiled with the **CONFIG_SOCK_CGROUP_DATA**
>   configuration option.
>
>   **Return** The id is returned or 0 in case the id could not be
>   retrieved.

u64 bpf_sk_ancestor_cgroup_id(void **sk***, int** *ancestor_level***)**

>   **Description**
>   Return id of cgroup v2 that is ancestor of cgroup
>   associated with the *sk* at the *ancestor_level*. The
>   root cgroup is at *ancestor_level* zero and each step
>   down the hierarchy increments the level. If
>   *ancestor_level* == level of cgroup associated with
>   *sk*, then return value will be same as that of
>   **bpf_sk_cgroup_id**().
>
>   The helper is useful to implement policies based on
>   cgroups that are upper in hierarchy than immediate
>   cgroup associated with *sk*.
>
>   The format of returned id and helper limitations
>   are same as in **bpf_sk_cgroup_id**().

> **Return** The id is returned or 0 in case the id could not be retrieved.

long bpf_ringbuf_output(void *_ringbuf_, void *_data_, u64 _size_, u64 _flags_)

> **Description**
>> Copy _size_ bytes from _data_ into a ring buffer _ringbuf_. If **BPF_RB_NO_WAKEUP** is specified in _flags_, no notification of new data availability is sent. If **BPF_RB_FORCE_WAKEUP** is specified in _flags_, notification of new data availability is sent unconditionally. If **0** is specified in _flags_, an adaptive notification of new data availability is sent.
>>
>> An adaptive notification is a notification sent whenever the user-space process has caught up and consumed all available payloads. In case the user-space process is still processing a previous payload, then no notification is needed as it will process the newly added payload automatically.
>
> **Return** 0 on success, or a negative error in case of failure.

void *bpf_ringbuf_reserve(void *_ringbuf_, u64 _size_, u64 _flags_)

> **Description**
>> Reserve _size_ bytes of payload in a ring buffer _ringbuf_. _flags_ must be 0.
>
> **Return** Valid pointer with _size_ bytes of memory available; NULL, otherwise.

void bpf_ringbuf_submit(void *_data_, u64 _flags_)

> **Description**
>> Submit reserved ring buffer sample, pointed to by _data_. If **BPF_RB_NO_WAKEUP** is specified in _flags_, no notification of new data availability is sent. If **BPF_RB_FORCE_WAKEUP** is specified in _flags_, notification of new data availability is sent unconditionally. If **0** is specified in _flags_, an adaptive notification of new data availability is sent.
>>
>> See 'bpf_ringbuf_output()' for the definition of

adaptive notification.

> **Return** Nothing. Always succeeds.

void bpf_ringbuf_discard(void *_data_, u64 _flags_)

> **Description**
>> Discard reserved ring buffer sample, pointed to by
>> _data_. If **BPF_RB_NO_WAKEUP** is specified in _flags_,
>> no notification of new data availability is sent.
>> If **BPF_RB_FORCE_WAKEUP** is specified in _flags_,
>> notification of new data availability is sent
>> unconditionally. If **0** is specified in _flags_, an
>> adaptive notification of new data availability is
>> sent.
>>
>> See 'bpf_ringbuf_output()' for the definition of
>> adaptive notification.

> **Return** Nothing. Always succeeds.

u64 bpf_ringbuf_query(void *_ringbuf_, u64 _flags_)

> **Description**
>> Query various characteristics of provided ring
>> buffer. What exactly is queries is determined by
>> _flags_:
>>
>> - **BPF_RB_AVAIL_DATA**: Amount of data not yet
>>   consumed.
>>
>> - **BPF_RB_RING_SIZE**: The size of ring buffer.
>>
>> - **BPF_RB_CONS_POS**: Consumer position (can wrap
>>   around).
>>
>> - **BPF_RB_PROD_POS**: Producer(s) position (can wrap
>>   around).
>>
>> Data returned is just a momentary snapshot of
>> actual values and could be inaccurate, so this
>> facility should be used to power heuristics and for
>> reporting, not to make 100% correct calculation.

> **Return** Requested value, or 0, if _flags_ are not recognized.

long bpf_csum_level(struct sk_buff *_skb_, u64 _level_)

Description

Change the skbs checksum level by one layer up or down, or reset it entirely to none in order to have the stack perform checksum validation. The level is applicable to the following protocols: TCP, UDP, GRE, SCTP, FCOE. For example, a decap of | ETH | IP | UDP | GUE | IP | TCP | into | ETH | IP | TCP | through **bpf_skb_adjust_room**() helper with passing in **BPF_F_ADJ_ROOM_NO_CSUM_RESET** flag would require one call to **bpf_csum_level**() with **BPF_CSUM_LEVEL_DEC** since the UDP header is removed. Similarly, an encap of the latter into the former could be accompanied by a helper call to **bpf_csum_level**() with **BPF_CSUM_LEVEL_INC** if the skb is still intended to be processed in higher layers of the stack instead of just egressing at tc.

There are three supported level settings at this time:

- **BPF_CSUM_LEVEL_INC**: Increases skb->csum_level for skbs with CHECKSUM_UNNECESSARY.

- **BPF_CSUM_LEVEL_DEC**: Decreases skb->csum_level for skbs with CHECKSUM_UNNECESSARY.

- **BPF_CSUM_LEVEL_RESET**: Resets skb->csum_level to 0 and sets CHECKSUM_NONE to force checksum validation by the stack.

- **BPF_CSUM_LEVEL_QUERY**: No-op, returns the current skb->csum_level.

Return 0 on success, or a negative error in case of failure. In the case of **BPF_CSUM_LEVEL_QUERY**, the current skb->csum_level is returned or the error code -EACCES in case the skb is not subject to CHECKSUM_UNNECESSARY.

struct tcp6_sock *bpf_skc_to_tcp6_sock(void *_sk_)

Description

Dynamically cast a _sk_ pointer to a _tcp6_sock_ pointer.

Return _sk_ if casting is valid, or **NULL** otherwise.

struct tcp_sock *bpf_skc_to_tcp_sock(void *_sk_)

Description

Dynamically cast a *sk* pointer to a *tcp_sock* pointer.

Return *sk* if casting is valid, or **NULL** otherwise.

struct tcp_timewait_sock *bpf_skc_to_tcp_timewait_sock(void **sk***)**

Description

Dynamically cast a *sk* pointer to a *tcp_timewait_sock* pointer.

Return *sk* if casting is valid, or **NULL** otherwise.

struct tcp_request_sock *bpf_skc_to_tcp_request_sock(void **sk***)**

Description

Dynamically cast a *sk* pointer to a *tcp_request_sock* pointer.

Return *sk* if casting is valid, or **NULL** otherwise.

struct udp6_sock *bpf_skc_to_udp6_sock(void **sk***)**

Description

Dynamically cast a *sk* pointer to a *udp6_sock* pointer.

Return *sk* if casting is valid, or **NULL** otherwise.

long bpf_get_task_stack(struct task_struct **task***, void ****buf***, u32 ***size***, u64 ***flags***)**

Description

Return a user or a kernel stack in bpf program provided buffer. Note: the user stack will only be populated if the *task* is the current task; all other tasks will return -EOPNOTSUPP. To achieve this, the helper needs *task*, which is a valid pointer to **struct task_struct**. To store the stacktrace, the bpf program provides *buf* with a nonnegative *size*.

The last argument, *flags*, holds the number of stack frames to skip (from 0 to 255), masked with **BPF_F_SKIP_FIELD_MASK**. The next bits can be used to set the following flags:

BPF_F_USER_STACK
> Collect a user space stack instead of a
> kernel stack. The *task* must be the current
> task.

BPF_F_USER_BUILD_ID
> Collect buildid+offset instead of ips for
> user stack, only valid if **BPF_F_USER_STACK**
> is also specified.

> **bpf_get_task_stack**() can collect up to
> **PERF_MAX_STACK_DEPTH** both kernel and user frames,
> subject to sufficient large buffer size. Note that
> this limit can be controlled with the **sysctl**
> program, and that it should be manually increased
> in order to profile long user stacks (such as
> stacks for Java programs). To do so, use:

> # sysctl kernel.perf_event_max_stack=<new value>

Return The non-negative copied *buf* length equal to or less
> than *size* on success, or a negative error in case
> of failure.

**long bpf_load_hdr_opt(struct bpf_sock_ops *_skops_, void
*_searchby_res_, u32 _len_, u64 _flags_)**

Description
> Load header option. Support reading a particular
> TCP header option for bpf program
> (**BPF_PROG_TYPE_SOCK_OPS**).

> If *flags* is 0, it will search the option from the
> *skops*->**skb_data**. The comment in **struct
> bpf_sock_ops** has details on what skb_data contains
> under different *skops*->**op**.

> The first byte of the *searchby_res* specifies the
> kind that it wants to search.

> If the searching kind is an experimental kind (i.e.
> 253 or 254 according to RFC6994). It also needs to
> specify the "magic" which is either 2 bytes or 4
> bytes. It then also needs to specify the size of
> the magic by using the 2nd byte which is
> "kind-length" of a TCP header option and the
> "kind-length" also includes the first 2 bytes

"kind" and "kind-length" itself as a normal TCP header option also does.

For example, to search experimental kind 254 with 2 byte magic 0xeB9F, the searchby_res should be [254, 4, 0xeB, 0x9F, 0, 0, ... 0].

To search for the standard window scale option (3), the *searchby_res* should be [3, 0, 0, ... 0]. Note, kind-length must be 0 for regular option.

Searching for No-Op (0) and End-of-Option-List (1) are not supported.

len must be at least 2 bytes which is the minimal size of a header option.

Supported flags:

- **BPF_LOAD_HDR_OPT_TCP_SYN** to search from the saved_syn packet or the just-received syn packet.

Return > 0 when found, the header option is copied to *searchby_res*. The return value is the total length copied. On failure, a negative error code is returned:

-EINVAL if a parameter is invalid.

-ENOMSG if the option is not found.

-ENOENT if no syn packet is available when **BPF_LOAD_HDR_OPT_TCP_SYN** is used.

-ENOSPC if there is not enough space. Only *len* number of bytes are copied.

-EFAULT on failure to parse the header options in the packet.

-EPERM if the helper cannot be used under the current *skops***->op**.

long bpf_store_hdr_opt(struct bpf_sock_ops **skops***, const void ****from***, u32** *len***, u64** *flags***)**

Description

Store header option. The data will be copied from

buffer *from* with length *len* to the TCP header.

The buffer *from* should have the whole option that
includes the kind, kind-length, and the actual
option data. The *len* must be at least kind-length
long. The kind-length does not have to be 4 byte
aligned. The kernel will take care of the padding
and setting the 4 bytes aligned value to th->doff.

This helper will check for duplicated option by
searching the same option in the outgoing skb.

This helper can only be called during
BPF_SOCK_OPS_WRITE_HDR_OPT_CB.

Return 0 on success, or negative error in case of failure:

> **-EINVAL** If param is invalid.
>
> **-ENOSPC** if there is not enough space in the header.
> Nothing has been written
>
> **-EEXIST** if the option already exists.
>
> **-EFAULT** on failure to parse the existing header
> options.
>
> **-EPERM** if the helper cannot be used under the
> current *skops->op*.

**long bpf_reserve_hdr_opt(struct bpf_sock_ops *_skops_, u32 _len_, u64
flags)**

> **Description**
>
> Reserve *len* bytes for the bpf header option. The
> space will be used by **bpf_store_hdr_opt**() later in
> **BPF_SOCK_OPS_WRITE_HDR_OPT_CB**.
>
> If **bpf_reserve_hdr_opt**() is called multiple times,
> the total number of bytes will be reserved.
>
> This helper can only be called during
> **BPF_SOCK_OPS_HDR_OPT_LEN_CB**.

Return 0 on success, or negative error in case of failure:

> **-EINVAL** if a parameter is invalid.

-ENOSPC if there is not enough space in the header.

-EPERM if the helper cannot be used under the current *skops->***op**.

void *bpf_inode_storage_get(struct bpf_map **map***, void ****inode***, void ****value***, u64** *flags***)**

> ### Description
>
> Get a bpf_local_storage from an *inode*.
>
> Logically, it could be thought of as getting the value from a *map* with *inode* as the **key**. From this perspective, the usage is not much different from **bpf_map_lookup_elem**(*map***, &***inode*) except this helper enforces the key must be an inode and the map must also be a **BPF_MAP_TYPE_INODE_STORAGE**.
>
> Underneath, the value is stored locally at *inode* instead of the *map*. The *map* is used as the bpf-local-storage "type". The bpf-local-storage "type" (i.e. the *map*) is searched against all bpf_local_storage residing at *inode*.
>
> An optional *flags* (**BPF_LOCAL_STORAGE_GET_F_CREATE**) can be used such that a new bpf_local_storage will be created if one does not exist. *value* can be used together with **BPF_LOCAL_STORAGE_GET_F_CREATE** to specify the initial value of a bpf_local_storage. If *value* is **NULL**, the new bpf_local_storage will be zero initialized.
>
> ### Return
> A bpf_local_storage pointer is returned on success.
>
> **NULL** if not found or there was an error in adding a new bpf_local_storage.

int bpf_inode_storage_delete(struct bpf_map **map***, void ****inode***)**

> ### Description
>
> Delete a bpf_local_storage from an *inode*.
>
> ### Return
> 0 on success.
>
> **-ENOENT** if the bpf_local_storage cannot be found.

long bpf_d_path(struct path **path***, char ****buf***, u32** *sz***)**

Description

Return full path for given **struct path** object, which needs to be the kernel BTF *path* object. The path is returned in the provided buffer *buf* of size *sz* and is zero terminated.

Return

On success, the strictly positive length of the string, including the trailing NUL character. On error, a negative value.

long bpf_copy_from_user(void **dst***, u32** *size***, const void ****user_ptr***)**

Description

Read *size* bytes from user space address *user_ptr* and store the data in *dst*. This is a wrapper of **copy_from_user()**.

Return

0 on success, or a negative error in case of failure.

long bpf_snprintf_btf(char **str***, u32** *str_size***, struct btf_ptr ****ptr***, u32** *btf_ptr_size***, u64** *flags***)**

Description

Use BTF to store a string representation of *ptr*->ptr in *str*, using *ptr*->type_id. This value should specify the type that *ptr*->ptr points to. LLVM __builtin_btf_type_id(type, 1) can be used to look up vmlinux BTF type ids. Traversing the data structure using BTF, the type information and values are stored in the first *str_size* - 1 bytes of *str*. Safe copy of the pointer data is carried out to avoid kernel crashes during operation. Smaller types can use string space on the stack; larger programs can use map data to store the string representation.

The string can be subsequently shared with userspace via bpf_perf_event_output() or ring buffer interfaces. bpf_trace_printk() is to be avoided as it places too small a limit on string size to be useful.

flags is a combination of

BTF_F_COMPACT
no formatting around type information

> **BTF_F_NONAME**
> > no struct/union member names/types
>
> **BTF_F_PTR_RAW**
> > show raw (unobfuscated) pointer values;
> > equivalent to printk specifier %px.
>
> **BTF_F_ZERO**
> > show zero-valued struct/union members; they
> > are not displayed by default

> **Return** The number of bytes that were written (or would
> > have been written if output had to be truncated due
> > to string size), or a negative error in cases of
> > failure.

**long bpf_seq_printf_btf(struct seq_file *_m_, struct btf_ptr *_ptr_,
u32 _ptr_size_, u64 _flags_)**

> **Description**
> > Use BTF to write to seq_write a string
> > representation of _ptr_->ptr, using _ptr_->type_id as
> > per bpf_snprintf_btf(). _flags_ are identical to
> > those used for bpf_snprintf_btf.

> **Return** 0 on success or a negative error in case of
> > failure.

u64 bpf_skb_cgroup_classid(struct sk_buff *_skb_)

> **Description**
> > See **bpf_get_cgroup_classid**() for the main
> > description. This helper differs from
> > **bpf_get_cgroup_classid**() in that the cgroup v1
> > net_cls class is retrieved only from the _skb_'s
> > associated socket instead of the current process.

> **Return** The id is returned or 0 in case the id could not be
> > retrieved.

**long bpf_redirect_neigh(u32 _ifindex_, struct bpf_redir_neigh
*_params_, int _plen_, u64 _flags_)**

> **Description**
> > Redirect the packet to another net device of index
> > _ifindex_ and fill in L2 addresses from neighboring
> > subsystem. This helper is somewhat similar to

bpf_redirect(), except that it populates L2
addresses as well, meaning, internally, the helper
relies on the neighbor lookup for the L2 address of
the nexthop.

The helper will perform a FIB lookup based on the
skb's networking header to get the address of the
next hop, unless this is supplied by the caller in
the *params* argument. The *plen* argument indicates
the len of *params* and should be set to 0 if *params*
is NULL.

The *flags* argument is reserved and must be 0. The
helper is currently only supported for tc BPF
program types, and enabled for IPv4 and IPv6
protocols.

Return The helper returns **TC_ACT_REDIRECT** on success or
TC_ACT_SHOT on error.

void *bpf_per_cpu_ptr(const void **percpu_ptr***, u32** *cpu***)**

Description
Take a pointer to a percpu ksym, *percpu_ptr*, and
return a pointer to the percpu kernel variable on
cpu. A ksym is an extern variable decorated with
'__ksym'. For ksym, there is a global var (either
static or global) defined of the same name in the
kernel. The ksym is percpu if the global var is
percpu. The returned pointer points to the global
percpu var on *cpu*.

bpf_per_cpu_ptr() has the same semantic as
per_cpu_ptr() in the kernel, except that
bpf_per_cpu_ptr() may return NULL. This happens if
cpu is larger than nr_cpu_ids. The caller of
bpf_per_cpu_ptr() must check the returned value.

Return A pointer pointing to the kernel percpu variable on
cpu, or NULL, if *cpu* is invalid.

void *bpf_this_cpu_ptr(const void **percpu_ptr***)**

Description
Take a pointer to a percpu ksym, *percpu_ptr*, and
return a pointer to the percpu kernel variable on
this cpu. See the description of 'ksym' in
bpf_per_cpu_ptr().

bpf_this_cpu_ptr() has the same semantic as
this_cpu_ptr() in the kernel. Different from
bpf_per_cpu_ptr(), it would never return NULL.

Return A pointer pointing to the kernel percpu variable on
this cpu.

long bpf_redirect_peer(u32 *ifindex,* **u64** *flags***)**

Description
Redirect the packet to another net device of index
ifindex. This helper is somewhat similar to
bpf_redirect(), except that the redirection happens
to the *ifindex*' peer device and the netns switch
takes place from ingress to ingress without going
through the CPU's backlog queue.

The *flags* argument is reserved and must be 0. The
helper is currently only supported for tc BPF
program types at the ingress hook and for veth and
netkit target device types. The peer device must
reside in a different network namespace.

Return The helper returns **TC_ACT_REDIRECT** on success or
TC_ACT_SHOT on error.

**void *bpf_task_storage_get(struct bpf_map ***map,* **struct
task_struct ***task,* **void ***value,* **u64** *flags***)**

Description
Get a bpf_local_storage from the *task*.

Logically, it could be thought of as getting the
value from a *map* with *task* as the **key**. From this
perspective, the usage is not much different from
bpf_map_lookup_elem(*map,* &*task*) except this helper
enforces the key must be a task_struct and the map
must also be a **BPF_MAP_TYPE_TASK_STORAGE**.

Underneath, the value is stored locally at *task*
instead of the *map*. The *map* is used as the
bpf-local-storage "type". The bpf-local-storage
"type" (i.e. the *map*) is searched against all
bpf_local_storage residing at *task*.

An optional *flags* (**BPF_LOCAL_STORAGE_GET_F_CREATE**)
can be used such that a new bpf_local_storage will

be created if one does not exist. *value* can be
used together with **BPF_LOCAL_STORAGE_GET_F_CREATE**
to specify the initial value of a
bpf_local_storage. If *value* is **NULL**, the new
bpf_local_storage will be zero initialized.

Return A bpf_local_storage pointer is returned on success.

NULL if not found or there was an error in adding a
new bpf_local_storage.

long bpf_task_storage_delete(struct bpf_map **map***, struct
task_struct ****task***)**

Description
Delete a bpf_local_storage from a *task*.

Return 0 on success.

-ENOENT if the bpf_local_storage cannot be found.

struct task_struct *bpf_get_current_task_btf(void)

Description
Return a BTF pointer to the "current" task. This
pointer can also be used in helpers that accept an
ARG_PTR_TO_BTF_ID of type *task_struct*.

Return Pointer to the current task.

long bpf_bprm_opts_set(struct linux_binprm **bprm***, u64** *flags***)**

Description
Set or clear certain options on *bprm*:

BPF_F_BPRM_SECUREEXEC Set the secureexec bit which
sets the **AT_SECURE** auxv for glibc. The bit is
cleared if the flag is not specified.

Return **-EINVAL** if invalid *flags* are passed, zero
otherwise.

u64 bpf_ktime_get_coarse_ns(void)

Description
Return a coarse-grained version of the time elapsed
since system boot, in nanoseconds. Does not include
time the system was suspended.

See: **clock_gettime**(**CLOCK_MONOTONIC_COARSE**)

Return Current *ktime*.

long bpf_ima_inode_hash(struct inode *inode, void *dst, u32 size)

Description
Returns the stored IMA hash of the *inode* (if it's available). If the hash is larger than *size*, then only *size* bytes will be copied to *dst*

Return The **hash_algo** is returned on success, **-EOPNOTSUP** if IMA is disabled or **-EINVAL** if invalid arguments are passed.

struct socket *bpf_sock_from_file(struct file *file)

Description
If the given file represents a socket, returns the associated socket.

Return A pointer to a struct socket on success or NULL if the file is not a socket.

long bpf_check_mtu(void *ctx, u32 ifindex, u32 *mtu_len, s32 len_diff, u64 flags)

Description
Check packet size against exceeding MTU of net device (based on *ifindex*). This helper will likely be used in combination with helpers that adjust/change the packet size.

The argument *len_diff* can be used for querying with a planned size change. This allows to check MTU prior to changing packet ctx. Providing a *len_diff* adjustment that is larger than the actual packet size (resulting in negative packet size) will in principle not exceed the MTU, which is why it is not considered a failure. Other BPF helpers are needed for performing the planned size change; therefore the responsibility for catching a negative packet size belongs in those helpers.

Specifying *ifindex* zero means the MTU check is performed against the current net device. This is practical if this isn't used prior to redirect.

On input *mtu_len* must be a valid pointer, else
verifier will reject BPF program. If the value
mtu_len is initialized to zero then the ctx packet
size is use. When value *mtu_len* is provided as
input this specify the L3 length that the MTU check
is done against. Remember XDP and TC length operate
at L2, but this value is L3 as this correlate to
MTU and IP-header tot_len values which are L3
(similar behavior as bpf_fib_lookup).

The Linux kernel route table can configure MTUs on
a more specific per route level, which is not
provided by this helper. For route level MTU
checks use the **bpf_fib_lookup**() helper.

ctx is either **struct xdp_md** for XDP programs or
struct sk_buff for tc cls_act programs.

The *flags* argument can be a combination of one or
more of the following values:

BPF_MTU_CHK_SEGS
        This flag will only works for *ctx* **struct
        sk_buff**. If packet context contains extra
        packet segment buffers (often knows as GSO
        skb), then MTU check is harder to check at
        this point, because in transmit path it is
        possible for the skb packet to get
        re-segmented (depending on net device
        features). This could still be a MTU
        violation, so this flag enables performing
        MTU check against segments, with a different
        violation return code to tell it apart.
        Check cannot use len_diff.

On return *mtu_len* pointer contains the MTU value of
the net device. Remember the net device configured
MTU is the L3 size, which is returned here and XDP
and TC length operate at L2. Helper take this into
account for you, but remember when using MTU value
in your BPF-code.

Return

- 0 on success, and populate MTU value in *mtu_len*
  pointer.

- < 0 if any input argument is invalid (*mtu_len* not
  updated)

  MTU violations return positive values, but also
  populate MTU value in *mtu_len* pointer, as this can
  be needed for implementing PMTU handing:

  - **BPF_MTU_CHK_RET_FRAG_NEEDED**

  - **BPF_MTU_CHK_RET_SEGS_TOOBIG**

long bpf_for_each_map_elem(struct bpf_map *map, **void
****callback_fn,* **void ****callback_ctx,* **u64** *flags***)**

> **Description**
>
> > For each element in **map**, call **callback_fn** function
> > with **map**, **callback_ctx** and other map-specific
> > parameters. The **callback_fn** should be a static
> > function and the **callback_ctx** should be a pointer
> > to the stack. The **flags** is used to control certain
> > aspects of the helper. Currently, the **flags** must
> > be 0.
> >
> > The following are a list of supported map types and
> > their respective expected callback signatures:
> >
> > BPF_MAP_TYPE_HASH, BPF_MAP_TYPE_PERCPU_HASH,
> > BPF_MAP_TYPE_LRU_HASH,
> > BPF_MAP_TYPE_LRU_PERCPU_HASH, BPF_MAP_TYPE_ARRAY,
> > BPF_MAP_TYPE_PERCPU_ARRAY
> >
> > long (*callback_fn)(struct bpf_map *map, const void
> > *key, void *value, void *ctx);
> >
> > For per_cpu maps, the map_value is the value on the
> > cpu where the bpf_prog is running.
> >
> > If **callback_fn** return 0, the helper will continue
> > to the next element. If return value is 1, the
> > helper will skip the rest of elements and return.
> > Other return values are not used now.
>
> **Return** The number of traversed map elements for success,
> > **-EINVAL** for invalid **flags**.

long bpf_snprintf(char *str, **u32** *str_size*, **const char ****fmt*, **u64
****data*, **u32** *data_len***)**

Description

Outputs a string into the **str** buffer of size **str_size** based on a format string stored in a read-only map pointed by **fmt**.

Each format specifier in **fmt** corresponds to one u64 element in the **data** array. For strings and pointers where pointees are accessed, only the pointer values are stored in the *data* array. The *data_len* is the size of *data* in bytes - must be a multiple of 8.

Formats **%s** and **%p{i,I}{4,6}** require to read kernel memory. Reading kernel memory may fail due to either invalid address or valid address but requiring a major memory fault. If reading kernel memory fails, the string for **%s** will be an empty string, and the ip address for **%p{i,I}{4,6}** will be 0. Not returning error to bpf program is consistent with what **bpf_trace_printk**() does for now.

Return The strictly positive length of the formatted string, including the trailing zero character. If the return value is greater than **str_size**, **str** contains a truncated string, guaranteed to be zero-terminated except when **str_size** is 0.

Or **-EBUSY** if the per-CPU memory copy buffer is busy.

long bpf_sys_bpf(u32 *cmd*, **void ****attr*, **u32** *attr_size***)**

Description

Execute bpf syscall with given arguments.

Return A syscall result.

long bpf_btf_find_by_name_kind(char *name, **int** *name_sz*, **u32** *kind*,
int *flags***)**

Description

Find BTF type with given name and kind in vmlinux BTF or in module's BTFs.

Return Returns btf_id and btf_obj_fd in lower and upper 32 bits.

long bpf_sys_close(u32 *fd***)**

       Description
            Execute close syscall for given FD.

       Return A syscall result.

long bpf_timer_init(struct bpf_timer **timer***, struct bpf_map ****map***,**
u64 *flags***)**

       Description
            Initialize the timer. First 4 bits of *flags*
            specify clockid. Only CLOCK_MONOTONIC,
            CLOCK_REALTIME, CLOCK_BOOTTIME are allowed. All
            other bits of *flags* are reserved. The verifier
            will reject the program if *timer* is not from the
            same *map*.

       Return 0 on success. **-EBUSY** if *timer* is already
            initialized. **-EINVAL** if invalid *flags* are passed.
            -EPERM if *timer* is in a map that doesn't have any
            user references. The user space should either hold
            a file descriptor to a map with timers or pin such
            map in bpffs. When map is unpinned or file
            descriptor is closed all timers in the map will be
            cancelled and freed.

long bpf_timer_set_callback(struct bpf_timer **timer***, void**
callback_fn*)**

       Description
            Configure the timer to call *callback_fn* static
            function.

       Return 0 on success. **-EINVAL** if *timer* was not initialized
            with bpf_timer_init() earlier. **-EPERM** if *timer* is
            in a map that doesn't have any user references.
            The user space should either hold a file descriptor
            to a map with timers or pin such map in bpffs. When
            map is unpinned or file descriptor is closed all
            timers in the map will be cancelled and freed.

long bpf_timer_start(struct bpf_timer **timer***, u64** *nsecs***, u64**
*flags***)**

       Description
            Set timer expiration N nanoseconds from the current
            time. The configured callback will be invoked in

soft irq context on some cpu and will not repeat
unless another bpf_timer_start() is made. In such
case the next invocation can migrate to a different
cpu. Since struct bpf_timer is a field inside map
element the map owns the timer. The
bpf_timer_set_callback() will increment refcnt of
BPF program to make sure that callback_fn code
stays valid. When user space reference to a map
reaches zero all timers in a map are cancelled and
corresponding program's refcnts are decremented.
This is done to make sure that Ctrl-C of a user
process doesn't leave any timers running. If map is
pinned in bpffs the callback_fn can re-arm itself
indefinitely. bpf_map_update/delete_elem() helpers
and user space sys_bpf commands cancel and free the
timer in the given map element. The map can
contain timers that invoke callback_fn-s from
different programs. The same callback_fn can serve
different timers from different maps if key/value
layout matches across maps. Every
bpf_timer_set_callback() can have different
callback_fn.

flags can be one of:

BPF_F_TIMER_ABS
        Start the timer in absolute expire value
        instead of the default relative one.

BPF_F_TIMER_CPU_PIN
        Timer will be pinned to the CPU of the
        caller.

Return 0 on success. **-EINVAL** if *timer* was not initialized
        with bpf_timer_init() earlier or invalid *flags* are
        passed.

long bpf_timer_cancel(struct bpf_timer *_timer_)

Description
        Cancel the timer and wait for callback_fn to finish
        if it was running.

Return 0 if the timer was not active. 1 if the timer was
        active. **-EINVAL** if *timer* was not initialized with
        bpf_timer_init() earlier. **-EDEADLK** if callback_fn
        tried to call bpf_timer_cancel() on its own timer
        which would have led to a deadlock otherwise.

u64 bpf_get_func_ip(void **ctx***)**

> **Description**
>> Get address of the traced function (for tracing and kprobe programs).
>>
>> When called for kprobe program attached as uprobe it returns probe address for both entry and return uprobe.
>
> **Return** Address of the traced function for kprobe. 0 for kprobes placed within the function (not at the entry). Address of the probe for uprobe and return uprobe.

u64 bpf_get_attach_cookie(void **ctx***)**

> **Description**
>> Get bpf_cookie value provided (optionally) during the program attachment. It might be different for each individual attachment, even if BPF program itself is the same. Expects BPF program context *ctx* as a first argument.
>>
>> **Supported for the following program types:**
>>
>>> • kprobe/uprobe;
>>>
>>> • tracepoint;
>>>
>>> • perf_event.
>
> **Return** Value specified by user at BPF link creation/attachment time or 0, if it was not specified.

long bpf_task_pt_regs(struct task_struct **task***)**

> **Description**
>> Get the struct pt_regs associated with **task**.
>
> **Return** A pointer to struct pt_regs.

long bpf_get_branch_snapshot(void **entries***, u32** *size***, u64** *flags***)**

> **Description**
>> Get branch trace from hardware engines like Intel

LBR. The hardware engine is stopped shortly after
the helper is called. Therefore, the user need to
filter branch entries based on the actual use case.
To capture branch trace before the trigger point of
the BPF program, the helper should be called at the
beginning of the BPF program.

The data is stored as struct perf_branch_entry into
output buffer *entries*. *size* is the size of *entries*
in bytes. *flags* is reserved for now and must be
zero.

Return On success, number of bytes written to *buf*. On
error, a negative value.

-EINVAL if *flags* is not zero.

-ENOENT if architecture does not support branch
records.

long bpf_trace_vprintk(const char **fmt***, u32** *fmt_size***, const void**
data*, u32** *data_len***)**

Description
Behaves like **bpf_trace_printk**() helper, but takes
an array of u64 to format and can handle more
format args as a result.

Arguments are to be used as in **bpf_seq_printf**()
helper.

Return The number of bytes written to the buffer, or a
negative error in case of failure.

struct unix_sock *bpf_skc_to_unix_sock(void **sk***)**

Description
Dynamically cast a *sk* pointer to a *unix_sock*
pointer.

Return *sk* if casting is valid, or **NULL** otherwise.

long bpf_kallsyms_lookup_name(const char **name***, int** *name_sz***, int**
*flags***, u64 ****res***)**

Description
Get the address of a kernel symbol, returned in
res. *res* is set to 0 if the symbol is not found.

Return On success, zero. On error, a negative value.

> **-EINVAL** if *flags* is not zero.
>
> **-EINVAL** if string *name* is not the same size as
> *name_sz*.
>
> **-ENOENT** if symbol is not found.
>
> **-EPERM** if caller does not have permission to obtain
> kernel address.

long bpf_find_vma(struct task_struct *****task***, u64** *addr*****, void**
*****callback_fn*****, void *****callback_ctx*****, u64** *flags*****)**

> Description
>
> > Find vma of *task* that contains *addr*, call
> > *callback_fn* function with *task*, *vma*, and
> > *callback_ctx*. The *callback_fn* should be a static
> > function and the *callback_ctx* should be a pointer
> > to the stack. The *flags* is used to control certain
> > aspects of the helper. Currently, the *flags* must
> > be 0.
> >
> > The expected callback signature is
> >
> > long (*callback_fn)(struct task_struct *task,
> > struct vm_area_struct *vma, void *callback_ctx);
>
> Return 0 on success. **-ENOENT** if *task->mm* is NULL, or no
> > vma contains *addr*. **-EBUSY** if failed to try lock
> > mmap_lock. **-EINVAL** for invalid **flags**.

long bpf_loop(u32 *nr_loops*****, void *****callback_fn*****, void**
*****callback_ctx*****, u64** *flags*****)**

> Description
>
> > For **nr_loops**, call **callback_fn** function with
> > **callback_ctx** as the context parameter. The
> > **callback_fn** should be a static function and the
> > **callback_ctx** should be a pointer to the stack. The
> > **flags** is used to control certain aspects of the
> > helper. Currently, the **flags** must be 0. Currently,
> > nr_loops is limited to 1 << 23 (~8 million) loops.
> >
> > long (*callback_fn)(u32 index, void *ctx);

where **index** is the current index in the loop. The
index is zero-indexed.

If **callback_fn** returns 0, the helper will continue
to the next loop. If return value is 1, the helper
will skip the rest of the loops and return. Other
return values are not used now, and will be
rejected by the verifier.

Return The number of loops performed, **-EINVAL** for invalid
flags, **-E2BIG** if **nr_loops** exceeds the maximum
number of loops.

long bpf_strncmp(const char *_s1_, u32 _s1_sz_, const char *_s2_)

Description
Do strncmp() between **s1** and **s2**. **s1** doesn't need to
be null-terminated and **s1_sz** is the maximum storage
size of **s1**. **s2** must be a read-only string.

Return An integer less than, equal to, or greater than
zero if the first **s1_sz** bytes of **s1** is found to be
less than, to match, or be greater than **s2**.

long bpf_get_func_arg(void *_ctx_, u32 _n_, u64 *_value_)

Description
Get **n**-th argument register (zero based) of the
traced function (for tracing programs) returned in
value.

Return 0 on success. **-EINVAL** if n >= argument register
count of traced function.

long bpf_get_func_ret(void *_ctx_, u64 *_value_)

Description
Get return value of the traced function (for
tracing programs) in **value**.

Return 0 on success. **-EOPNOTSUPP** for tracing programs
other than BPF_TRACE_FEXIT or BPF_MODIFY_RETURN.

long bpf_get_func_arg_cnt(void *_ctx_)

Description
Get number of registers of the traced function (for
tracing programs) where function arguments are

stored in these registers.

Return The number of argument registers of the traced
function.

int bpf_get_retval(void)

Description
Get the BPF program's return value that will be
returned to the upper layers.

This helper is currently supported by cgroup
programs and only by the hooks where BPF program's
return value is returned to the userspace via
errno.

Return The BPF program's return value.

int bpf_set_retval(int *retval*)

Description
Set the BPF program's return value that will be
returned to the upper layers.

This helper is currently supported by cgroup
programs and only by the hooks where BPF program's
return value is returned to the userspace via
errno.

Note that there is the following corner case where
the program exports an error via bpf_set_retval but
signals success via 'return 1':
    bpf_set_retval(-EPERM); return 1;

In this case, the BPF program's return value will
use helper's -EPERM. This still holds true for
cgroup/bind{4,6} which supports extra 'return 3'
success case.

Return 0 on success, or a negative error in case of
failure.

u64 bpf_xdp_get_buff_len(struct xdp_buff *_xdp_md_)

Description
Get the total size of a given xdp buff (linear and
paged area)

          Return The total size of a given xdp buffer.

long bpf_xdp_load_bytes(struct xdp_buff **xdp_md**, u32 *offset***, void
****buf**, u32 *len***)**

          Description
                 This helper is provided as an easy way to load data
                 from a xdp buffer. It can be used to load *len* bytes
                 from *offset* from the frame associated to *xdp_md*,
                 into the buffer pointed by *buf*.

          Return 0 on success, or a negative error in case of
                 failure.

long bpf_xdp_store_bytes(struct xdp_buff **xdp_md**, u32 *offset***,
void ****buf**, u32 *len***)**

          Description
                 Store *len* bytes from buffer *buf* into the frame
                 associated to *xdp_md*, at *offset*.

          Return 0 on success, or a negative error in case of
                 failure.

long bpf_copy_from_user_task(void **dst**, u32 *size***, const void
****user_ptr**, struct task_struct ****tsk**, u64 *flags***)**

          Description
                 Read *size* bytes from user space address *user_ptr* in
                 tsk's address space, and stores the data in *dst*.
                 flags is not used yet and is provided for future
                 extensibility. This helper can only be used by
                 sleepable programs.

          Return 0 on success, or a negative error in case of
                 failure. On error *dst* buffer is zeroed out.

long bpf_skb_set_tstamp(struct sk_buff **skb**, u64 *tstamp***, u32
*tstamp_type***)**

          Description
                 Change the __sk_buff->tstamp_type to *tstamp_type*
                 and set *tstamp* to the __sk_buff->tstamp together.

                 If there is no need to change the
                 __sk_buff->tstamp_type, the tstamp value can be
                 directly written to __sk_buff->tstamp instead.

BPF_SKB_TSTAMP_DELIVERY_MONO is the only tstamp
that will be kept during bpf_redirect_*(). A non
zero *tstamp* must be used with the
BPF_SKB_TSTAMP_DELIVERY_MONO *tstamp_type*.

A BPF_SKB_TSTAMP_UNSPEC *tstamp_type* can only be
used with a zero *tstamp*.

Only IPv4 and IPv6 skb->protocol are supported.

This function is most useful when it needs to set a
mono delivery time to __sk_buff->tstamp and then
bpf_redirect_*() to the egress of an iface. For
example, changing the (rcv) timestamp in
__sk_buff->tstamp at ingress to a mono delivery
time and then bpf_redirect_*() to *sch_fq@phy-dev*.

> **Return** 0 on success. **-EINVAL** for invalid input
> **-EOPNOTSUPP** for unsupported protocol

long bpf_ima_file_hash(struct file *_file_, void *_dst_, u32 _size_)

> **Description**
>> Returns a calculated IMA hash of the _file_. If the
>> hash is larger than _size_, then only _size_ bytes will
>> be copied to _dst_

> **Return** The **hash_algo** is returned on success, **-EOPNOTSUP** if
>> the hash calculation failed or **-EINVAL** if invalid
>> arguments are passed.

void *bpf_kptr_xchg(void *_map_value_, void *_ptr_)

> **Description**
>> Exchange kptr at pointer _map_value_ with _ptr_, and
>> return the old value. _ptr_ can be NULL, otherwise it
>> must be a referenced pointer which will be released
>> when this helper is called.

> **Return** The old value of kptr (which can be NULL). The
>> returned pointer if not NULL, is a reference which
>> must be released using its corresponding release
>> function, or moved into a BPF map before program
>> exit.

**void *bpf_map_lookup_percpu_elem(struct bpf_map *_map_, const void
*_key_, u32 _cpu_)**

> **Description**
>> Perform a lookup in *percpu map* for an entry
>> associated to *key* on *cpu*.
>
> **Return** Map value associated to *key* on *cpu*, or **NULL** if no
>> entry was found or *cpu* is invalid.

struct mptcp_sock *bpf_skc_to_mptcp_sock(void ***sk)**

> **Description**
>> Dynamically cast a *sk* pointer to a *mptcp_sock*
>> pointer.
>
> **Return** *sk* if casting is valid, or **NULL** otherwise.

long bpf_dynptr_from_mem(void ***data, u32 ***size**, u64 *flags*, struct
bpf_dynptr ***ptr**)**

> **Description**
>> Get a dynptr to local memory *data*.
>>
>> *data* must be a ptr to a map value. The maximum
>> *size* supported is DYNPTR_MAX_SIZE. *flags* is
>> currently unused.
>
> **Return** 0 on success, -E2BIG if the size exceeds
>> DYNPTR_MAX_SIZE, -EINVAL if flags is not 0.

long bpf_ringbuf_reserve_dynptr(void ***ringbuf, u32 ***size**, u64
flags, struct bpf_dynptr ***ptr**)**

> **Description**
>> Reserve *size* bytes of payload in a ring buffer
>> *ringbuf* through the dynptr interface. *flags* must be
>> 0.
>>
>> Please note that a corresponding
>> bpf_ringbuf_submit_dynptr or
>> bpf_ringbuf_discard_dynptr must be called on *ptr*,
>> even if the reservation fails. This is enforced by
>> the verifier.
>
> **Return** 0 on success, or a negative error in case of
>> failure.

void bpf_ringbuf_submit_dynptr(struct bpf_dynptr ***ptr, u64 *flags*)**

> **Description**

Submit reserved ring buffer sample, pointed to by
data, through the dynptr interface. This is a no-op
if the dynptr is invalid/null.

For more information on *flags*, please see
'bpf_ringbuf_submit'.

Return Nothing. Always succeeds.

void bpf_ringbuf_discard_dynptr(struct bpf_dynptr **ptr***, u64**
*flags***)**

Description
Discard reserved ring buffer sample through the
dynptr interface. This is a no-op if the dynptr is
invalid/null.

For more information on *flags*, please see
'bpf_ringbuf_discard'.

Return Nothing. Always succeeds.

long bpf_dynptr_read(void **dst***, u32** *len***, const struct bpf_dynptr**
src*, u32** *offset***, u64** *flags***)**

Description
Read *len* bytes from *src* into *dst*, starting from
offset into *src*. *flags* is currently unused.

Return 0 on success, -E2BIG if *offset* + *len* exceeds the
length of *src*'s data, -EINVAL if *src* is an invalid
dynptr or if *flags* is not 0.

long bpf_dynptr_write(const struct bpf_dynptr **dst***, u32** *offset***,**
void **src***, u32** *len***, u64** *flags***)**

Description
Write *len* bytes from *src* into *dst*, starting from
offset into *dst*.

flags must be 0 except for skb-type dynptrs.

For skb-type dynptrs:

• All data slices of the dynptr are
automatically invalidated after
bpf_dynptr_write(). This is because
writing may pull the skb and change the

underlying packet buffer.

- For *flags*, please see the flags accepted
  by **bpf_skb_store_bytes**().

Return 0 on success, -E2BIG if *offset* + *len* exceeds the
length of *dst*'s data, -EINVAL if *dst* is an invalid
dynptr or if *dst* is a read-only dynptr or if *flags*
is not correct. For skb-type dynptrs, other errors
correspond to errors returned by
bpf_skb_store_bytes().

**void *bpf_dynptr_data(const struct bpf_dynptr *ptr, u32 offset,
u32 len)**

Description
Get a pointer to the underlying dynptr data.

len must be a statically known value. The returned
data slice is invalidated whenever the dynptr is
invalidated.

skb and xdp type dynptrs may not use
bpf_dynptr_data. They should instead use
bpf_dynptr_slice and bpf_dynptr_slice_rdwr.

Return Pointer to the underlying dynptr data, NULL if the
dynptr is read-only, if the dynptr is invalid, or
if the offset and length is out of bounds.

**s64 bpf_tcp_raw_gen_syncookie_ipv4(struct iphdr *iph, struct
tcphdr *th, u32 th_len)**

Description
Try to issue a SYN cookie for the packet with
corresponding IPv4/TCP headers, *iph* and *th*, without
depending on a listening socket.

iph points to the IPv4 header.

th points to the start of the TCP header, while
th_len contains the length of the TCP header (at
least **sizeof**(**struct tcphdr**)).

Return On success, lower 32 bits hold the generated SYN
cookie in followed by 16 bits which hold the MSS
value for that cookie, and the top 16 bits are
unused.

On failure, the returned value is one of the following:

-EINVAL if *th_len* is invalid.

s64 bpf_tcp_raw_gen_syncookie_ipv6(struct ipv6hdr **iph**, struct tcphdr ****th**, u32 **th_len**)**

>   **Description**
>   >   Try to issue a SYN cookie for the packet with corresponding IPv6/TCP headers, *iph* and *th*, without depending on a listening socket.
>   >
>   >   *iph* points to the IPv6 header.
>   >
>   >   *th* points to the start of the TCP header, while *th_len* contains the length of the TCP header (at least **sizeof**(**struct tcphdr**)).
>
>   **Return** On success, lower 32 bits hold the generated SYN cookie in followed by 16 bits which hold the MSS value for that cookie, and the top 16 bits are unused.
>   >
>   >   On failure, the returned value is one of the following:
>   >
>   >   **-EINVAL** if *th_len* is invalid.
>   >
>   >   **-EPROTONOSUPPORT** if CONFIG_IPV6 is not builtin.

long bpf_tcp_raw_check_syncookie_ipv4(struct iphdr **iph**, struct tcphdr ****th**)**

>   **Description**
>   >   Check whether *iph* and *th* contain a valid SYN cookie ACK without depending on a listening socket.
>   >
>   >   *iph* points to the IPv4 header.
>   >
>   >   *th* points to the TCP header.
>
>   **Return** 0 if *iph* and *th* are a valid SYN cookie ACK.
>   >
>   >   On failure, the returned value is one of the following:

-**EACCES** if the SYN cookie is not valid.

long bpf_tcp_raw_check_syncookie_ipv6(struct ipv6hdr *_iph_, struct tcphdr *_th_)

>   Description
>   Check whether _iph_ and _th_ contain a valid SYN cookie
>   ACK without depending on a listening socket.
>
>   _iph_ points to the IPv6 header.
>
>   _th_ points to the TCP header.
>
>   Return 0 if _iph_ and _th_ are a valid SYN cookie ACK.
>
>   On failure, the returned value is one of the
>   following:
>
>   -**EACCES** if the SYN cookie is not valid.
>
>   -**EPROTONOSUPPORT** if CONFIG_IPV6 is not builtin.

u64 bpf_ktime_get_tai_ns(void)

>   Description
>   A nonsettable system-wide clock derived from
>   wall-clock time but ignoring leap seconds. This
>   clock does not experience discontinuities and
>   backwards jumps caused by NTP inserting leap
>   seconds as CLOCK_REALTIME does.
>
>   See: **clock_gettime(CLOCK_TAI)**
>
>   Return Current _ktime_.

long bpf_user_ringbuf_drain(struct bpf_map *_map_, void *_callback_fn_, void *_ctx_, u64 _flags_)

>   Description
>   Drain samples from the specified user ring buffer,
>   and invoke the provided callback for each such
>   sample:
>
>   long (*callback_fn)(const struct bpf_dynptr
>   *dynptr, void *ctx);
>
>   If **callback_fn** returns 0, the helper will continue
>   to try and drain the next sample, up to a maximum

of BPF_MAX_USER_RINGBUF_SAMPLES samples. If the
return value is 1, the helper will skip the rest of
the samples and return. Other return values are not
used now, and will be rejected by the verifier.

Return The number of drained samples if no error was
encountered while draining samples, or 0 if no
samples were present in the ring buffer. If a
user-space producer was epoll-waiting on this map,
and at least one sample was drained, they will
receive an event notification notifying them of
available space in the ring buffer. If the
BPF_RB_NO_WAKEUP flag is passed to this function,
no wakeup notification will be sent. If the
BPF_RB_FORCE_WAKEUP flag is passed, a wakeup
notification will be sent even if no sample was
drained.

On failure, the returned value is one of the
following:

-EBUSY if the ring buffer is contended, and another
calling context was concurrently draining the ring
buffer.

-EINVAL if user-space is not properly tracking the
ring buffer due to the producer position not being
aligned to 8 bytes, a sample not being aligned to 8
bytes, or the producer position not matching the
advertised length of a sample.

-E2BIG if user-space has tried to publish a sample
which is larger than the size of the ring buffer,
or which cannot fit within a struct bpf_dynptr.

void *bpf_cgrp_storage_get(struct bpf_map *map**, struct cgroup
***cgroup**, void ***value**, u64** flags**)**

Description
Get a bpf_local_storage from the *cgroup*.

Logically, it could be thought of as getting the
value from a *map* with *cgroup* as the **key**. From this
perspective, the usage is not much different from
bpf_map_lookup_elem(*map*, &*cgroup*) except this
helper enforces the key must be a cgroup struct and
the map must also be a **BPF_MAP_TYPE_CGRP_STORAGE**.

In reality, the local-storage value is embedded
directly inside of the *cgroup* object itself, rather
than being located in the **BPF_MAP_TYPE_CGRP_STORAGE**
map. When the local-storage value is queried for
some *map* on a *cgroup* object, the kernel will
perform an O(n) iteration over all of the live
local-storage values for that *cgroup* object until
the local-storage value for the *map* is found.

An optional *flags* (**BPF_LOCAL_STORAGE_GET_F_CREATE**)
can be used such that a new bpf_local_storage will
be created if one does not exist. *value* can be
used together with **BPF_LOCAL_STORAGE_GET_F_CREATE**
to specify the initial value of a
bpf_local_storage. If *value* is **NULL**, the new
bpf_local_storage will be zero initialized.

Return A bpf_local_storage pointer is returned on success.

NULL if not found or there was an error in adding a
new bpf_local_storage.

long bpf_cgrp_storage_delete(struct bpf_map **map***, struct cgroup
****cgroup***)**

Description
Delete a bpf_local_storage from a *cgroup*.

Return 0 on success.

-ENOENT if the bpf_local_storage cannot be found.

EXAMPLES   top

Example usage for most of the eBPF helpers listed in this manual
page are available within the Linux kernel sources, at the
following locations:

- *samples/bpf/*

- *tools/testing/selftests/bpf/*

LICENSE   top

eBPF programs can have an associated license, passed along with
the bytecode instructions to the kernel when the programs are

loaded. The format for that string is identical to the one in use
for kernel modules (Dual licenses, such as "Dual BSD/GPL", may be
used). Some helper functions are only accessible to programs that
are compatible with the GNU General Public License (GNU GPL).

In order to use such helpers, the eBPF program must be loaded
with the correct license string passed (via **attr**) to the **bpf**()
system call, and this generally translates into the C source code
of the program containing a line similar to the following:

    char ____license[] __attribute__((section("license"), used)) = "GPL";

IMPLEMENTATION   top

This manual page is an effort to document the existing eBPF
helper functions. But as of this writing, the BPF sub-system is
under heavy development. New eBPF program or map types are added,
along with new helper functions. Some helpers are occasionally
made available for additional program types. So in spite of the
efforts of the community, this page might not be up-to-date. If
you want to check by yourself what helper functions exist in your
kernel, or what types of programs they can support, here are some
files among the kernel tree that you may be interested in:

- *include/uapi/linux/bpf.h* is the main BPF header. It contains
  the full list of all helper functions, as well as many other
  BPF definitions including most of the flags, structs or
  constants used by the helpers.

- *net/core/filter.c* contains the definition of most
  network-related helper functions, and the list of program types
  from which they can be used.

- *kernel/trace/bpf_trace.c* is the equivalent for most tracing
  program-related helpers.

- *kernel/bpf/verifier.c* contains the functions used to check that
  valid types of eBPF maps are used with a given helper function.

- *kernel/bpf/* directory contains other files in which additional
  helpers are defined (for cgroups, sockmaps, etc.).

- The bpftool utility can be used to probe the availability of
  helper functions on the system (as well as supported program
  and map types, and a number of other parameters). To do so, run
  bpftool feature probe (see **bpftool-feature**(8) for details). Add
  the **unprivileged** keyword to list features available to

unprivileged users.

Compatibility between helper functions and program types can generally be found in the files where helper functions are defined. Look for the **struct bpf_func_proto** objects and for functions returning them: these functions contain a list of helpers that a given program type can call. Note that the **default:** label of the **switch ... case** used to filter helpers can call other functions, themselves allowing access to additional helpers. The requirement for GPL license is also in those **struct bpf_func_proto**.

Compatibility between helper functions and map types can be found in the **check_map_func_compatibility**() function in file *kernel/bpf/verifier.c*.

Helper functions that invalidate the checks on **data** and **data_end** pointers for network processing are listed in function **bpf_helper_changes_pkt_data**() in file *net/core/filter.c*.

SEE ALSO   top

bpf(2), **bpftool**(8), cgroups(7), ip(8), perf_event_open(2), sendmsg(2), socket(7), tc-bpf(8)

COLOPHON   top

This page is part of the *man-pages* (Linux kernel and C library user-space interface documentation) project. Information about the project can be found at ⟨https://www.kernel.org/doc/man-pages/⟩. If you have a bug report for this manual page, see ⟨https://git.kernel.org/pub/scm/docs/man-pages/man-pages.git/tree/CONTRIBUTING⟩. This page was obtained from the tarball man-pages-6.9.1.tar.gz fetched from ⟨https://mirrors.edge.kernel.org/pub/linux/docs/man-pages/⟩ on 2024-06-26. If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date source for the page, or you have corrections or improvements to the information in this COLOPHON (which is *not* part of the original manual page), send a mail to man-pages@man7.org

Linux v6.9   2024-01-23   *BPF-HELPERS*(7)

Pages that refer to this page: bpf(2), capabilities(7)