



pathtofile / **bad-bpf** Public

Notifications

Fork **79**

Star **544**

Code

Issues

Pull requests **1**

Actions

Projects

Security

Insights

main ▾

Go to file <> Code ▾

About

A collection of eBPF programs demonstrating bad behavior, presented at DEF CON 29

- Readme
- BSD-3-Clause license
- Activity
- 544 stars
- 9 watching
- 79 forks

[Report repository](#)

.github/workflows		
blazesym @ 53253d4		
bpftool @ 8485b9f		
libbpf @ f11758a		
src		
tools		
vmlinux		
.clang-format		
.gitignore		
.gitmodules		
LICENSE		
README.md		

README BSD-3-Clause license

Releases **3**

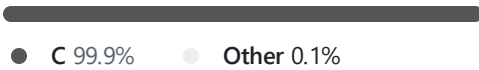
Release v1.2 Latest
on Jul 7

[+ 2 releases](#)

Packages

No packages published

Languages



Bad BPF

A collection of malicious eBPF programs that make use of eBPF's ability to read and write user data in between the usermode program and the kernel.

- [Overview](#)
- [To Build](#)
- [To Run](#)
- [Available Programs](#)

Overview

See my [blog](#) and my [DEF CON talk](#) for an overview on how these programs work and why this is interesting.

Examples have been tested on:

- Ubuntu 22.04

Build

To use pre-build binaries, grab them from the [Releases](#) page.

To build from source, do the following:

Dependencies

To build and run all the examples, you will need a Linux kernel version of at least [4.7](#).

As this code makes use of CO-RE, it requires a recent version of Linux that has BTF Type information. See [these notes in the libbpf README](#) for more information. For example Ubuntu requires `Ubuntu 20.10` +.

To build it requires these dependencies:

- zlib
- libelf
- libbfd
- clang and llvm 14
- make

On Ubuntu these can be installed by

```
sudo apt install build-essential clang-14 llvm-14
```

Build

To Build from source, recursively clone the repository then run `make` in the `src` directory to build:

```
# --recursive is needed to also get the libbpf :  
git clone --recursive https://github.com/pathtofile/bad-bpf  
cd bad-bpf/src  
make
```

The binaries will be built into `bad-bpf/src/`. If you encounter issues related to `vmlinux.h`, try remaking the file for your specific kernel and distribution:

```
cd bad-bpf/tools  
./bpftool btf dump file /sys/kernel/btf/vmlinux
```

Run

To run, launch each program as `root`. Every program has a `-help` option that has required arguments and examples.

Programs

Common Arguments

As well as `--help`, every program also has a `--target-pid / -t`. This option restricts the programs' operation to only programs that are children of the process matching this PID. This demonstrates to how affect some programs, but not others.

- [Bad BPF](#)
- [Overview](#)
- [Build](#)
 - [Dependencies](#)
 - [Build](#)
- [Run](#)
- [Programs](#)
 - [Common Arguments](#)
 - [BPF-Dos](#)
 - [Exec-Hijack](#)
 - [Pid-Hide](#)
 - [Sudo-Add](#)
 - [Write-Blocker](#)
 - [Text-Replace](#)
 - [Text-Replace2](#)
 - [Altering Configuration](#)
 - [Running Detached](#)

BPF-Dos

```
sudo ./bpfdos
```



This program raises a `SIG_KILL` signal to any program attempting to use the `ptrace` syscall, e.g. `strace`. Once bpf-dos starts you can test it by running:

```
strace /bin/whoami
```



Exec-Hijack

```
sudo ./exechejack
```



This program intercepts all `execve` calls (used to create new processes) and instead makes then call `/a`. To run, first ensure there is a program in the root dir `/a` (probably best to make is executable by all). `bad-bpf` builds a simple program `hijackee` that simply prints out the `uid` and `argv[0]`, so you can use that:

```
make  
sudo cp ./hijackee /a  
sudo chmod ugo+rx /a
```



Then just run `sudo ./exechejack`.

Pid-Hide

```
sudo ./pidhide --pid-to-hide 2222
```



This program hides the process matching this pid from tools such as `ps`.

It works by hooking the `getdents64` syscall, as `ps` works by looking for every sub-folder of `/proc/`. `PidHide` unlinks the folder matching the PID, so `ps` only sees the folders before and after it.

Sudo-Add

```
sudo ./sudoadd --username lowpriv-user
```



This program allows a normally low-privileged user to use `sudo` to become root.

It works by intercepting `sudo`'s reading of the `/etc/sudoers` file, and overwriting the first line with `<username> ALL=(ALL:ALL) NOPASSWD:ALL #`. This tricks `sudo` into thinking the user is allowed to become root. Other programs such as `cat` or `sudoedit` are unaffected, so to those programs the file is unchanged and the user does not have those privileges. The `#` at the end of the line ensures the rest of the line is treated as a comment, so it doesn't corrupt the file's logic.

Write-Blocker

```
sudo ./writeblocker --pid 508
```



This program intercepts all write syscall for a given process PID. Instead of passing the data to the actual write syscall, `writeblocker` will instead fake the call, returning the same number of bytes that the userspace program expects to be written.

Only File Descriptors > 2 will be blocked, so `stdin`, `stdout`, and `stderr` still work.

For example, if you block the writes for the `rsyslogd` process, ssh logins will not be written to `/var/log/auth.log`.


Text-Replace

```
sudo ./textreplace --filename /path/to/file --input
```



This program replaces all text matching `input` in the file with the `replace` text. This has a number of uses, for example:

To hide kernel module `joydev` from tools such as `lsmod` :

```
./textreplace -f /proc/modules -i 'joydev' -r ' ' 
```

Spoof the MAC address of the `eth0` interface:

```
/textreplace -f /sys/class/net/eth0/address -i 
```

Malware conducting anti-sandbox checks might check the MAC address to look for signs it is running inside a Virtual Machine or Sandbox, and not on a 'real' machine.

NOTE: Both `input` and `replace` must be the same length, to avoid adding NULL characters to the middle of a block of text.

To enter a newline from a bash prompt, use `'$\n'` , e.g. `--replace '$text\n'` .


Text-Replace2

This program works the same as `Text-Replace` , however it has two extra features:

- The program's configuration is alterable at runtime using eBPF Maps.
- The userspace loader can detach and exit

Altering Configuration

The filename is stored in the eBPF Map `map_filename` . The Key is always `0` , and the value matches this struct:

```
struct tr_file {  
    char filename[50];  
    unsigned int filename_len;  
}; 
```

That is, 50 ascii characters, then an unsigned int matching the length of the actual filename string.

The easiest way to view and alter eBPF maps is using `bpftool` :

```
# List current config
$> bpftool map dump name map_filename
[{"key": 0,
  "value": {
    "filename": "/proc/modules",
    "filename_len": 13
  }
}]

# Alter filename to be 'AAAA'
$> bpftool map update name map_filename \
  key hex 00 00 00 00 \
  value hex 61 61 61 61 00 00 00 00 00 00 00 00 (

# Confirm change config
$> bpftool map dump name map_filename
[{"key": 0,
  "value": {
    "filename": "aaaa",
    "filename_len": 4
  }
}]
```

To alter the text to find and replace, alter the items in the Map `map_text` . The text to find is at key `0` , and the text to replace is key `1` . The values will each match this struct:

```
struct tr_text {
  char text[20];
  unsigned int text_len;
};
```

Running Detached

By running the program with `--detach`, the userspace loader can exit without stopping the eBPF Programs. Before running, first make sure the bpf filesystem is mounted:

```
sudo mount bpffs -t bpf /sys/fs/bpf
```



Then you can run text-replace2 detached:

```
./textreplace2 -f /proc/modules -i 'joydev' -r
```



This will create a number of eBPF Link files under `/sys/fs/bpf/textreplace`. Once loader has successfully run, you can check the logs by running:

```
sudo cat /sys/kernel/debug/tracing/trace_pipe  
# confirm link files are there  
sudo ls -l /sys/fs/bpf/textreplace
```



Then to stop, simply delete the link files:

